



27/08/2020



**Barcelona  
Supercomputing  
Center**  
*Centro Nacional de Supercomputación*



EXCELENCIA  
SEVERO  
OCHOA



**esiwace**

CENTRE OF EXCELLENCE IN SIMULATION OF WEATHER  
AND CLIMATE IN EUROPE

# Computational Profiling Analysis for Climate and Weather

Mario C. Acosta and Xavier Yépes

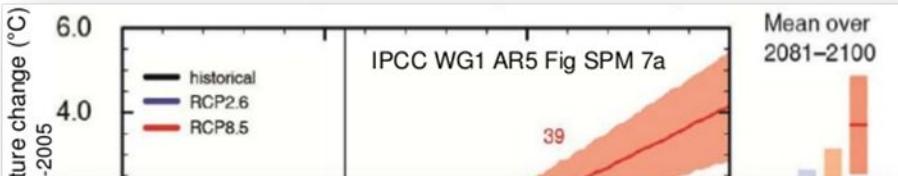
Summer School on Effective HPC

# Computational Profiling Analysis for Climate and Weather

- Objectives

- Define performance analysis fundamentals (objectives, methods, metrics, hardware counters, etc.)
- Define a methodology to study HPC performance for numerical models, know your enemy.
- Describe the BSC performance analysis tools suite (Exrae, Paraver, Dimemas)
- Interpret uses cases from Earth System Models (HARMONIE, IFS, NEMO, etc.) that illustrate how to identify and solve performance issues
- Apply profiling techniques to identify performance bottlenecks in your code
- Summarise typical performance problems
- Discuss specific knowledge about performance analysis applied to earth system modelling

# Introduction



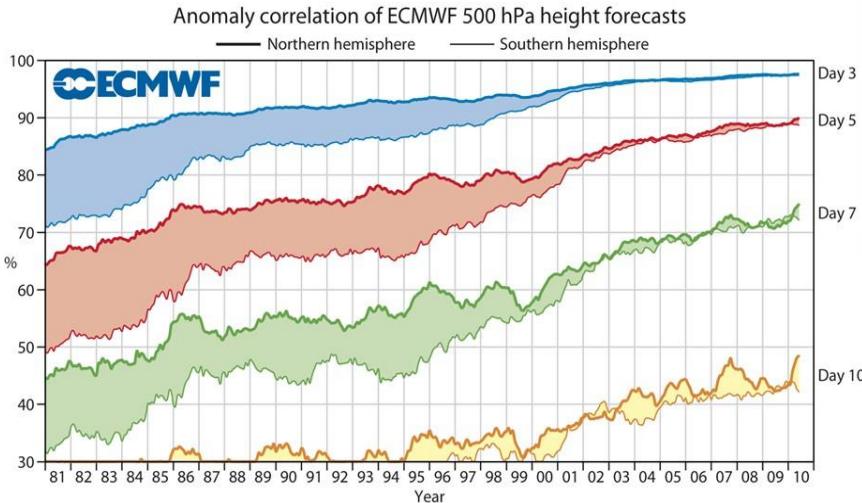
- « Projections
- « Impact analysis
- « Adaptation to climate change.

Currently, **only computational models have the potential** to provide geographically and physically consistent estimates.

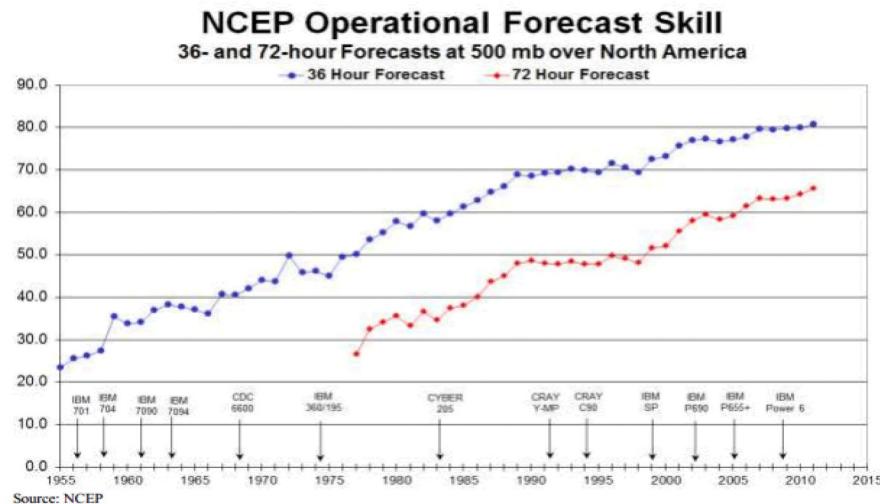
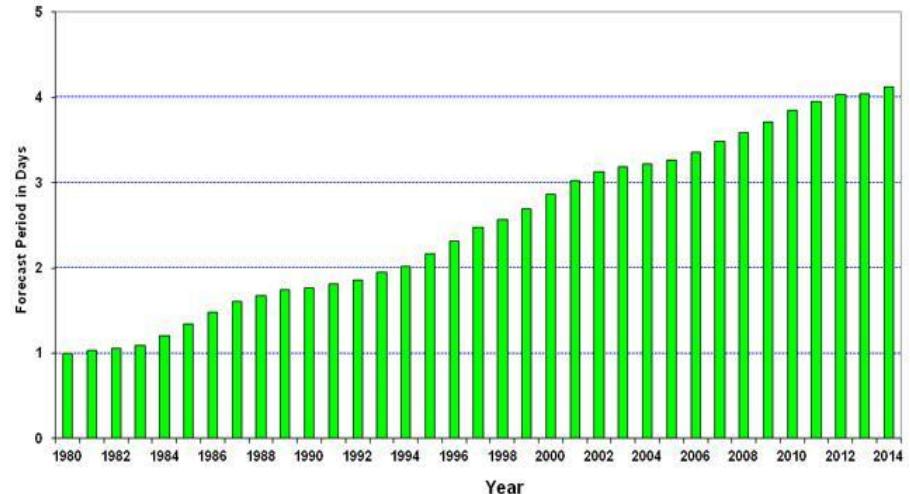


# Introduction

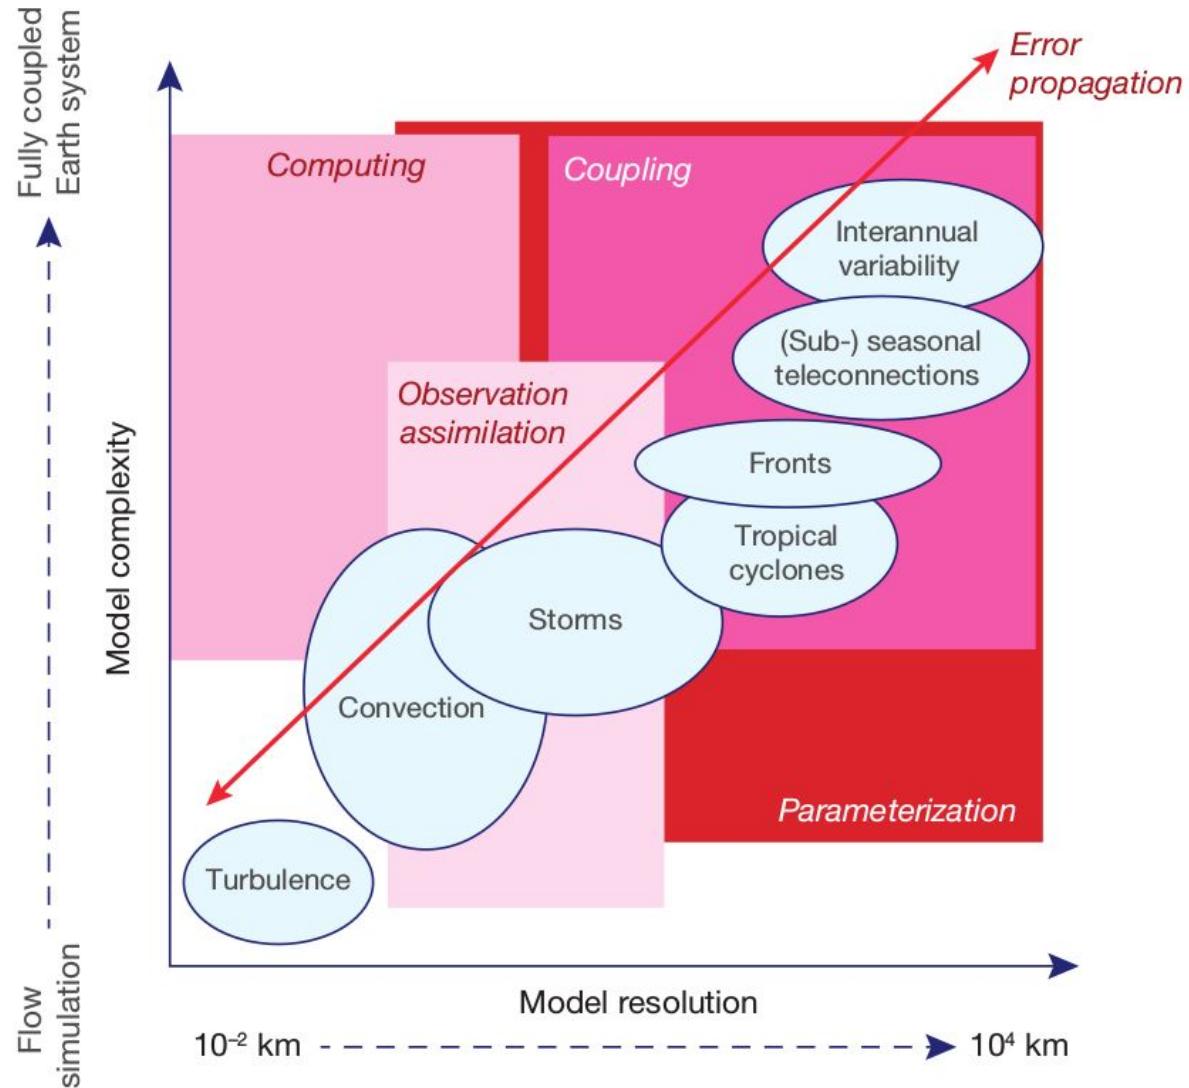
## Advances in Global and Regional Weather Forecasts



Accuracy of PMSL forecast (in days) compared to baseline of 1-day forecast in 1980



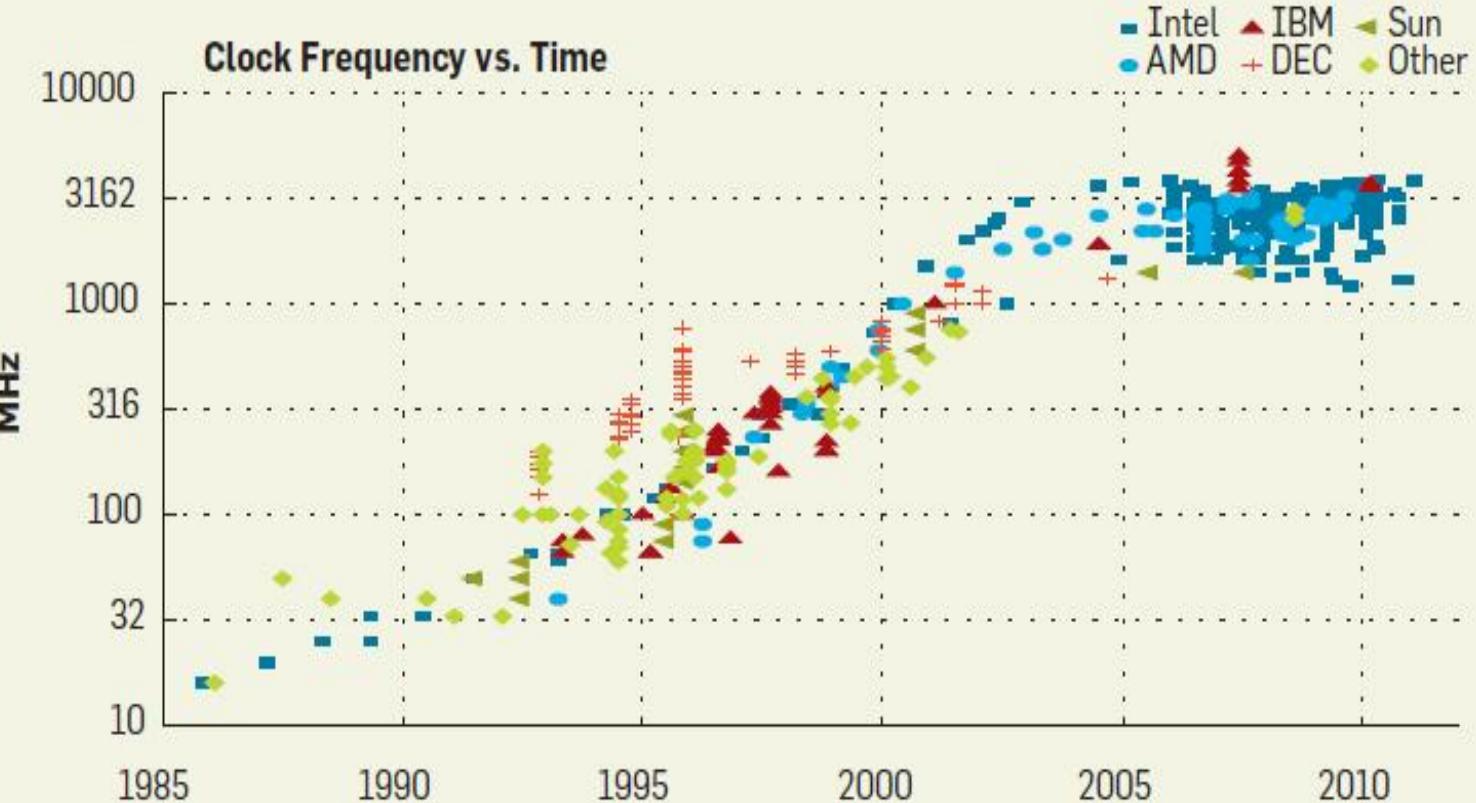
# Introduction



# Introduction



# Introduction

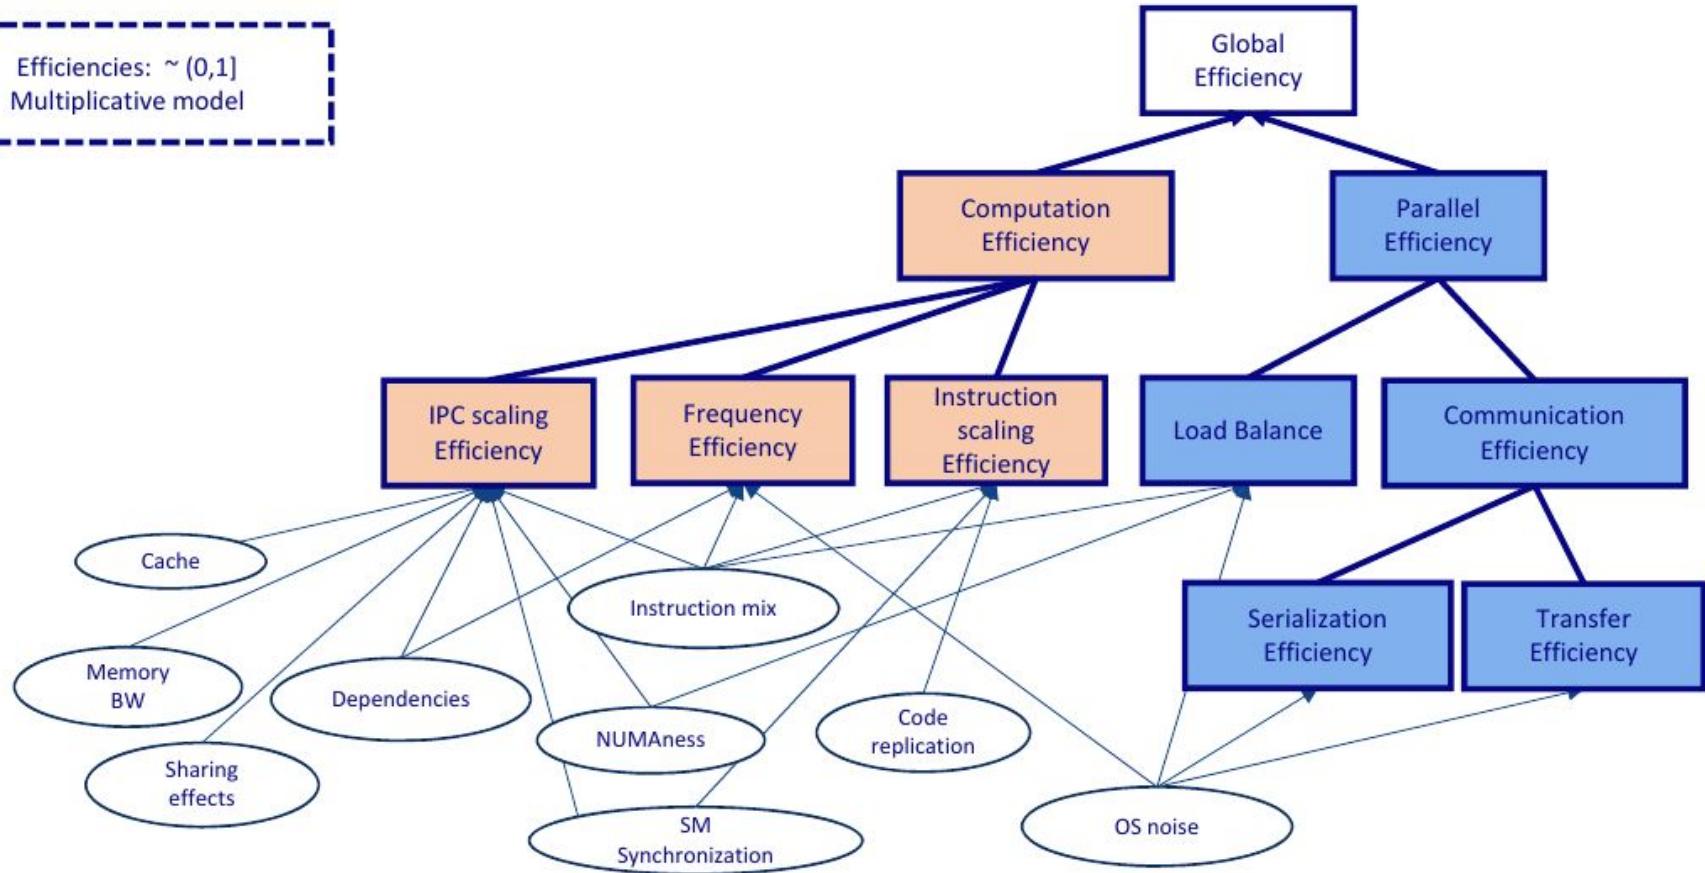


# Introduction

- To be able to use the computing power of modern supercomputers, applications must exploit parallelism.
- Parallelism produce overhead (extra computation and communications)
  - “*Overhead does not look a problem in my model*” → But if the needs increase (i.e. higher resolutions), a bad implementation will be a problem in some point.
  - We need a method to evaluate the parallelism efficiency of our computational models.
    - When the hardware change
    - When the number of resources change
    - When the model complexity change
    - When the resolution change
    - ...

# Introduction

Efficiencies:  $\sim (0,1]$   
Multiplicative model



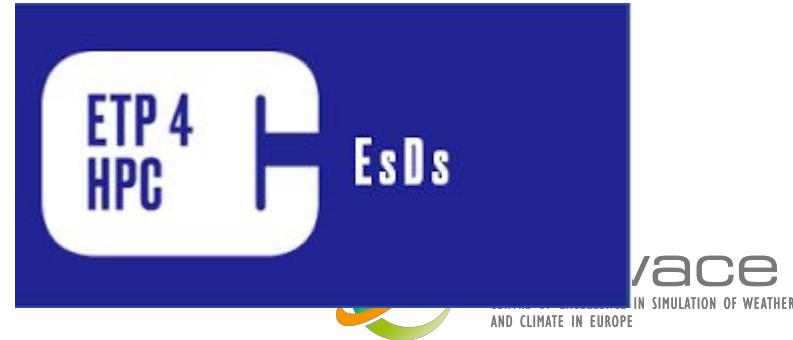
# Introduction

- The necessary refactoring of numerical codes is given a lot of attention and is stirring a number of discussions.
  - Computational performance analysis and new optimizations are needed for actual numerical models.
  - Study new algorithms for the new generation of high performance platforms (path to exascale).
- Several European institutions and projects working together on the same direction (ESCAPE2, ESiWACE2, IS-ENES3, ETP4HPC...)



**esiwace**  
CENTRE OF EXCELLENCE IN SIMULATION OF WEATHER  
AND CLIMATE IN EUROPE

**is-enes**  
INFRASTRUCTURE FOR THE EUROPEAN NETWORK  
FOR EARTH SYSTEM MODELLING

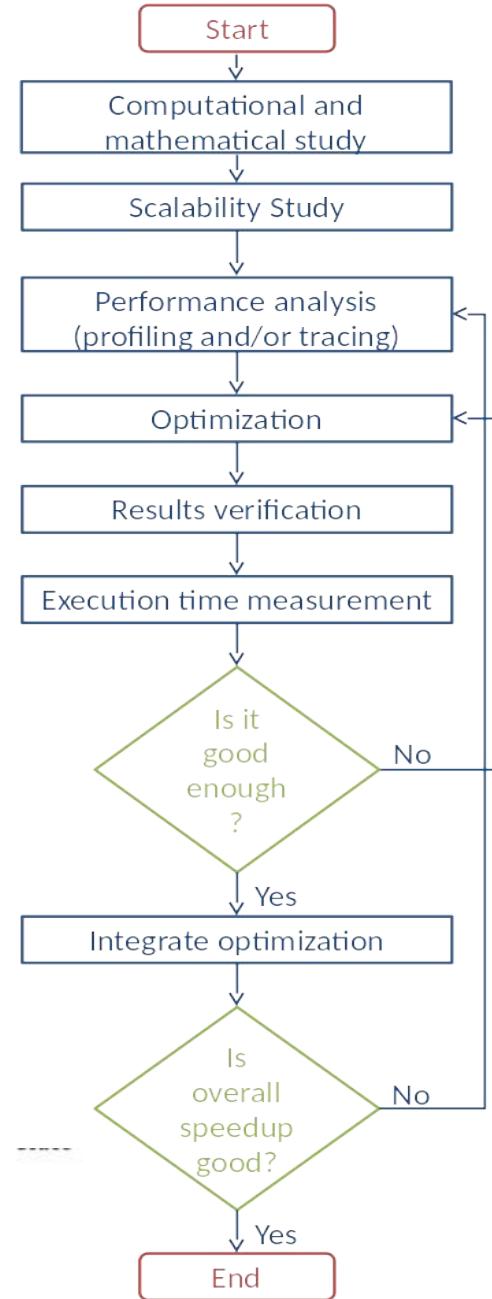


# CES-Performance Team & ES Department

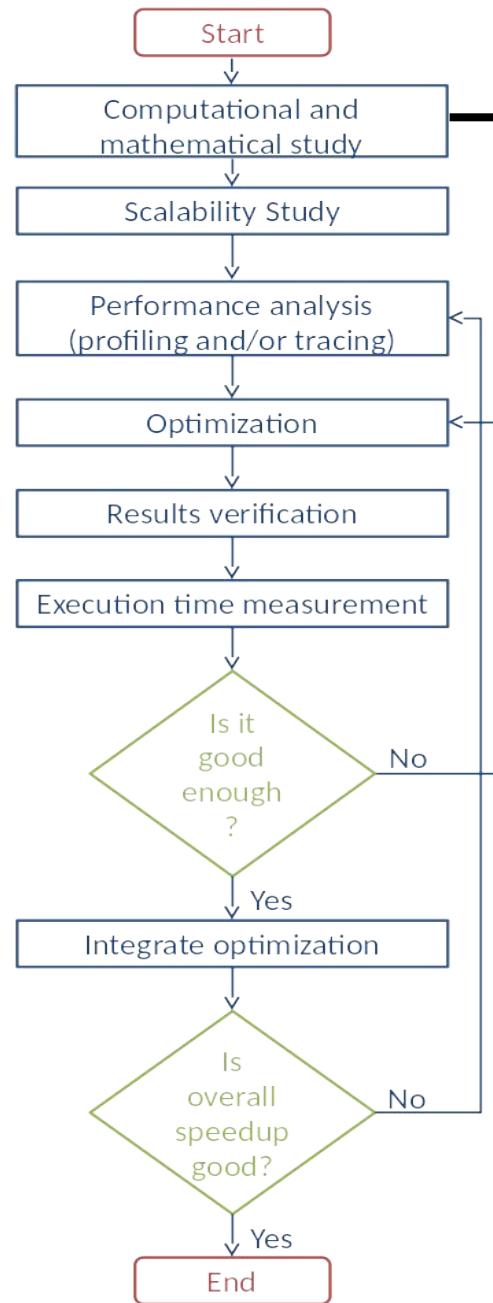


- Knowledge about the mathematical and computational side of Earth System Applications
- Knowledge about the specific needs in HPC of the Earth System Applications
- Researching about HPC methods specifically used for Earth System Applications

# Methodology

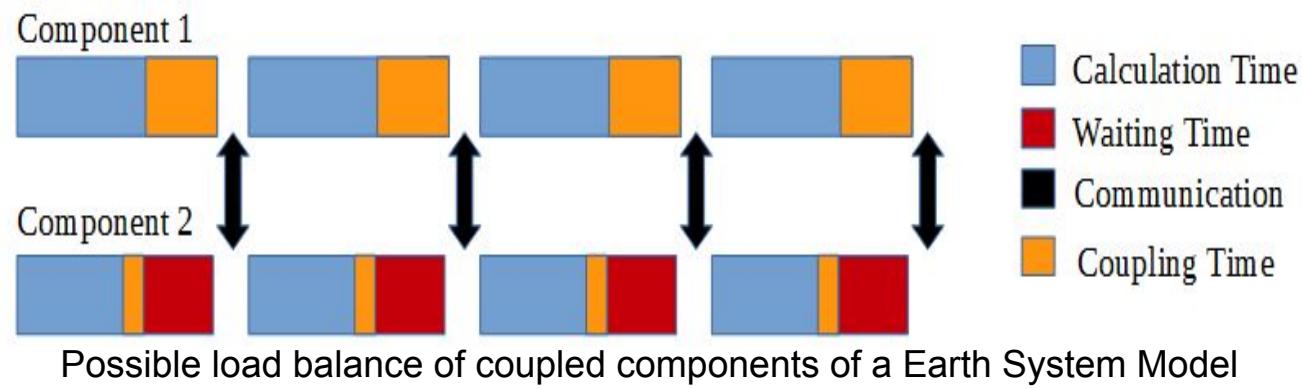
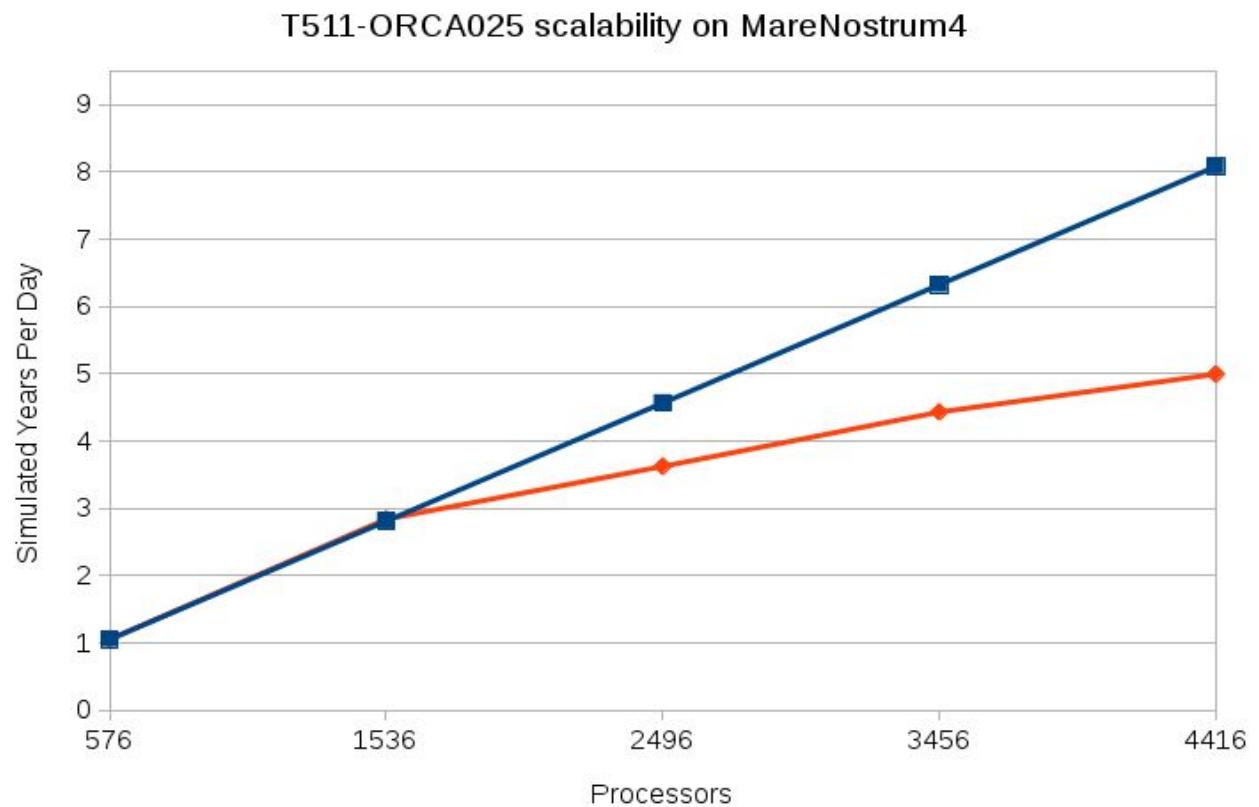
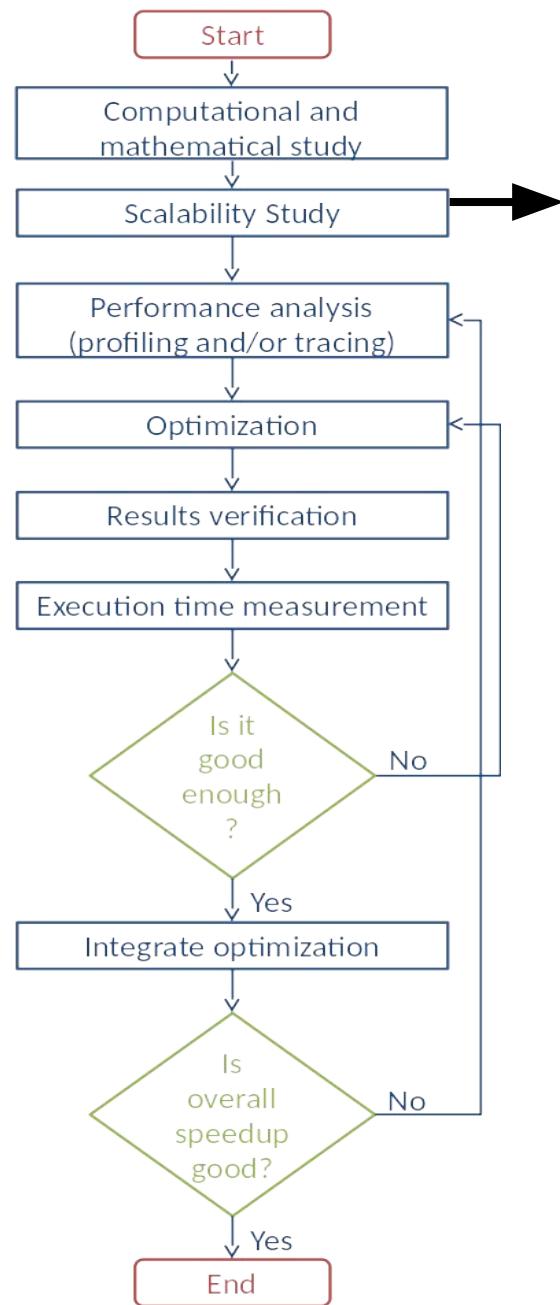


# Methodology

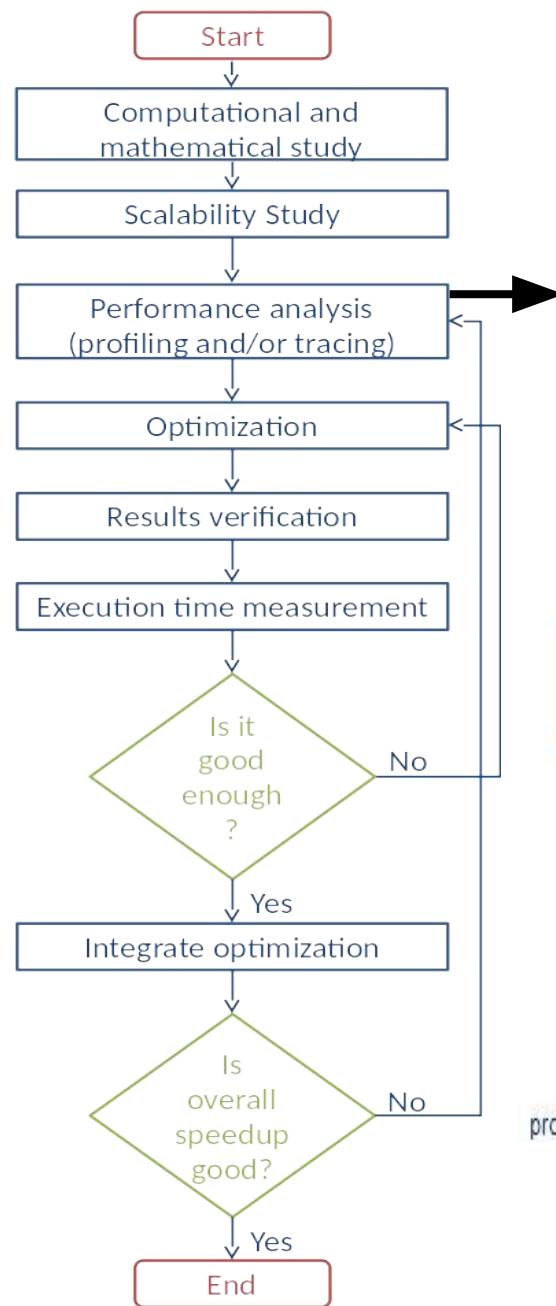


- Mathematical study
  - Some methods could be better than others
    - Discretization used (explicit, implicit, semi-implicit...)
    - Parallel adaptation (solvers, preconditioners...)
  - How to implement new algorithms for new architectures
- Computational study
  - Achieve load balance among components
  - Reduce overhead introduced by parallel applications
  - Assure that the computational algorithm takes advantage of the architecture

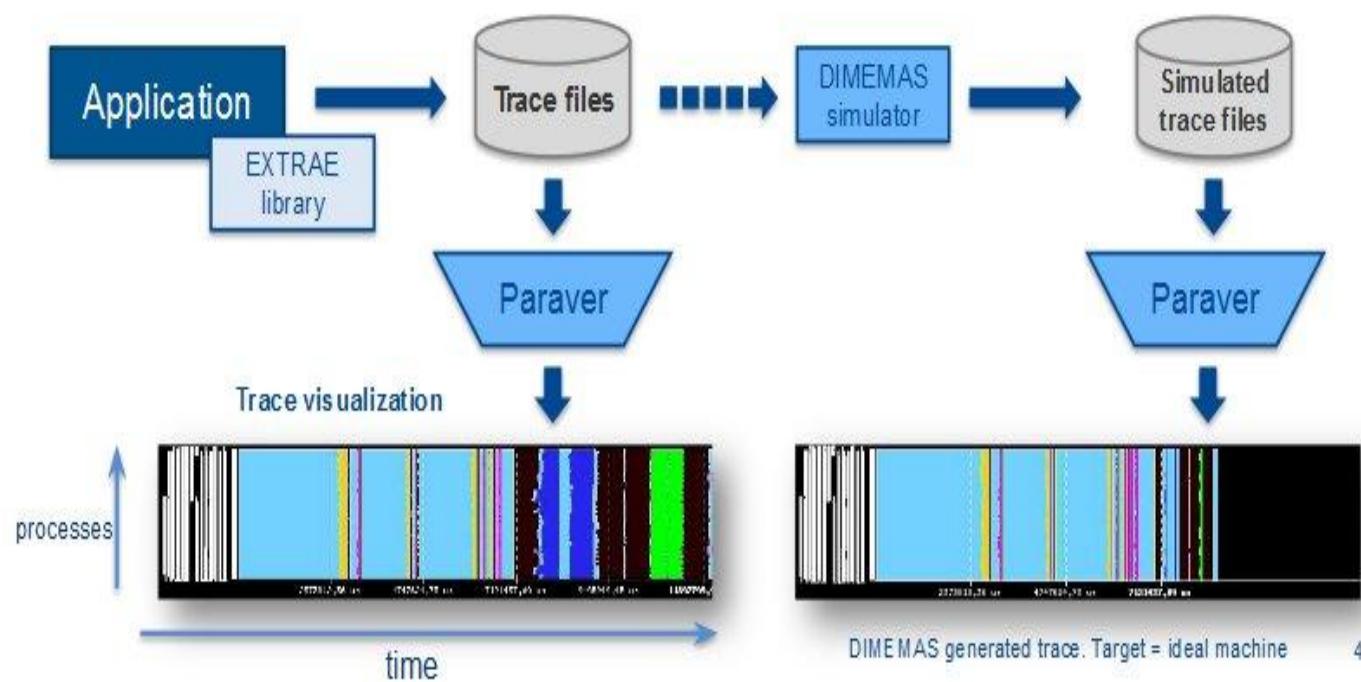
# Methodology



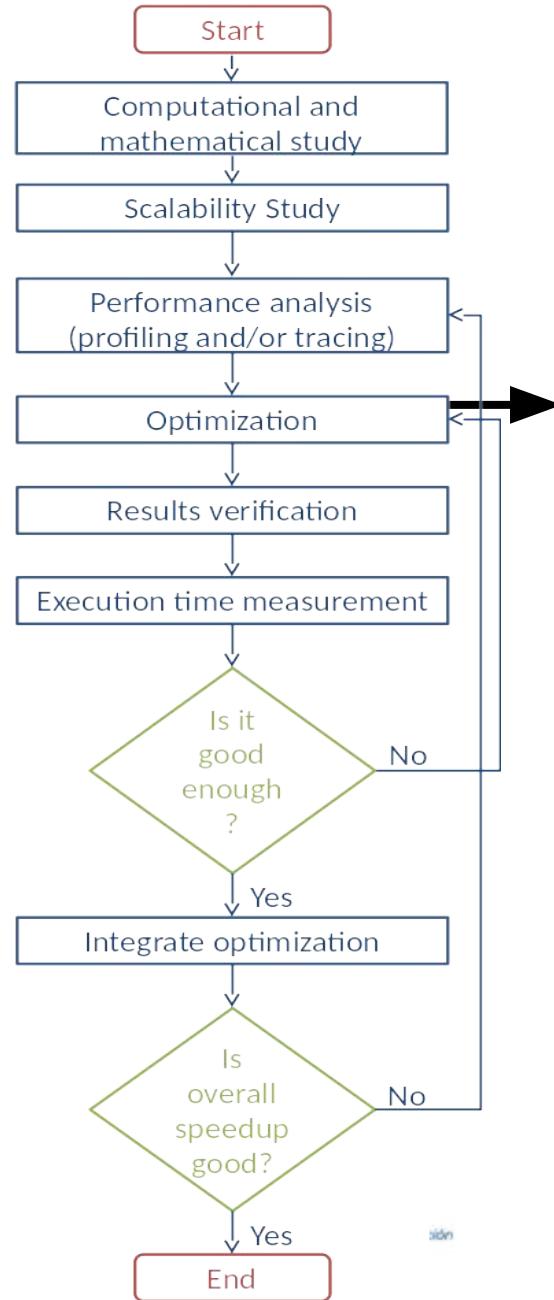
# Methodology



- Since 1991
- Based on traces
- Open Source: <http://www.bsc.es/paraver>
- **Exrae:** Package that generates Paraver trace-files for a post-mortem analysis
- **Paraver:** Trace visualization and analysis browser
  - Includes trace manipulation: Filter, cut traces
- **DIMEMAS:** Message passing simulator



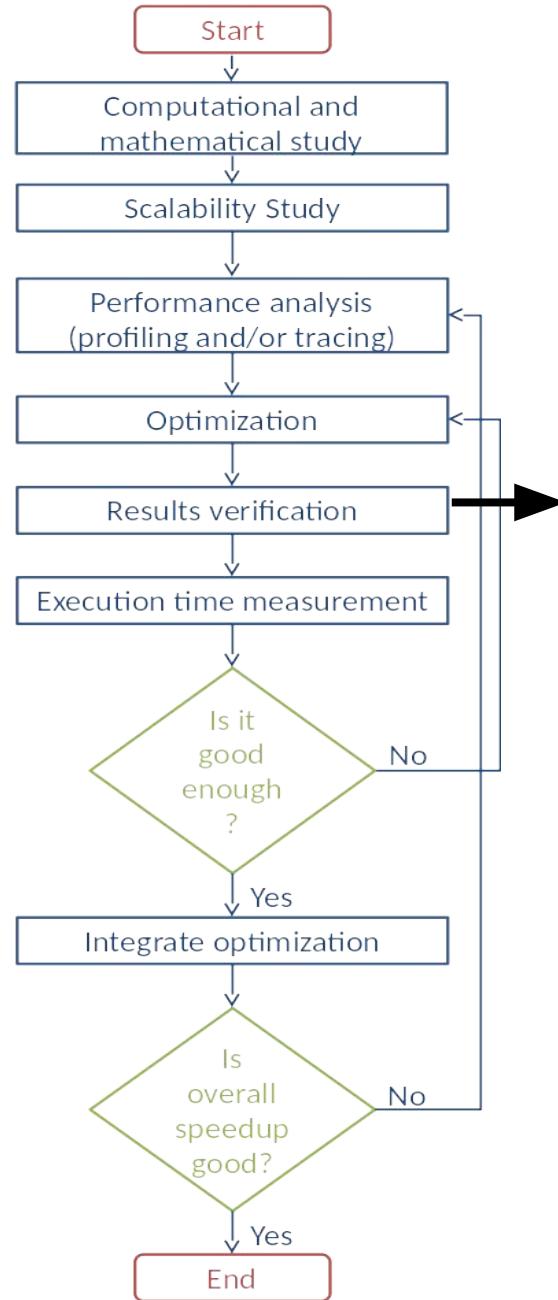
# Methodology



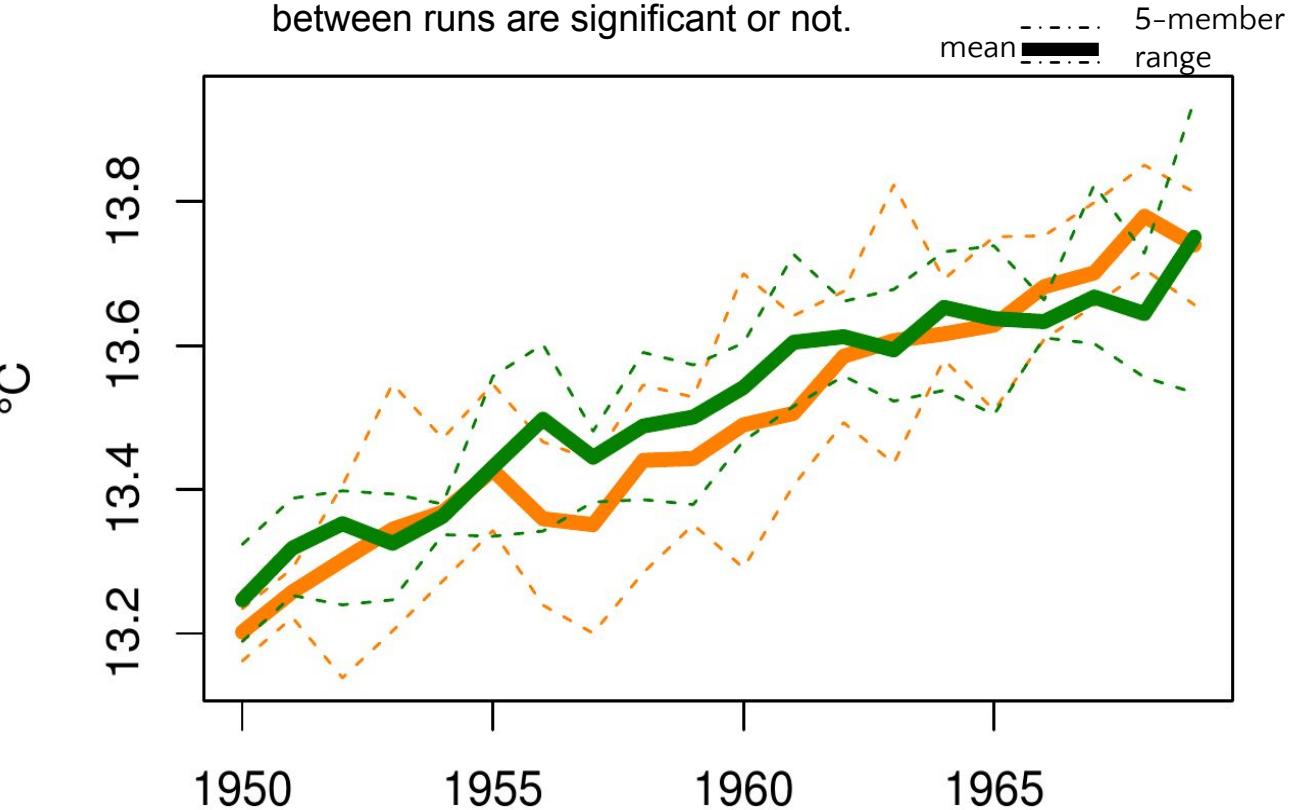
- Introducing optimizations
  - Improvement of the mathematical and/or computational algorithm
    - Apply scientific methods which are found in the literature
    - Improve the method with a new approach
  - Revolution: Create a new (and better) algorithm taking into account the research line followed



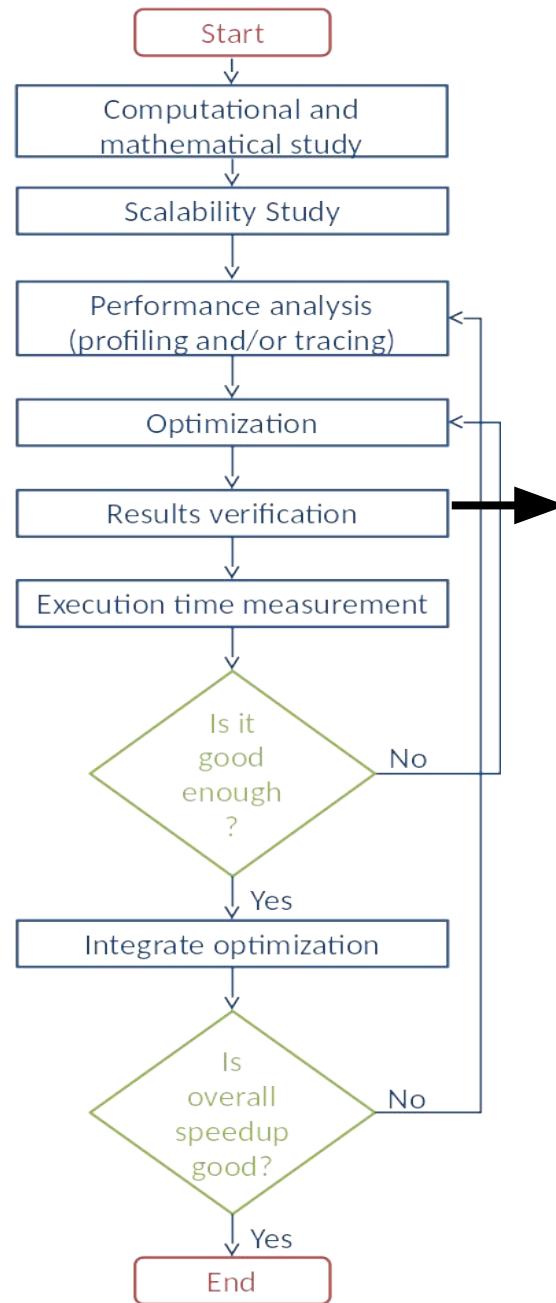
# Methodology



- Reproducibility study
  - Evaluate if the accuracy and reproducibility of the model is similar using or not the optimizations proposed
  - Take into account the nature of climate models
    - How to evaluate, in parallel executions, if the differences between runs are significant or not.

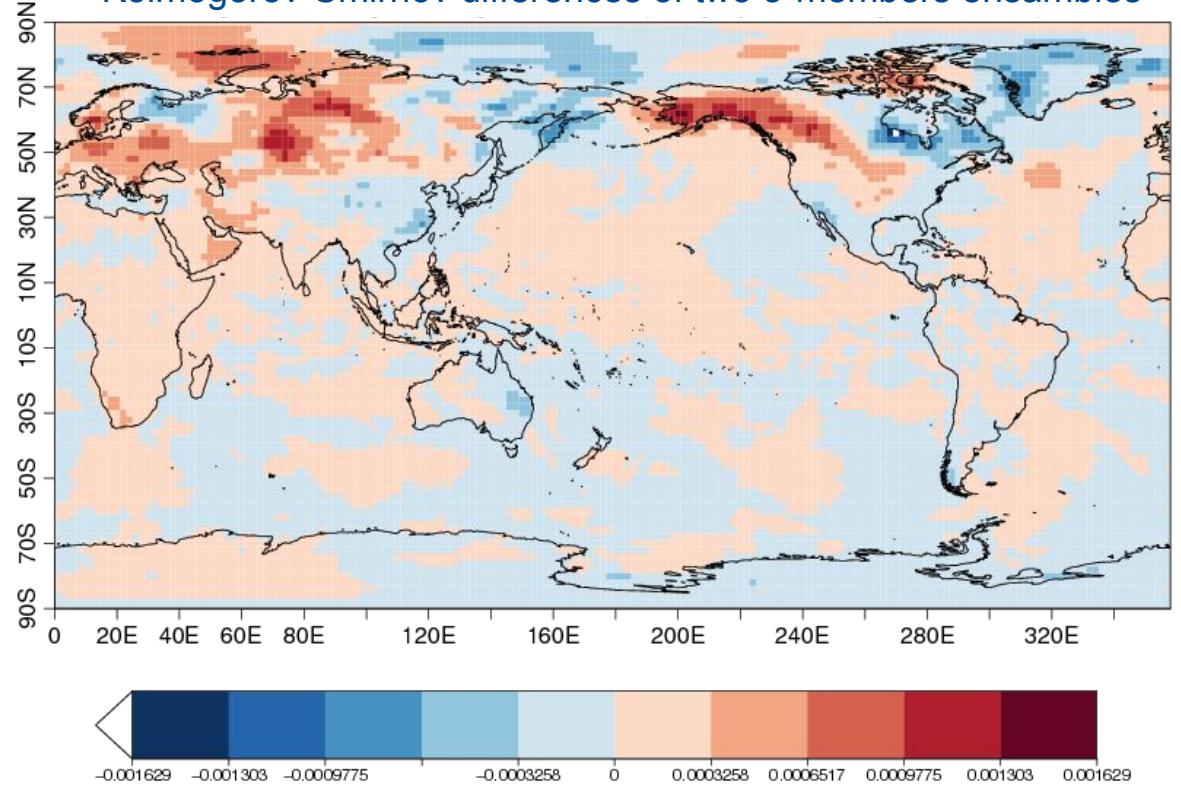


# Methodology



- Reproducibility study
  - Evaluate if the accuracy and reproducibility of the model is similar using or not the optimizations proposed
  - Take into account the nature of climate models
    - How to evaluate, in parallel executions, if the differences between runs are significant or not.

Kolmogorov-Smirnov differences of two 5-members ensambles

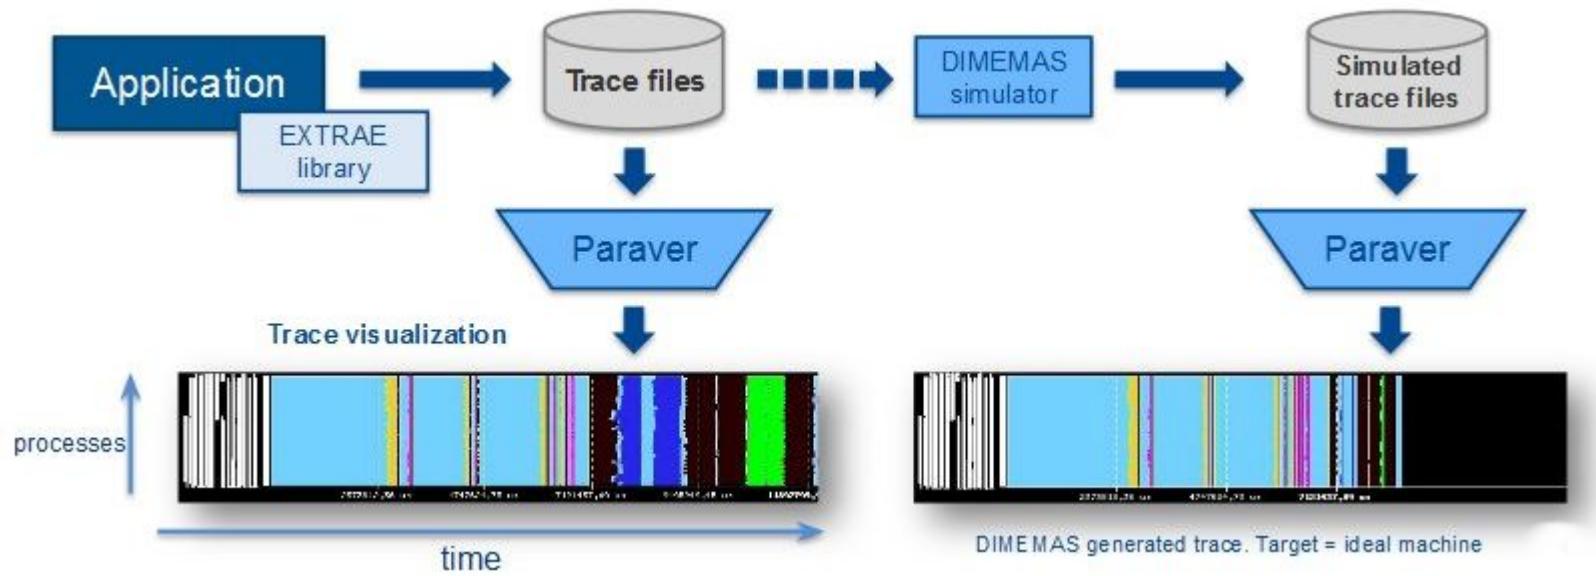


# Profiling Analysis: BSC Tools

- BSC Tools
  - General description
  - Extrae
    - General description
    - How to use it
  - Paraver
    - General description
    - How to use it
    - Configurations available
  - Dimemas
    - General description
  - How to work with large traces
    - Filtering/Burst mode
    - Cutting

# BSC Tools

- Since 1991
- Based on traces
- Open Source → <http://www.bsc.es/paraver>
- Extrae: Package that generates Paraver-trace files for a post-mortem analysis
- Paraver: Trace visualization and analysis browser
- Dimemas: Message passing simulator
  - Include traces manipulation: Filter, cut traces...



# BSC Tools

## CORE TOOLS

### EXTRAE

Instrumentation framework to generate execution traces of the most used parallel runtimes.

[Get EXTRAE](#)

Version 3.4.3 • 2.32 MB



### PARAVER

Expressive powerful and flexible trace visualizer for post-mortem trace analysis.

[Get PARAVER](#)

Version 4.6.3 • 1.56 MB



### DIMEMAS

High-abstracted network simulator for message-passing programs.

[Get DIMEMAS](#)

Version 5.3.0 • 0.93 MB



## PERFORMANCE ANALYTICS

### CLUSTERING

Automatically expose the main performance trends in applications' computation structure.

[Get CLUSTERING](#)

Version 2.6.6 • 1.97 MB



### TRACKING

Analyze how the behavior of a parallel application evolves through different scenarios.

[Get TRACKING](#)

Version 2.6.5 • 1.88 MB



### FOLDING

Combined instrumentation and sampling for instantaneous metric evolution with low overhead.

[Get FOLDING](#)

Version 1.3.1 • 12.67 MB



### SPECTRAL

Signal processing techniques to select representative regions from Paraver traces.

[Get SPECTRAL](#)

Version 3.4.0 • 0.3 MB



### BASIC ANALYSIS

Framework for automatic extraction of fundamental factors for Paraver traces.

[Get BASIC ANALYSIS](#)

Version 0.2 • 66.41 MB



<https://tools.bsc.es/downloads>

# BSC Tools

## Home » Documentation » Tutorial Guidelines

These six tutorials can be opened with wxParaver versions newer than 4.3.0, and you'll be able to follow the steps within the tool. To install them, download and untar the package and follow the instructions of the Help/Tutorial option on the Paraver main window. You can download them in a single package either in [.tar.gz format](#) (127 Mb) or [.zip format](#) (127 Mb).

- [Paraver introduction \(MPI\)](#) Start here to familiarize with Paraver basic commands and the first steps of a performance analysis.
- [Dimemas introduction](#) The basic steps to learn how to configure and run the Dimemas simulator and to start looking at the results.
- [Introduction to Paraver and Dimemas methodology](#) This tutorial presents different ways to analyze a MPI application through well-known rules, their diagnosis and how they impact on your exploration (no traces included).
- [Methodology](#) This tutorial shows some examples of the analysis that can be done using the provided configuration files.
- [Tutorial on HydroC analysis](#) (MPI, Dimemas, CUDA) One example of performance analysis of the MPI application Hydro and further simulations with Dimemas.
- [Trace preparation](#) Look at this tutorial to select a representative region for a large trace that cannot be loaded into memory.
- [Trace alignment tutorial](#). If you identify some unexpected unalignment or backwards communications, use this tutorial to learn how to correct shifts between processors.

## Methodology of analysis

[MPI+OpenMP Performance Analysis tips](#)

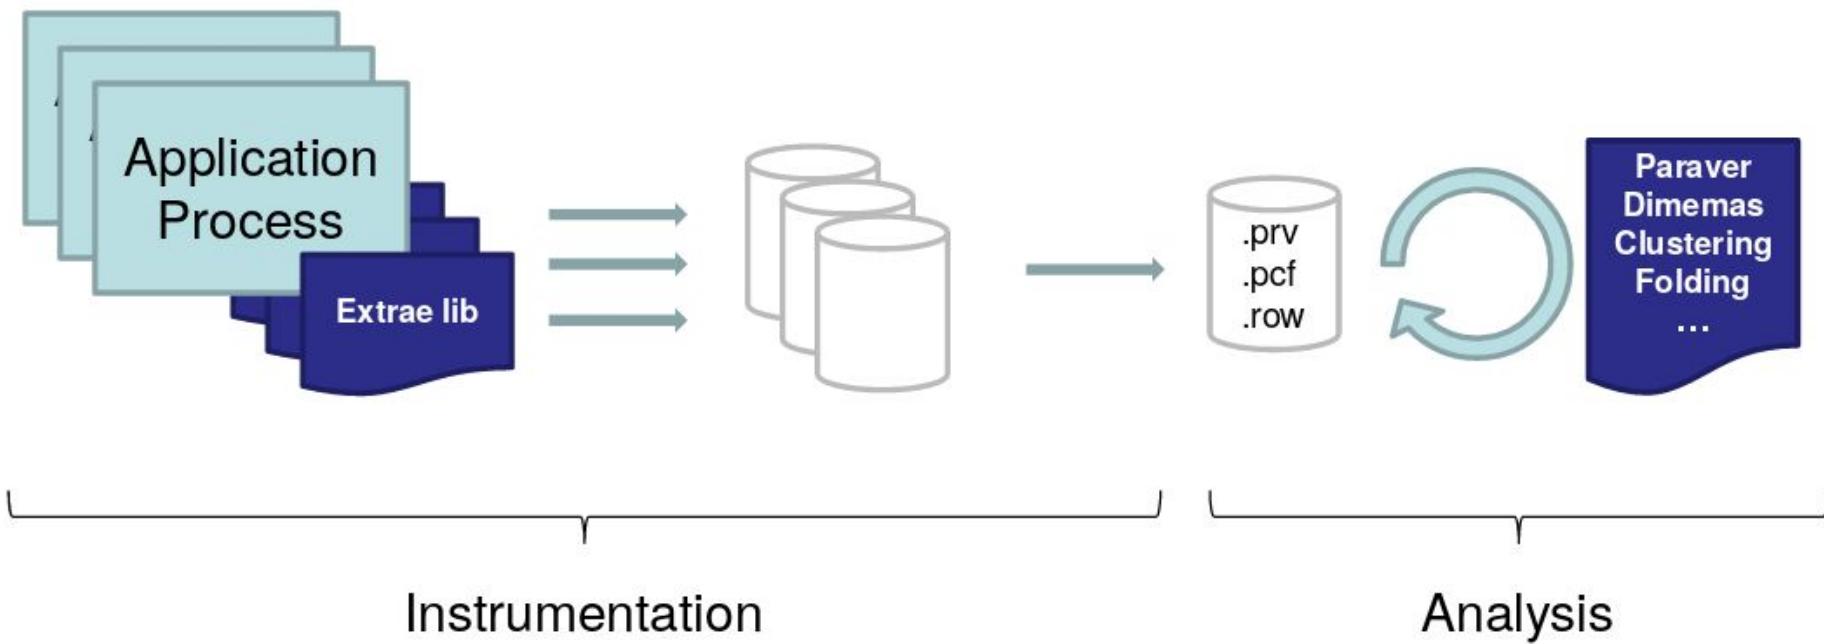
## Tutorial slides

[Introduction](#)

Core tools	Advanced features
Paraver, Detailed material	Tools scalability
Dimemas	Clustering
Extrae	Sampling

# BSC Tools:Extrace

- Trace generation



# BSC Tools:Extrae

- Trace Generation: Set Environment
  - Module load extrae
    - load extrae 3.X.0 (PATH, EXTRAE\_DIR, EXTRAE\_ROOT, EXTRAE\_LIB)
  - Job script → trace-fortran.sh | trace-c.sh
  - Extrae config → extraeMPI.xml | extraeMPI+OMP.xml
  - Files modified for model→ run\_parallel.sh

# BSC Tools:Extrae

- Job script:trace-fortran.sh
    - Available loading extrae module

```
#!/bin/bash

#Workaround for tracing in MN3, make TMPDIR point to an existing dir
#if [ ! -z "${TMPDIR}" ]; then
#    export TMPDIR=$TMPDIR/extrae
#    mkdir -p $TMPDIR
#fi

EXRAE_ROOT=/usr/local/apps/extrae

if [ -z "$EXRAE_DIR" ]; then
    echo "ERROR: EXRAE_DIR not set, maybe extrae module not loaded?"
    exit -1
fi

if [ -z "$OMP_TRACE" ]; then
    export LD_PRELOAD=${EXRAE_DIR}/lib/libmpitracef.so
    if [ -z "$EXRAE_CONFIG_FILE" ]; then
        export EXRAE_CONFIG_FILE=${EXRAE_ROOT}/xml/MPI/extrae.xml
    fi
else
    export LD_PRELOAD=${EXRAE_DIR}/lib/libompitracef.so # For Fortran apps
    if [ -z "$EXRAE_CONFIG_FILE" ]; then
        export EXRAE_CONFIG_FILE=${EXRAE_ROOT}/xml/MPI+OMP/extrae.xml
    fi
fi
```

True if openmp is used

Exrae config by default

# BSC Tools:Extrae

- Extrae config:extrae.xml
  - Available using nama\_CY43R1\_IFS\_traces branch

```
<?xml version='1.0'?>

<trace enabled="yes"
  home="/usr/local/apps/extrae/3.3.0rc/MPICH2-6.3.1"
  initial-mode="detail"
  type="paraver"
  xml-parser-id="Id: xml-parse.c 3893 2016-03-04 12:30:23Z harald $"
>
  <mpi enabled="yes">
    <counters enabled="yes" />
  </mpi>
  <openmp enabled="no">
    <locks enabled="no" />
    <counters enabled="yes" />
  </openmp>
  <pthread enabled="no">
    <locks enabled="no" />
    <counters enabled="yes" />
  </pthread>
  <callers enabled="yes">
    <mpi enabled="yes">1-3</mpi>
    <sampling enabled="yes">1-5</sampling>
    <dynamic-memory enabled="no">1-3</dynamic-memory>
  </callers>
  <user-functions enabled="yes" list="/perm/rd/nama/extrae/xml/MPI/functions_for_xml.txt" exclude-automatic-functions="no">
    <counters enabled="yes" />
  </user-functions>
  <counters enabled="yes">
    <cpu enabled="yes" starting-set-distribution="1">
      <set enabled="yes" domain="all" changeat-time="0">
        PAPI_TOT_INS,PAPI_TOT_CYC,PAPI_L1_DCM,PAPI_L2_DCM,PAPI_L3_TCM,PAPI_FP_INS,PAPI_BR_MSP,PAPI_FP_OPS
      </set>
      <set enabled="yes" domain="all" changeat-time="0">
        PAPI_TOT_INS,PAPI_TOT_CYC,PAPI_LD_INS,PAPI_SR_INS,PAPI_BR_UCN,PAPI_BR_CN,PAPI_VEC_SP,RESOURCE_STALLS
        <sampling enabled="no" period="1000000000">PAPI_TOT_CYC</sampling>
      </set>
    </cpu>
    <network enabled="no" />
    <resource-usage enabled="no" />
    <memory-usage enabled="no" />
  </counters>
```

Activate MPI tracing and emit hardware counters at MPI calls

Activate OpenMP tracing

Emit call stack information (number of levels) at acquisition point

Add instrumentation at specific user functions

PAPI counters used

# BSC Tools:Extrae

- Extrae config:extrae.xml
  - Available using nama\_CY43R1\_IFS\_traces branch

```
<storage enabled="no">
  <trace-prefix enabled="yes">TRACE</trace-prefix>
  <size enabled="no">5</size>
  <temporal-directory enabled="yes">/scratch</temporal-directory>
  <final-directory enabled="yes">/gpfs/scratch/bsc41/bsc41273</final-directory>
</storage>

<buffer enabled="yes">
  <size enabled="yes">500000</size>
  <circular enabled="no" />
</buffer>

<trace-control enabled="no">
  <file enabled="no" frequency="5M">/gpfs/scratch/bsc41/bsc41273/control</file>
  <global-ops enabled="no"></global-ops>
  <remote-control enabled="no">
    <signal enabled="no" which="USR1"/>
  </remote-control>
</trace-control>

<others enabled="yes">
  <minimum-time enabled="no">10M</minimum-time>
  <finalize-on-signal enabled="yes">
    SIGUSR1="no" SIGUSR2="no" SIGINT="yes"
    SIGQUIT="yes" SIGTERM="yes" SIGXCPU="yes"
    SIGFPE="yes" SIGSEGV="yes" SIGABRT="yes"
  />
  <flush-sampling-buffer-at-instrumentation-point enabled="yes" />
</others>

<bursts enabled="no">
  <threshold enabled="yes">500u</threshold>
  <mpi-statistics enabled="yes" />
</bursts>

<sampling enabled="no" type="default" period="50m" variability="10m" />

<dynamic-memory enabled="no">
  <alloc enabled="yes" threshold="32768" />
  <free enabled="yes" />
</dynamic-memory>

<input-outfit enabled="no" />

<merge enabled="yes">
  synchronization="default"
  tree-fan-out="16"
  max-memory="32768"
  joint-states="yes"
  keep-mpits="yes"
  sort-addresses="yes"
  overwrite="yes"
/>
```

Emit computation burst of a minimal duration

Plus summarized MPI events

Merge individual traces automatically

# BSC Tools:Extrae

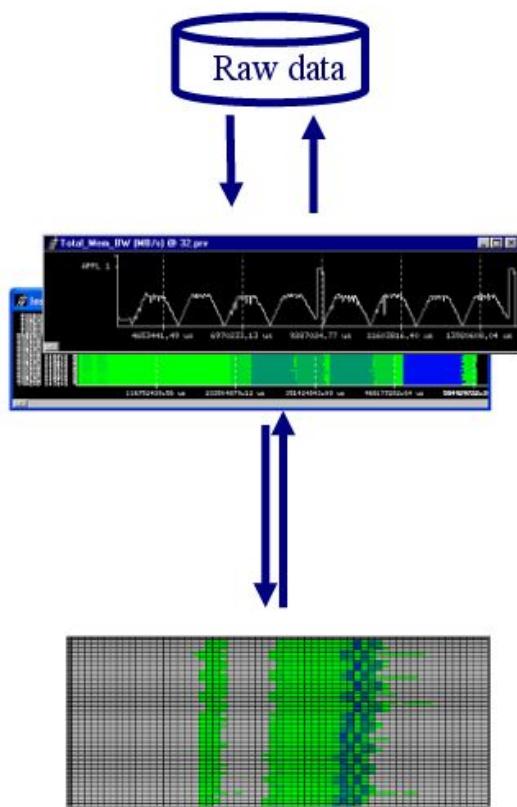
- Files modified for run\_parallel
  - If BSCTRACE=1 → Extrae is used

```
if [[ %BSCTRACE:0% = 1 ]]; then
    command="aprun -cc cpu $marg -n $submit_total_tasks -N $submit_tasks_per_node $Sarg -j $submit_cpus_per_compute_unit -d $omp_num_threads $submit_force numa_memory_affinity trace-fortran.sh $(whence $cmd) ${cmd##* }"
d) $args"
else
    command="aprun -cc cpu $marg -n $submit_total_tasks -N $submit_tasks_per_node $Sarg -j $submit_cpus_per_compute_unit -d $omp_num_threads $submit_force numa_memory_affinity $(whence $cmd) ${args}"
fi
fi
```

```
Parameter to activate profiling
export EC_LD_LIBRARY_PATH=${PIFS_LD_LIBRARY_PATH:-$LIBS}
if [[ %BSCTRACE:0% = 1 ]]; then
    LOAD_MODULE extrae
    export EXTRAE_CONFIG_FILE=/perm/rd/nama/extrae/xml/MPI/extrae.xml
    export TRACEDIR=${WDIR}/bsctrace.${TASK}.${ECF_TRYNO}.$$
    export EC_LD_LIBRARY_PATH=${EC_LD_LIBRARY_PATH}:$EXTRAELIB
    if [[ %OMPTRACE:0% = 1 ]]; then
        export OMP_TRACE=1
    fi
    mkdir $TRACEDIR
fi
(eval $command)
if [[ %BSCTRACE:0% = 1 ]]; then
    mv *.prv *.pcf *.row $TRACEDIR
fi
```

Trace files generated

# BSC Tools:Paraver



Trace visualization/analysis  
+ trace manipulation

## Timelines

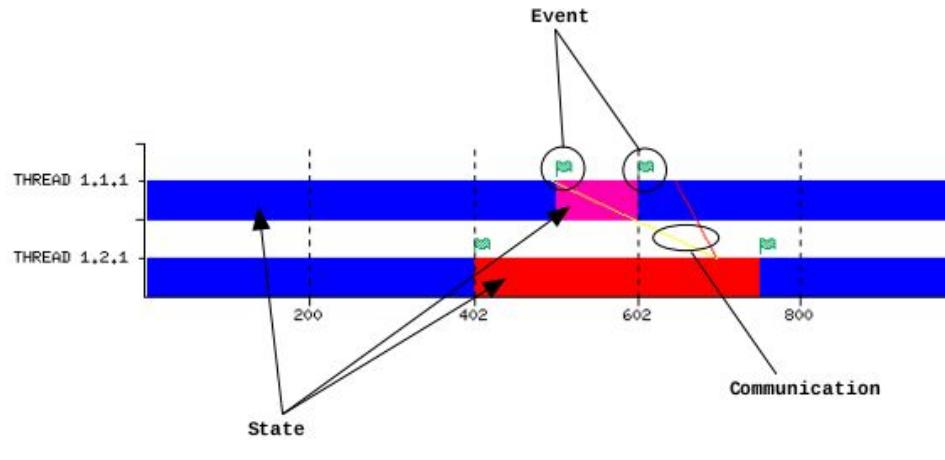
Goal = Flexibility  
No semantics  
Programmable

## 2/3D tables (Statistics)

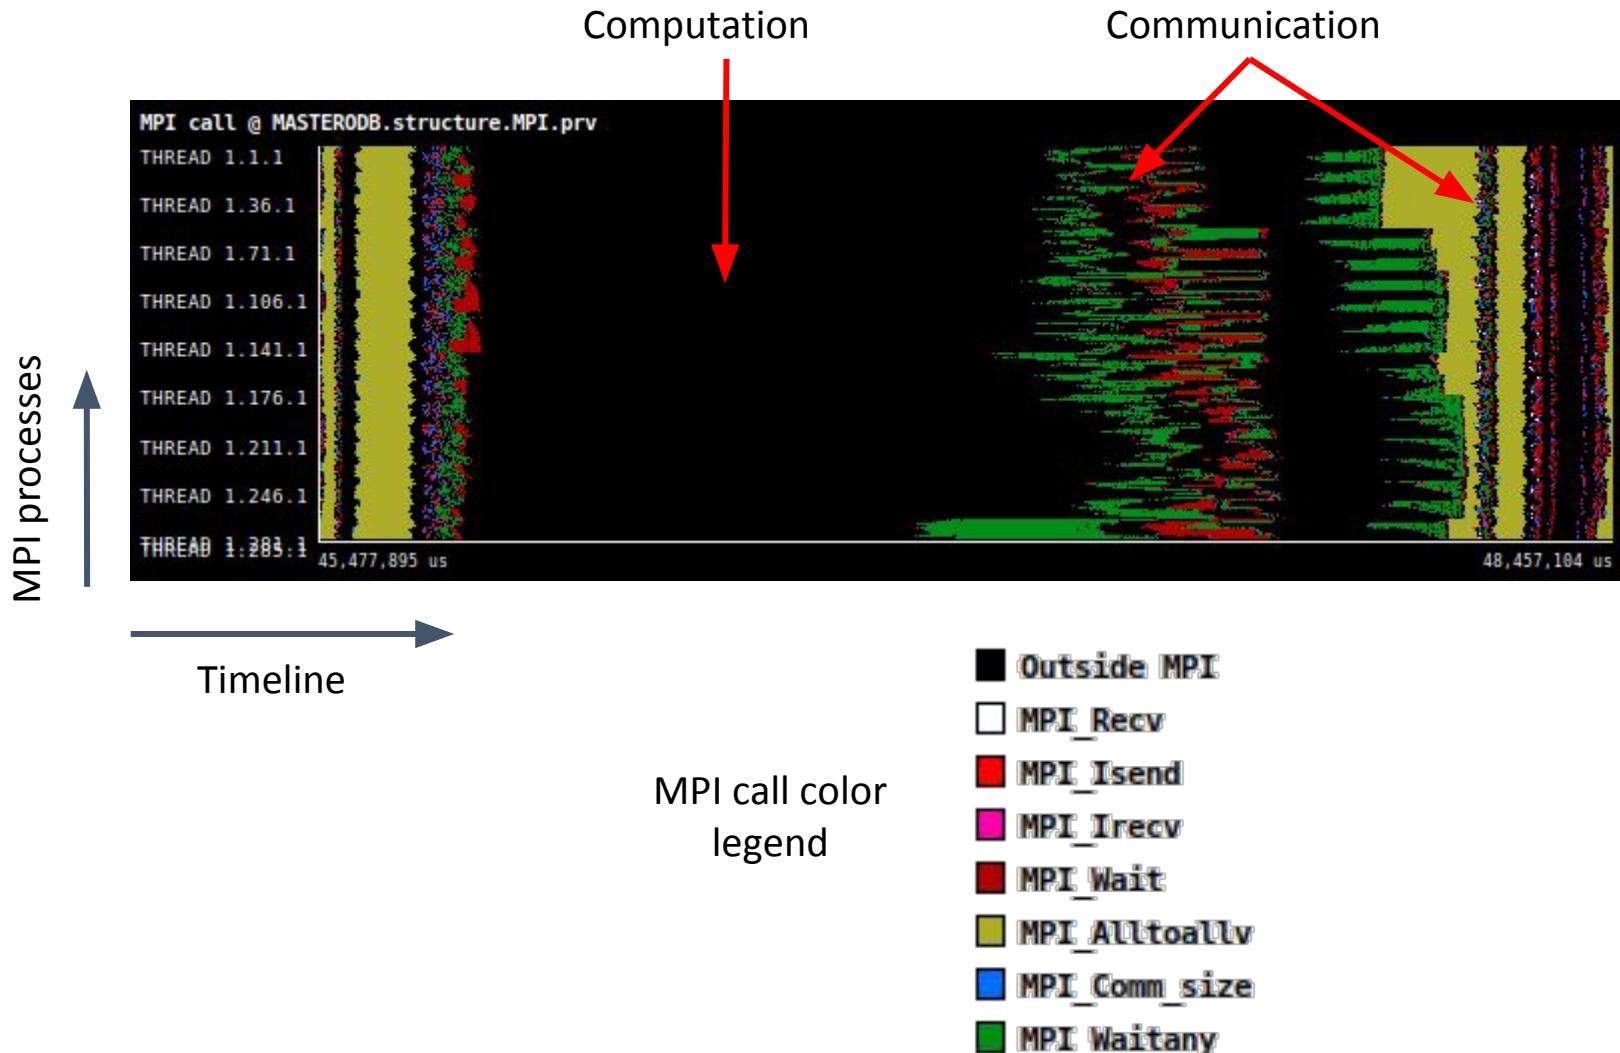
Comparative analyses  
Multiple traces  
Synchronize scales

# BSC Tools:Paraver

- **Paraver traces:** made up from records (timestamp + event or activity) of three different kind:
  - **State records:** intervals of thread status, i.e, waiting in a barrier (either MPI or OpenMP), waiting for a message, computing...
  - **Event records:** punctual event occurred in a given timestamp, as entry & exit points of user functions, MPI routines, OpenMP parallel regions...
  - **Communication records:** relationship between two objects, as communication between two processes (MPI), task movement among threads (OpenMP/OmpSs) or memory transfers (CUDA/OpenCL).

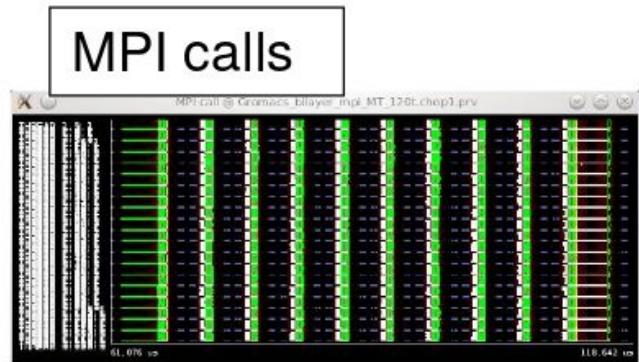


# How a trace looks like: basic overview



# BSC Tools:Paraver

- From timelines to tables

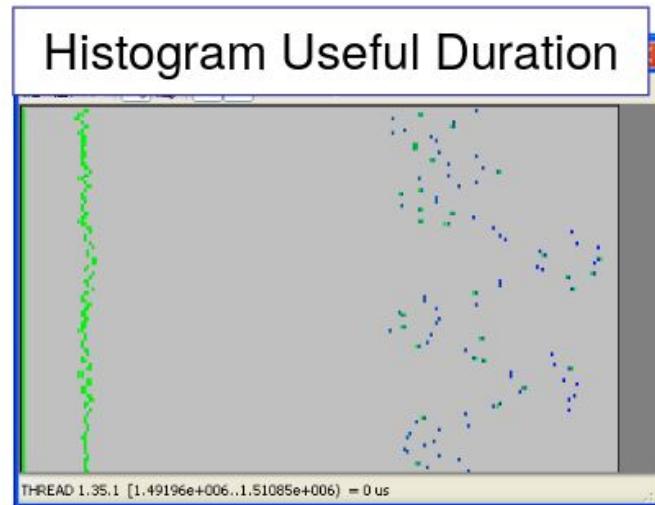


MPI calls profile

This figure displays the MPI calls profile for the Gromacs layer\_mpi\_MT\_120t.chop1.prv simulation. It includes a table of MPI operation percentages for individual threads and a histogram of MPI\_Sendrecv durations.

	Outside MPI	MPI_Send	MPI_Recv	MPI_Isend	MPI_Irecv	MPI_Waitall	MPI_Bcast	MPI_Reduce	MPI_Alloca
THREAD 1.113.1	87.6081 %	0.0682 %	9.9182 %	2.5777 %	1.7698 %	5.1676 %	0.5934 %	0.1465 %	-
THREAD 1.114.1	42.8434 %	-	20.5621 %	1.1847 %	1.0000 %	7.7056 %	-	-	-
THREAD 1.115.1	68.6127 %	0.0707 %	9.6223 %	2.2589 %	2.0177 %	5.9625 %	0.3249 %	0.0297 %	-
THREAD 1.116.1	74.6029 %	0.0531 %	9.6084 %	2.8913 %	2.5593 %	2.9296 %	0.5095 %	0.0482 %	-
THREAD 1.117.1	74.3733 %	0.0691 %	9.7012 %	2.8917 %	2.5240 %	-	-	-	-
THREAD 1.118.1	72.7770 %	0.0545 %	9.5489 %	2.8489 %	2.5353 %	-	-	-	-
THREAD 1.119.1	66.7994 %	0.0682 %	10.0674 %	2.4206 %	1.9741 %	-	-	-	-
THREAD 1.120.1	43.7224 %	-	20.5273 %	1.1912 %	1.0175 %	-	-	-	-
Total	8,012.4546 %	7.3174 %	1,370.5276 %	288.6168 %	253.0137 %	54	-	-	-
Average	66.7705 %	0.0690 %	11.4211 %	2.4051 %	2.1084 %	-	-	-	-
Maximum	75.6621 %	0.4390 %	21.2505 %	2.9706 %	2.6369 %	-	-	-	-
Minimum	40.5200 %	0.0129 %	8.8583 %	1.1489 %	1.0077 %	-	-	-	-
StdDev	11.3695 %	0.0474 %	4.0612 %	0.5984 %	0.5406 %	-	-	-	-
Avg/Max	0.8822	0.1572	0.5374	0.8096	0.7996	-	-	-	-

## Useful Duration

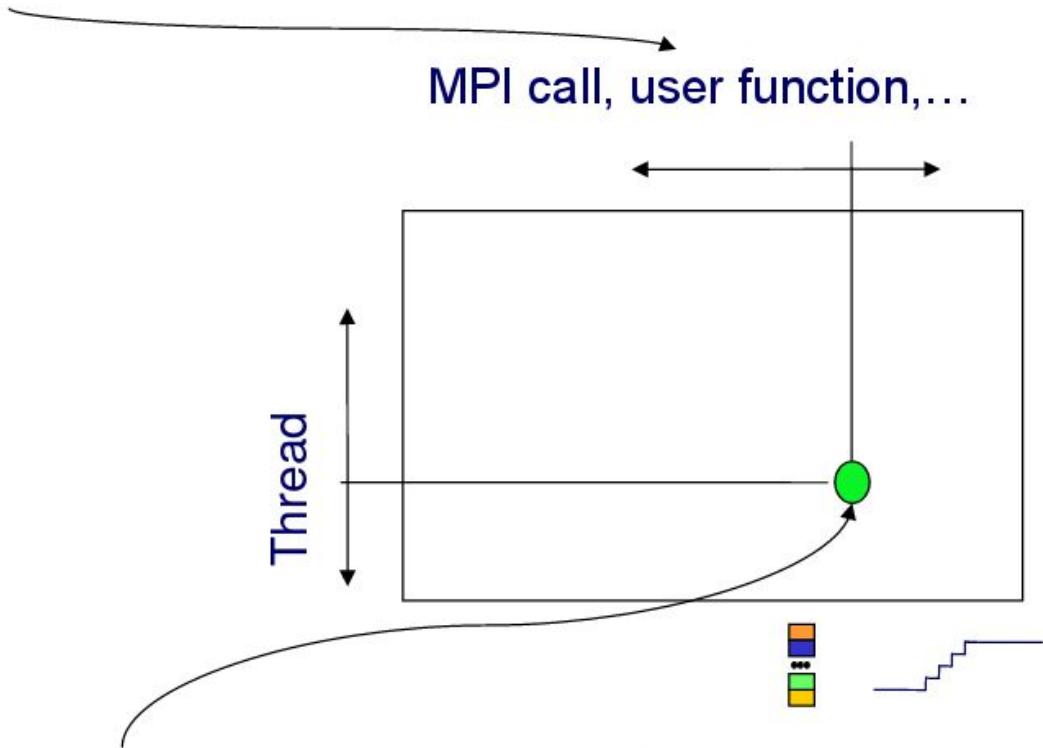


# BSC Tools:Paraver

One columns per specific value of categorical **Control window**

	End	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Allreduce	MPI_Comm
THREAD 1.1.1	86,98 %	0,06 %	0,08 %	11,12 %	1,75 %	
THREAD 1.2.1	88,29 %	0,10 %	0,10 %	9,95 %	1,56 %	
THREAD 1.3.1	88,33 %	0,13 %	0,10 %	9,92 %	1,51 %	
THREAD 1.4.1	89,75 %	0,10 %	0,09 %	8,62 %	1,44 %	
THREAD 1.5.1	89,47 %	0,11 %	0,10 %	8,85 %	1,46 %	
THREAD 1.6.1	88,76 %	0,12 %	0,09 %	9,54 %	1,48 %	
THREAD 1.7.1	91,77 %	0,13 %	0,10 %	6,51 %	1,49 %	
THREAD 1.8.1	90,23 %	0,06 %	0,08 %	8,13 %	1,50 %	
THREAD 1.9.1	91,88 %	0,13 %	0,09 %	6,73 %	1,17 %	
THREAD 1.10.1	93,24 %	0,18 %	0,11 %	5,41 %	1,05 %	
THREAD 1.11.1	93,25 %	0,18 %	0,11 %	5,45 %	1,00 %	
THREAD 1.12.1	94,63 %	0,17 %	0,11 %	4,16 %	0,93 %	
THREAD 1.13.1	93,40 %	0,17 %	0,11 %	5,35 %	0,96 %	
THREAD 1.14.1	94,99 %	0,20 %	0,11 %	3,77 %	0,93 %	
THREAD 1.15.1	96,80 %	0,22 %	0,11 %	1,92 %	0,95 %	
THREAD 1.16.1	95,73 %	0,12 %	0,09 %	2,99 %	1,06 %	

MPI call, user function, ...



Value/color is a statistic computed for the specific thread  
when control window had the value corresponding to the column

**Relevant statistics:**

Time, %time, #bursts, Avg. burst time  
Average of **Data window**

# MPI calls and profile

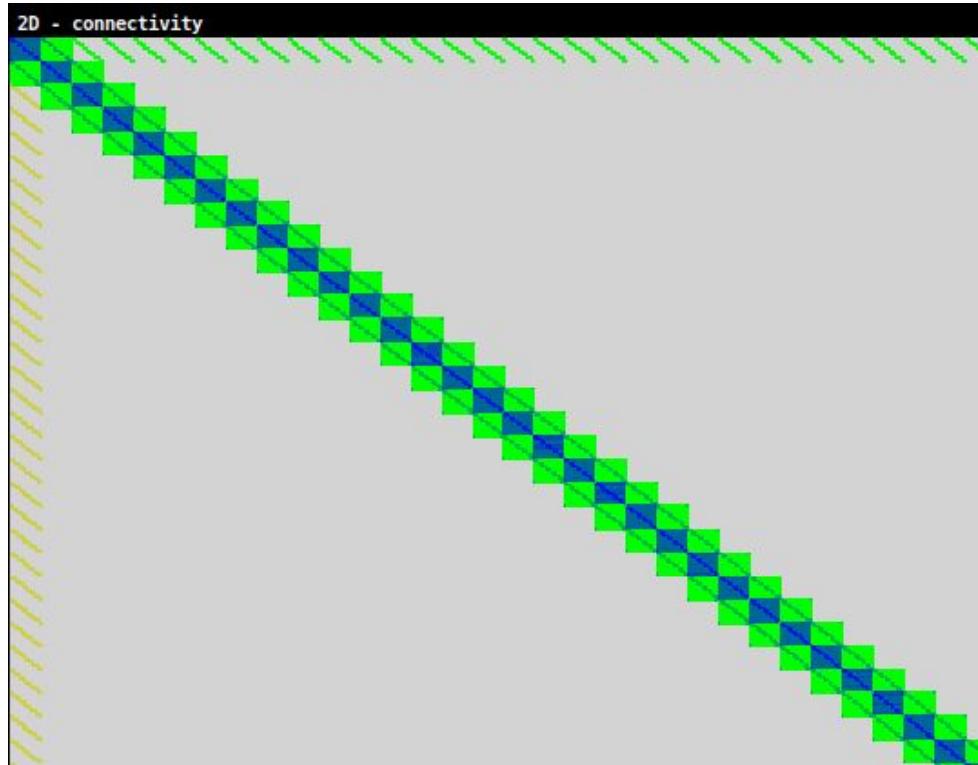
- Different types of MPI functions are quantified
- In this case, only the MPI\_Alltoallv and MPI\_Waitany functions represent a significant amount of time with 14.65% and 9.29% respectively.



	Outside MPI	MPI_Recv	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Alltoallv	MPI_Comm_size	MPI_Waitany
<b>Total</b>	41,154.02 %	37.32 %	260.82 %	172.16 %	2,023.25 %	8,438.25 %	163.09 %	5,351.09 %
<b>Average</b>	71.45 %	0.07 %	0.45 %	0.30 %	3.51 %	14.65 %	0.28 %	9.29 %
<b>Maximum</b>	84.51 %	0.18 %	1.15 %	0.81 %	9.70 %	20.07 %	0.31 %	31.12 %
<b>Minimum</b>	51.79 %	0.00 %	0.13 %	0.12 %	0.49 %	8.59 %	0.23 %	1.31 %
<b>StDev</b>	4.91 %	0.04 %	0.18 %	0.14 %	1.79 %	3.07 %	0.01 %	4.33 %
<b>Avg/Max</b>	0.85	0.38	0.39	0.37	0.36	0.73	0.93	0.30

# Point-to-point connectivity matrix

- It indicates who communicates with whom
- Almost all point-to-point communications are locally performed between MPI processes neighbours



# Collective communications

- Four calls to MPI\_Alltoallv each time step
- The most significant in terms of size and duration is the second one

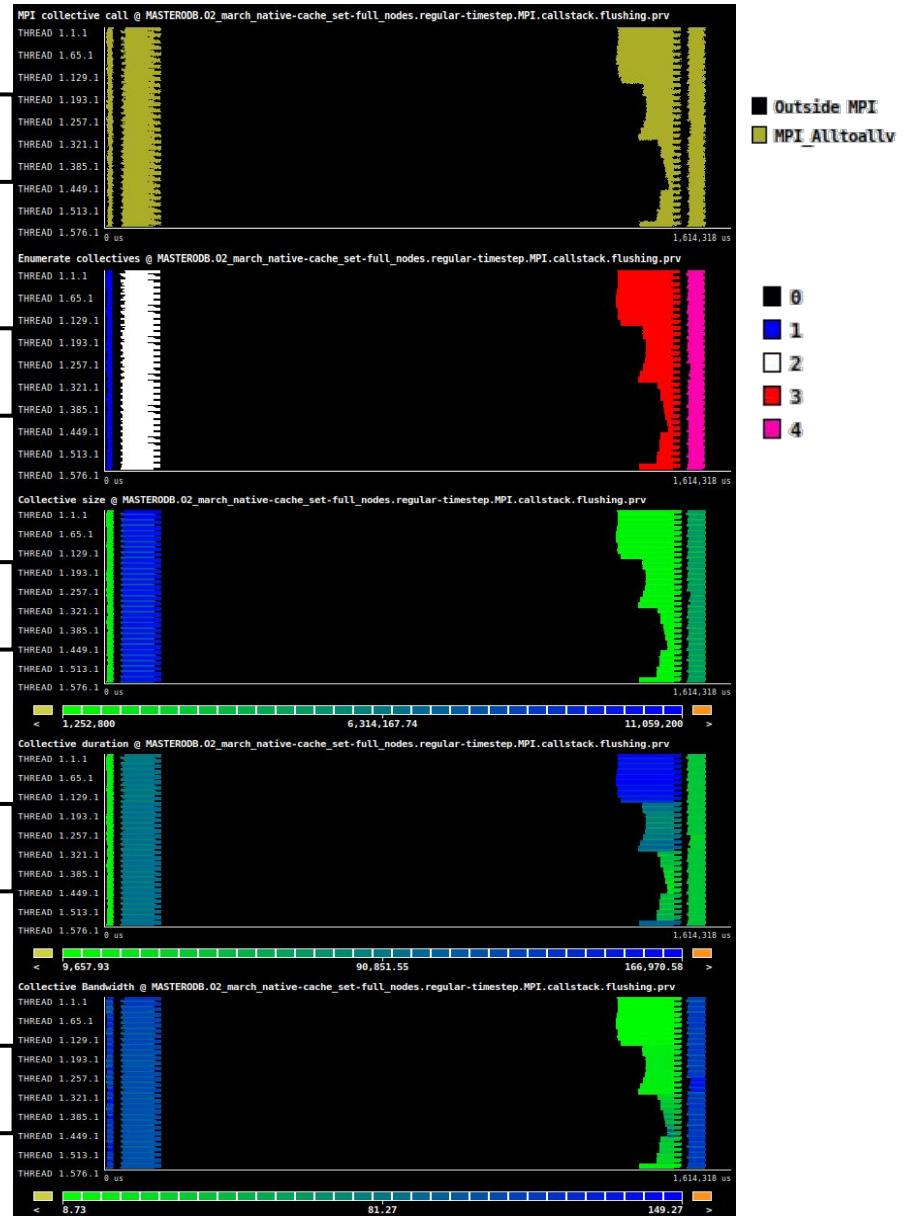
Call

Enumeration

Size

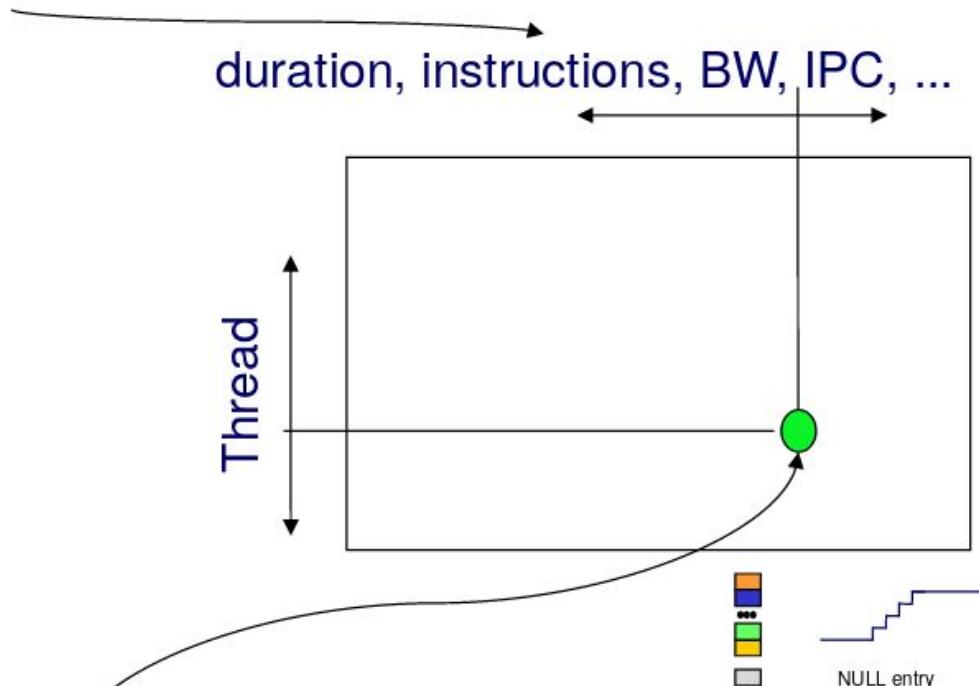
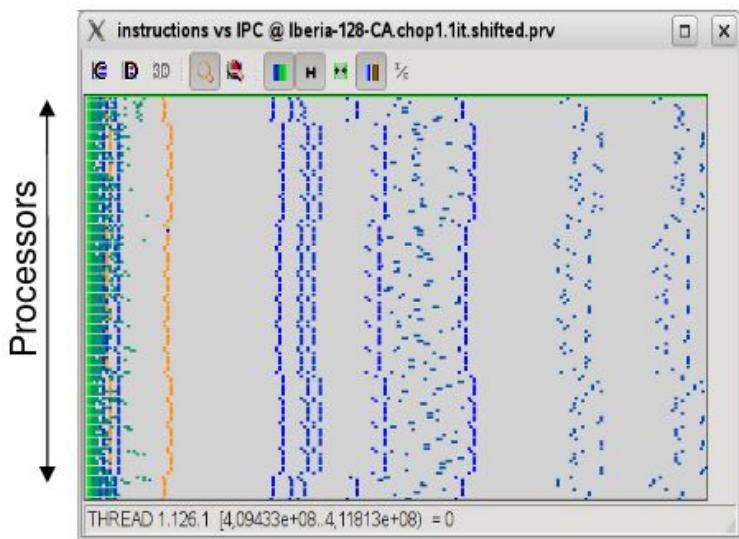
Duration

Bandwidth



# BSC Tools:Paraver

Columns correspond to bins of values of a numeric **Control window**

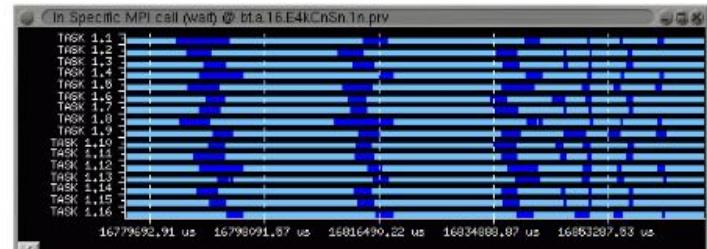
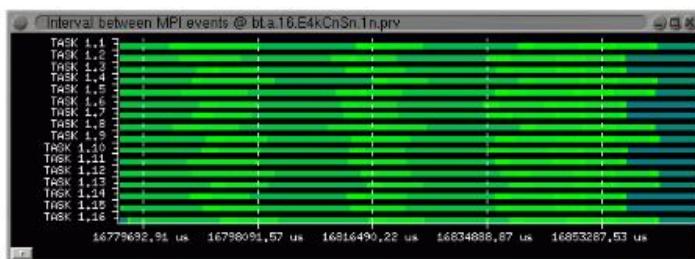


Value/color is a statistic computed for the specific thread  
when control window had the value corresponding to the column

**Relevant statistics:**  
Time, %time, #bursts, Avg. burst time  
Average of **Data window**

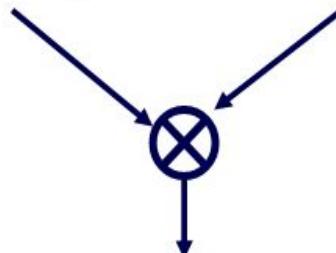
# BSC Tools:Paraver

- Semantic functionality
  - Derived windows
    - Point wise operation
      - $S = \alpha * S^a <\text{op}> \beta * S^b$
      - $<\text{op}> : +, -, *, /, \dots$



Interval between MPI events

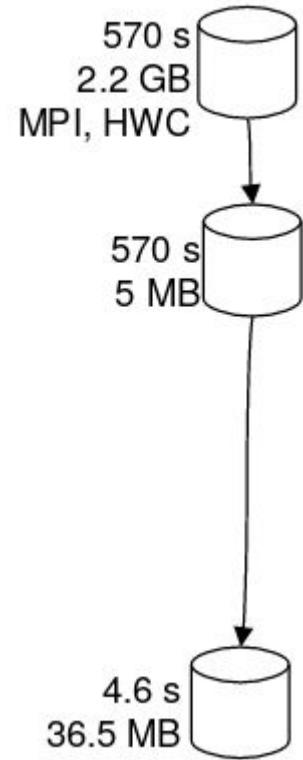
In MPI call



MPI call duration

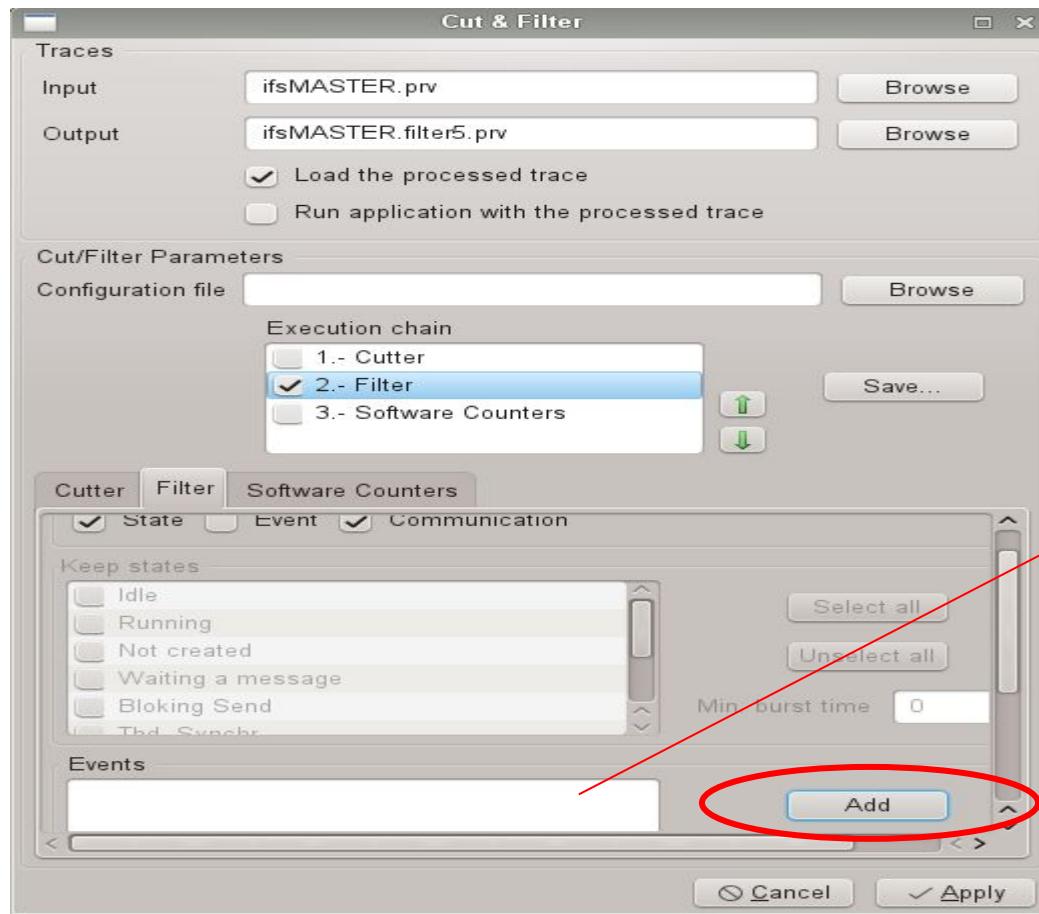
# BSC Tools:Paraver

- Data handling capability
  - Original trace containing all the events
  - Filtering/Burst mode
    - Subset of records in original trace
    - By duration, time, value, event type
    - Trace filtered can be analysed in the same way
    - Also using burst mode from xml file
      - Save only computation bursts longer than a value
  - Cutting
    - All records in a given time interval
    - Only some processes



# BSC Tools:Paraver

- Filtering
  - Filter original trace discarding most of the records, only keeping most relevant information (MPI events can be used for this purpose)

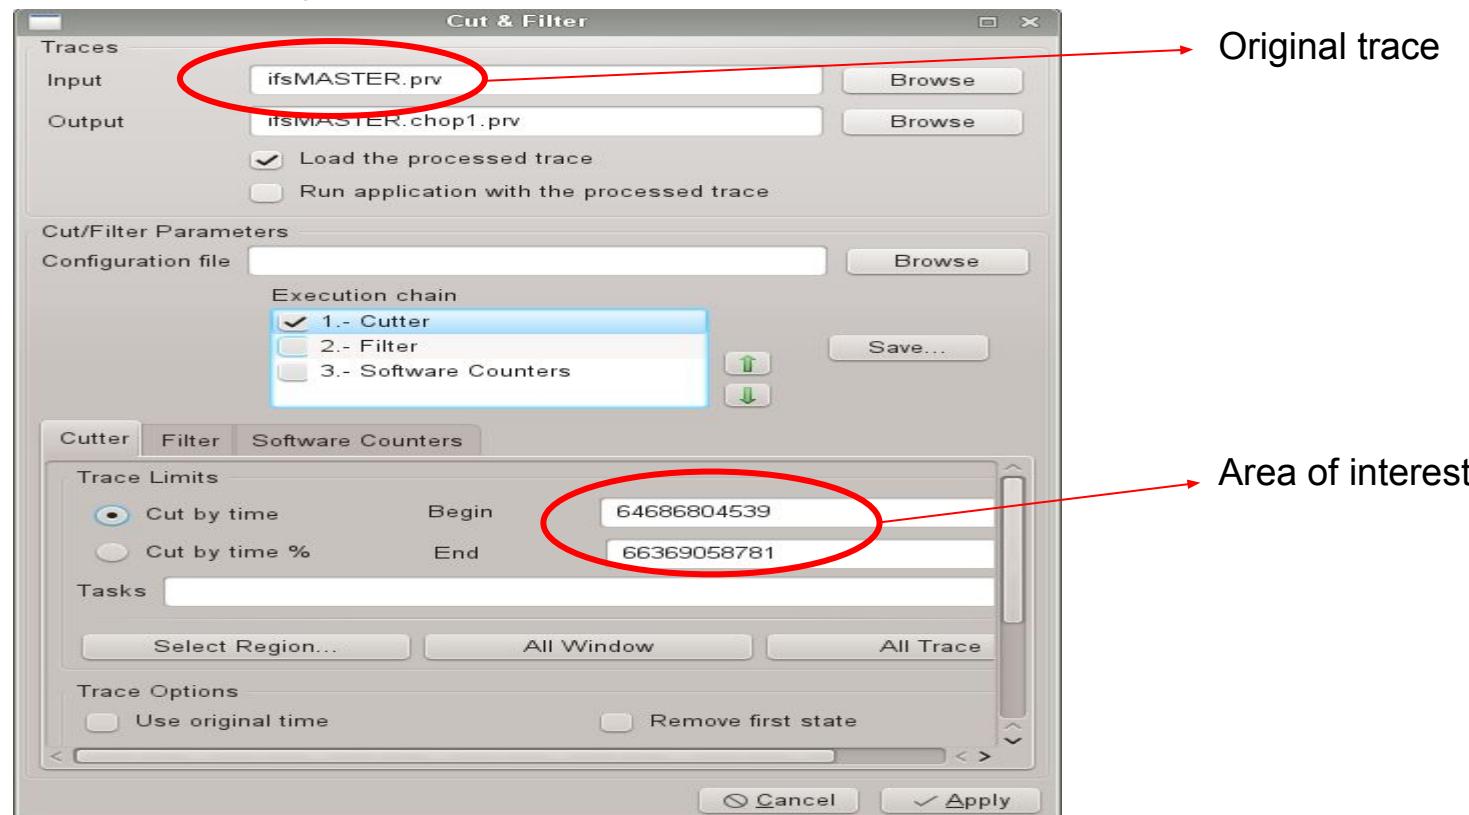


50000001  
50000002  
50000003  
50100001  
50100002  
50100003  
50100004

MPI events

# BSC Tools:Paraver

- Cutting
  - Cut original trace to obtain a fully detailed trace for the time interval considered representative or of interest
  - Use filtered trace to know the area of interest (remember that input must be the original trace)
    - Right click → run → cutter



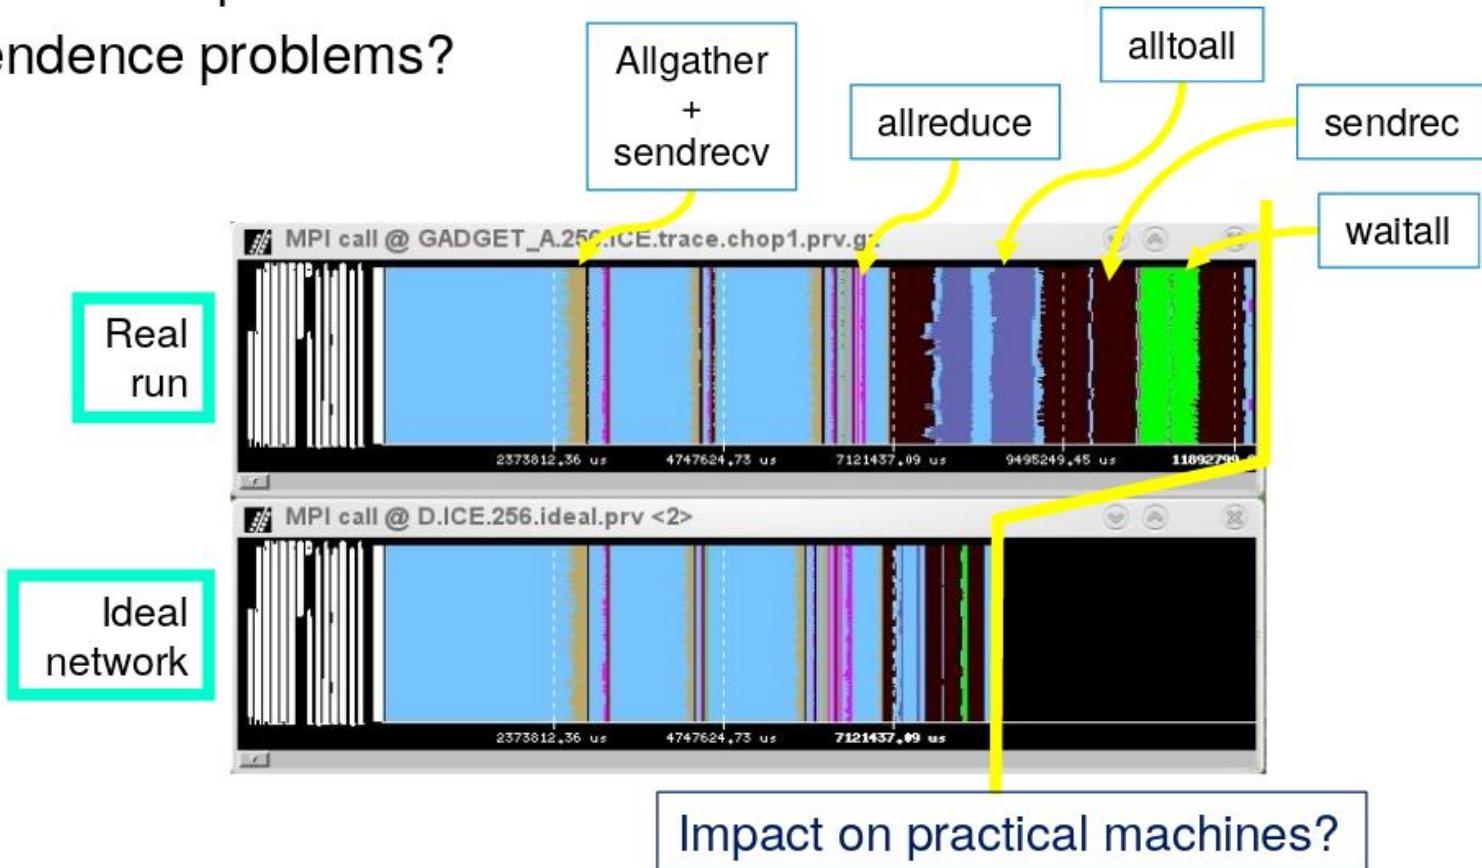
# BSC Tools:Paraver

- Configurations for analysis (usr/local/apps/paraver/X.X.X/cfgs)
  - General
    - Including basic views (timelines) and analysis (2D/3D profiles)
  - Counters\_PAPI
    - Hardware counters derived metrics
      - Program: related to algorithm/compilation (instructions, FP ops...)
      - Architecture:related to execution on specific architectures (cache misses...)
      - Performance: metrics reporting rating per time (MIPS, IPC...)
  - MPI → Views and analysis of MPI events
  - OpenMP → Views and analysis of OpenMP events
  - Complete Profile (general\_cfgs)

# BSC Tools:Dimemas

The impossible machine:  $BW = \infty$ ,  $L = 0$

- Actually describes/characterizes intrinsic application behavior
  - Load balance problems?
  - Dependence problems?



# Profiling Methodology

- Area of study
- Deployment efficiency
- Benchmarking
- Profiling analysis
- Validation

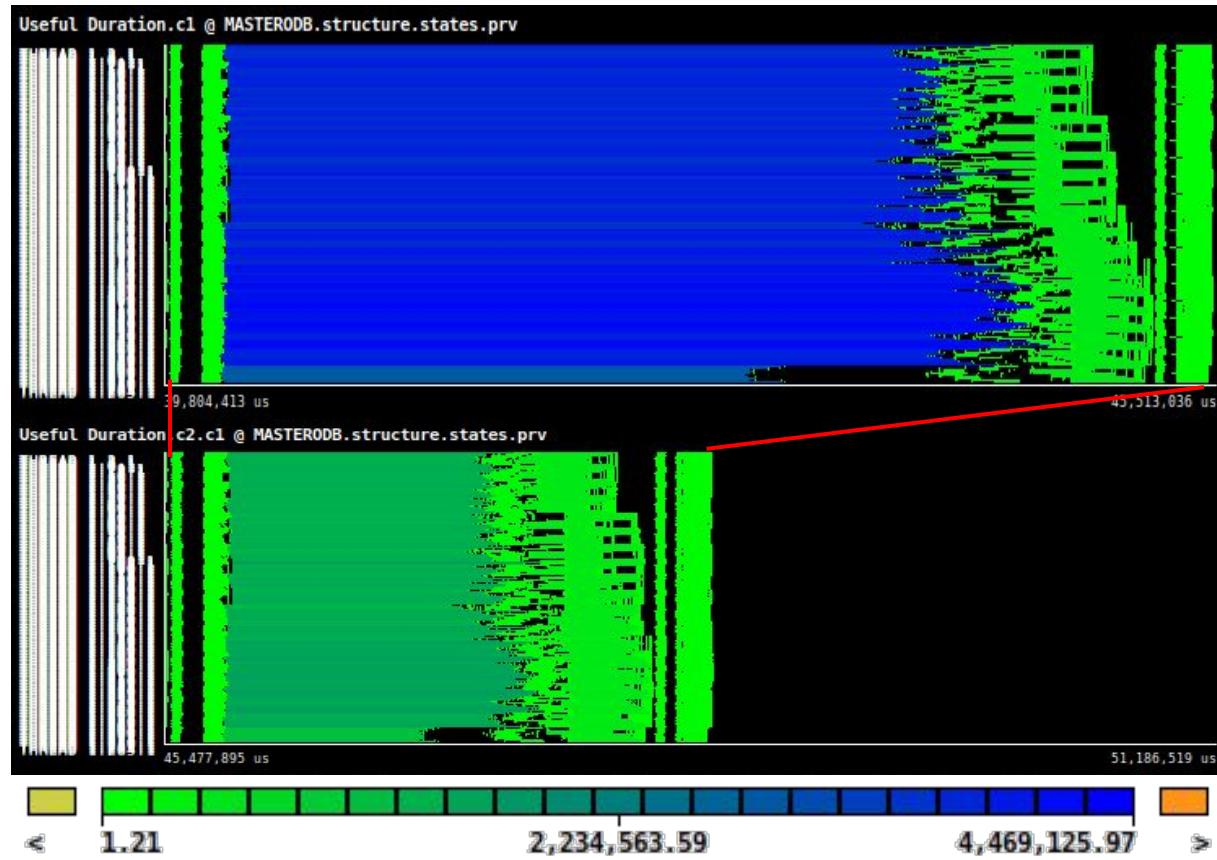
# Profiling Methodology

- **Area of study**
  - Configuration used (Operational, New algorithms, Global, Parallelization paradigm...)
  - Components activated and cyclic patterns
    - IO, ICE, Radiation, MPI, OpenMP
  - Area of study
    - 1 complete time step
- Deployment efficiency
- Benchmarking
- Profiling analysis
- Validation

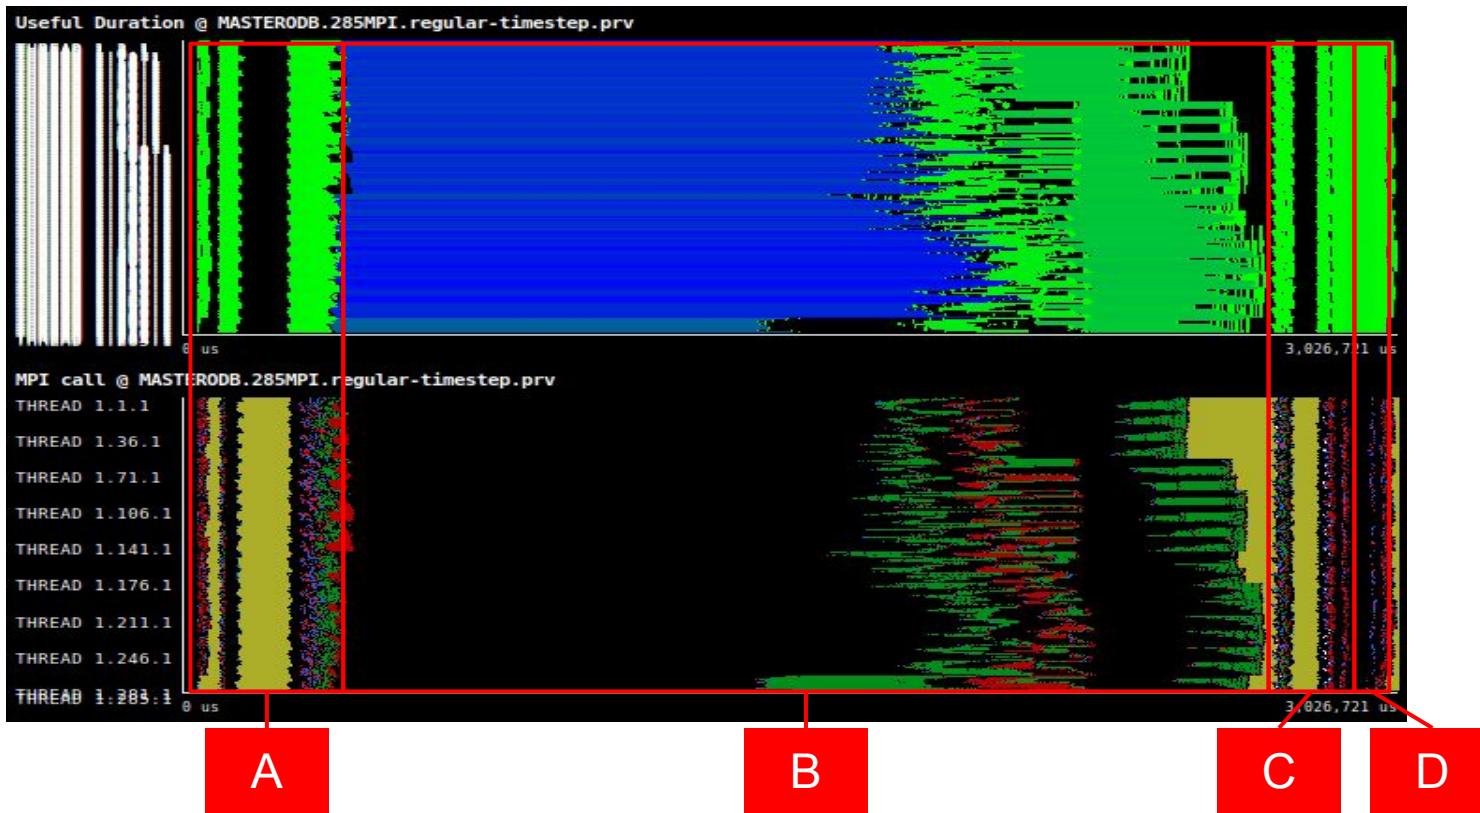
# Types of time step for the practical example

Time steps with radiation are much more expensive due to the extra computation in the grid-point part

Regular time step  
plus radiation



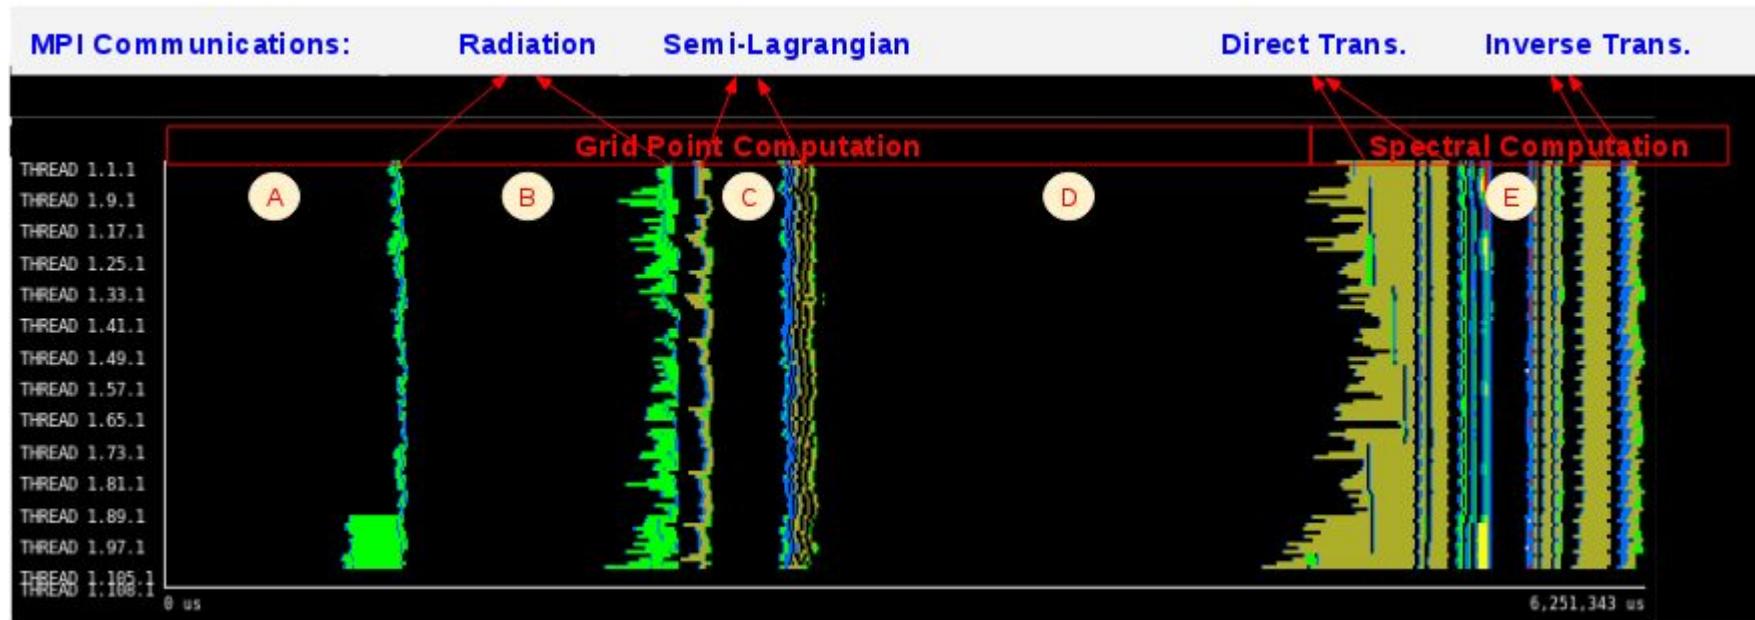
# Structure of a regular time step



- A - Inverse transformations
- B - Grid-point computations
- C - Direct transformations
- D - Spectral computations

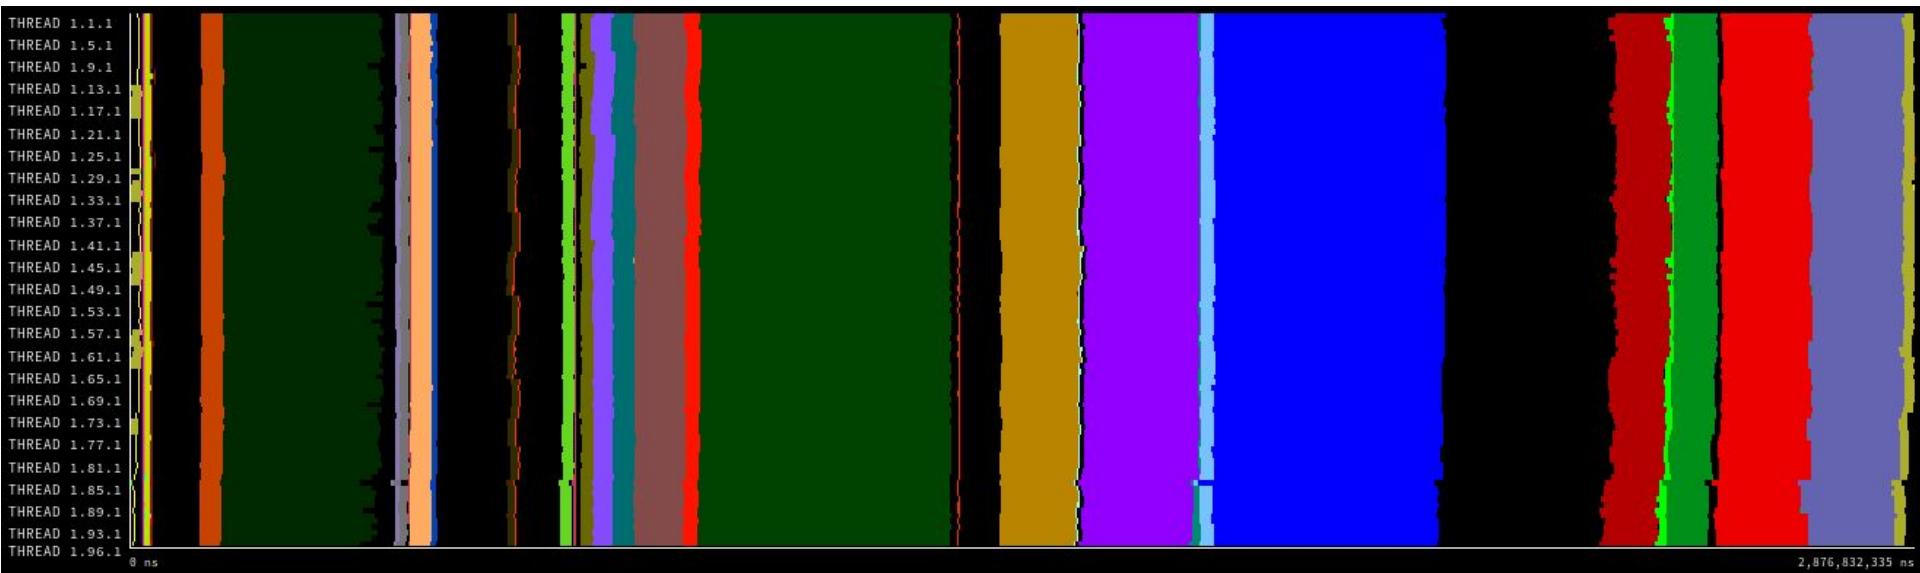
# Profiling Methodology

- Area of study (IFS)
  - 24 hours of simulation, T511L137 on CCA (ECMWF)
  - Selected 1 time step: 104 MPI processes + 4 IO (No OpenMP)
  - Metrics collected for large areas of computation automatically

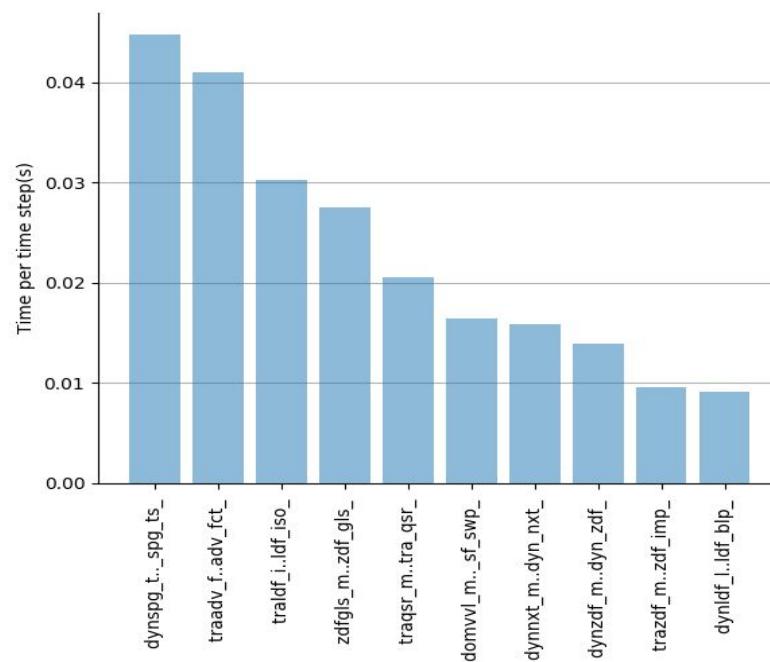
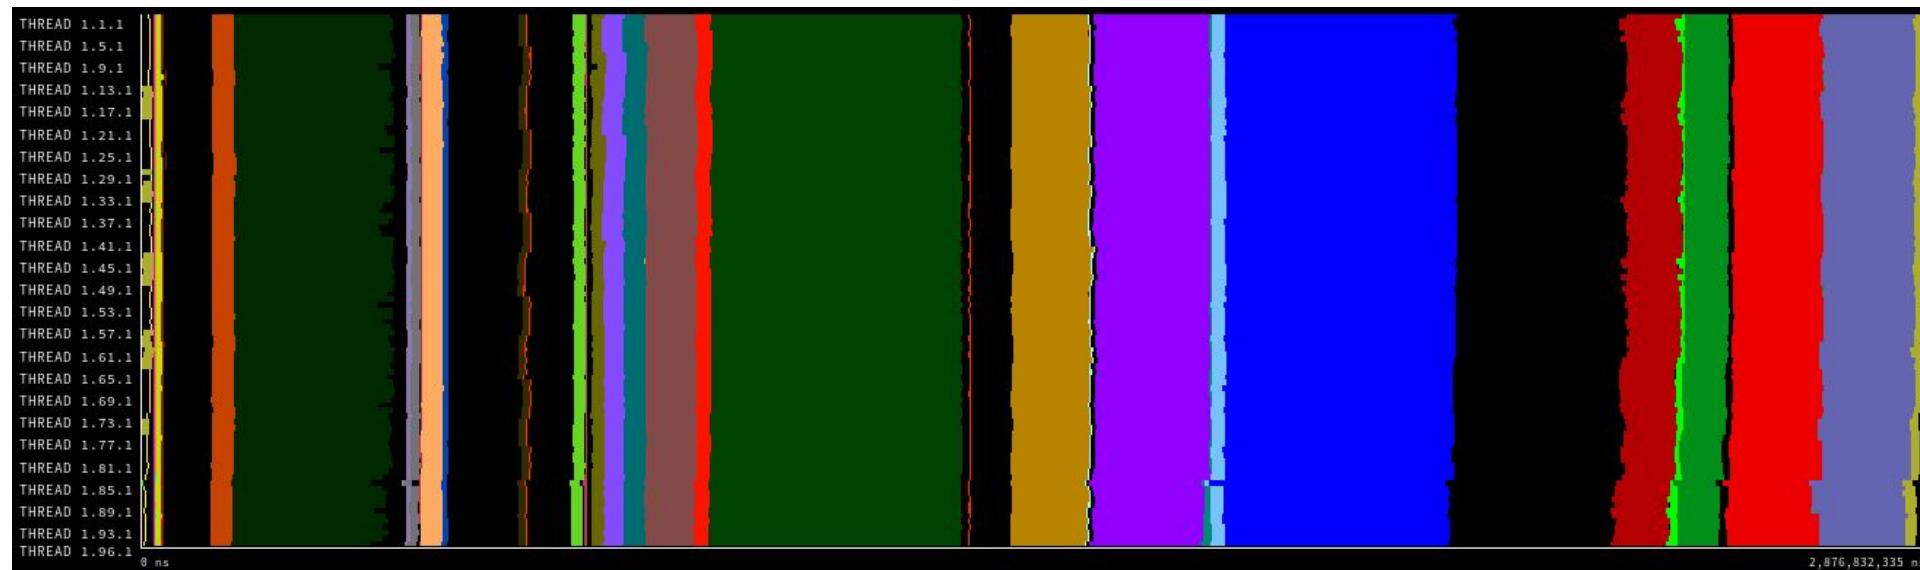


# Profiling Methodology

- Area of study (NEMO)
  - 1 day of simulation, ORCA025L91 on MN4 (BSC)
  - Selected the fastest time step automatically
  - 1 time step: 72 MPI processes (No IO, No OpenMP, No SI3)
  - Metrics collected for User functions manually



# Profiling Methodology



# Profiling Methodology

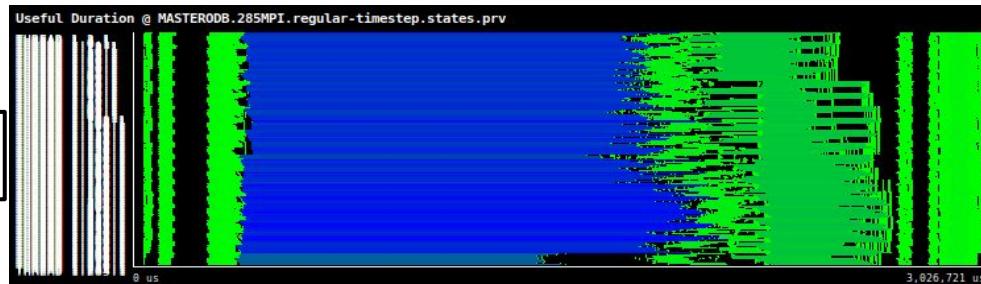
- Area of study
- Deployment efficiency
  - Compilation flags
    - Comparing fp options (fast, precise, strict...) and optimization options (OX, vectorization, approximations...)
    - Checking external libraries compilation
    - Debug flags (-g, Optimization reports, -f-instrument-functions...)
- Benchmarking
- Profiling analysis
- Validation

# Profiling Methodology

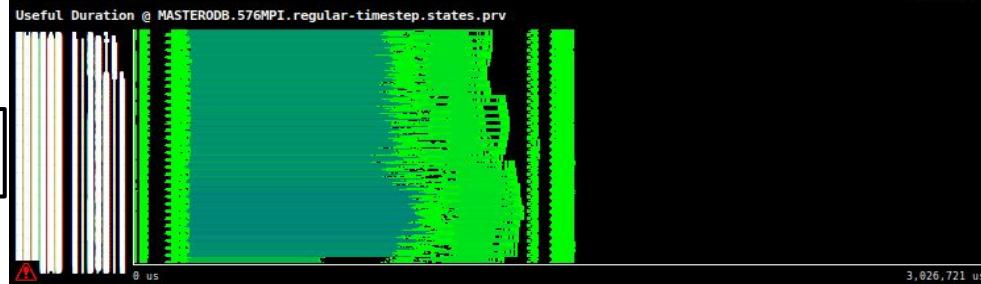
- Area of study
- Deployment efficiency
- **Benchmarking**
  - Basic Tests to collect Hardware metrics
    - Communications (Latency, Bandwidth, CPU, Parallel Efficiency...)
    - Weak and Strong scaling (MPI, OpenMP, Block processing and Hybrid sets)
    - Comparing optimizations (Double VS Single Precision...)
    - Extrae metrics collection and trace production
- Profiling analysis
- Validation

# MPI strong scaling: trace views

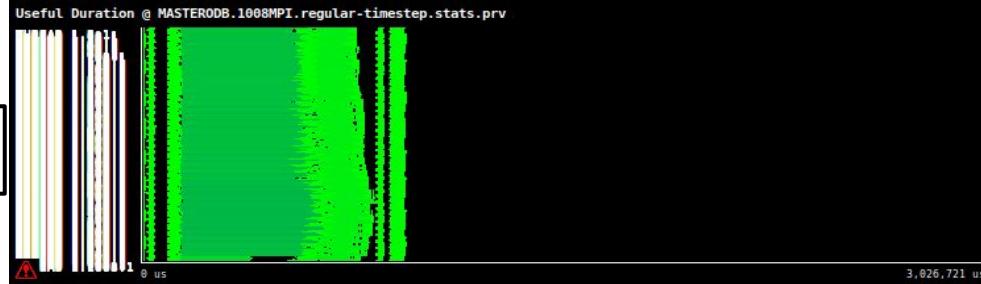
285 MPI (8 nodes)



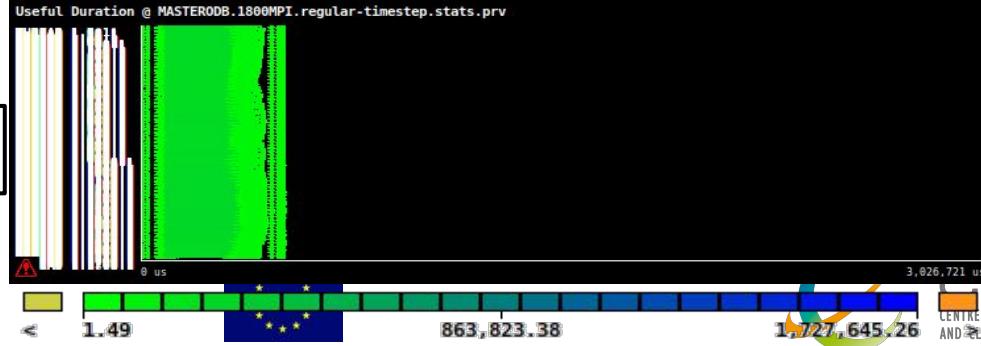
576 MPI (16 nodes)



1008 MPI (28 nodes)

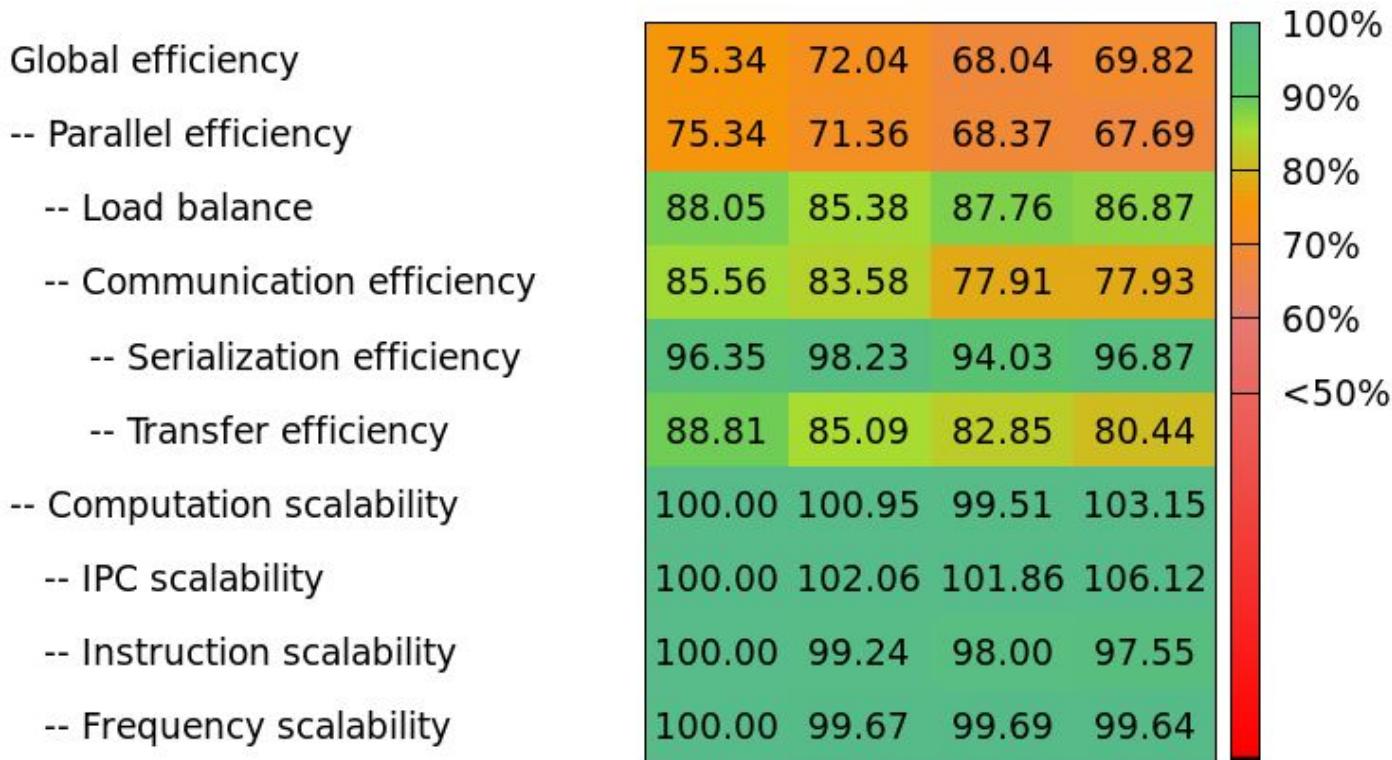


1800 MPI (50 nodes)



# Basic Analysis: MPI Strong Scaling

- Computation and parallel efficiency factors for MPI only:
  - Good computation scalability and serialization efficiency
  - Not very good load balance neither transfer efficiency



# Basic Analysis: Double P VS Single P

Overview of the collected raw data:

	108	108
Runtime (us)	110741508.76	71238767.9
Runtime (ideal)	105675625.64	68396939.23
Useful duration (average)	88427932.03	57382830.24
Useful duration (maximum)	94410288.2	61484222.58
Useful duration (total)	9196504931.3	5967814345.21
Useful duration (ideal, max)	94410288.2	61484222.58
Useful instructions (total)	26798422515714	23201423473963
Useful cycles (total)	21985000332874	14299301515415

Overview of the computed model factors:

	108	108
Parallel efficiency	79.85%	80.55%
Load balance	93.66%	93.33%
Communication efficiency	85.25%	86.31%
Serialization efficiency	89.34%	89.89%
Transfer efficiency	95.43%	96.01%
Computation scalability	100.00%	154.10%
Global efficiency	79.85%	124.13%
IPC scalability	100.00%	133.11%
Instruction scalability	100.00%	115.50%
Frequency scalability	100.00%	100.23%
Speedup	1.00	1.55
Average IPC	1.22	1.62
Average frequency (GHz)	2.39	2.40

# Profiling Methodology

- Area of study
- Deployment efficiency
- Benchmarking
- **Profiling analysis**
  - MPI and OpenMP profile summary and Basic Analysis Tool
  - PAPI counters
  - MPI and OpenMP evaluation in detail
  - Clustering and Tracking Tools
  - Sampling and Folding Tools
  - Connection to the code
  - Dimemas Tool
- Validation

# MPI Profile Summary

Parallel and Communication efficiency, Global load balance → less than 85%?

Parallel Efficiency

IFS

The screenshot shows a table from the IFS MPI profile summary. The columns represent various MPI operations and their percentages of total execution time. Red arrows point from specific cells to annotations: one arrow points from the 'Avg/Max' row to 'Global Load Balance', another from the 'MPI\_Send' column to 'Communication Efficiency', and a third from the 'MPI\_Send' cell in the 'Avg/Max' row to the 'MPI\_Send' column header.

	Outside MPI	MPI_Send	MPI_Recv	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_BARRIER	MPI_Alltoallv	MPI_Gatherv	MPI_Comm_rank	MPI_Comm_size	MPI_Bsend	MPI_Waitany
Total	66,578.44 %	1.71 %	773.76 %	646.21 %	239.35 %	12,362.37 %	806.93 %	10,757.31 %	35.56 %	2.49 %	448.23 %	0.81 %	7,746.82 %
Average	66.31 %	0.00 %	0.77 %	0.64 %	0.24 %	12.31 %	0.80 %	10.71 %	0.04 %	0.00 %	0.45 %	0.81 %	7.72 %
Maximum	72.93 %	0.01 %	2.98 %	1.60 %	0.80 %	18.56 %	1.84 %	25.06 %	1.12 %	0.01 %	1.88 %	0.81 %	19.25 %
Minimum	57.05 %	0.00 %	0.01 %	0.08 %	0.07 %	3.11 %	0.00 %	5.25 %	0.00 %	0.00 %	0.16 %	0.81 %	0.31 %
StDev	2.03 %	0.00 %	0.57 %	0.36 %	0.06 %	2.52 %	0.41 %	3.57 %	0.12 %	0.00 %	0.10 %	0 %	3.18 %
Avg/Max	0.91	0.31	0.26	0.40	0.30	0.66	0.44	0.43	0.03	0.34	0.24	1	0.40

Global Load Balance

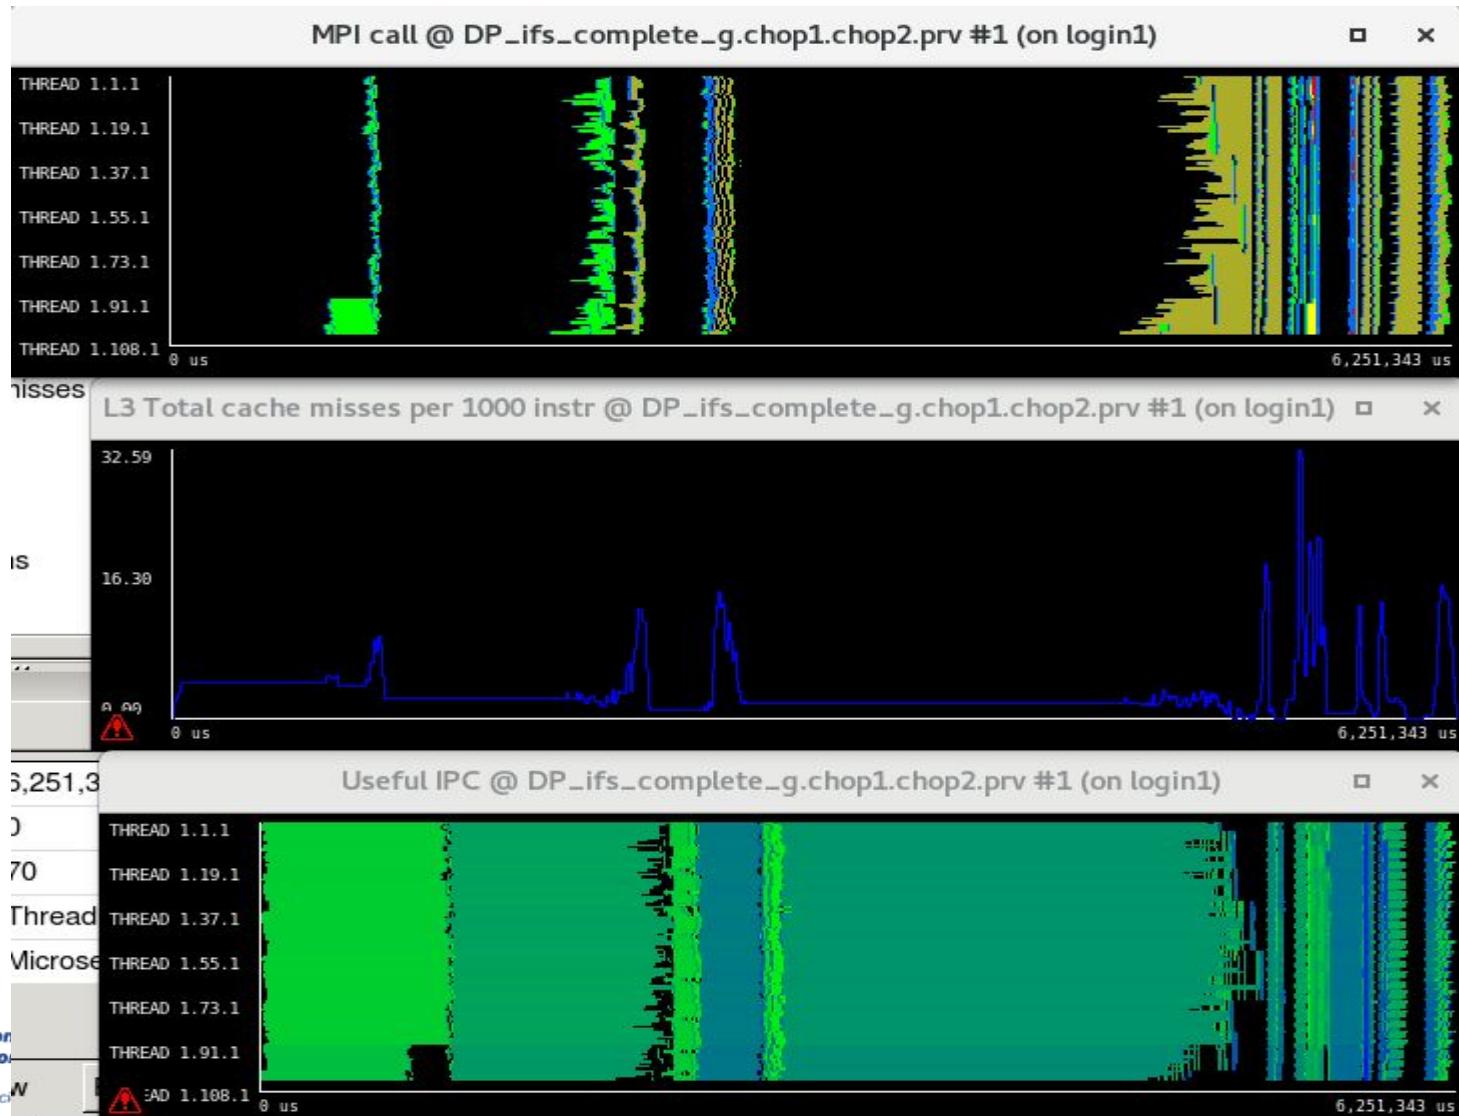
Communication Efficiency

# PAPI Counters

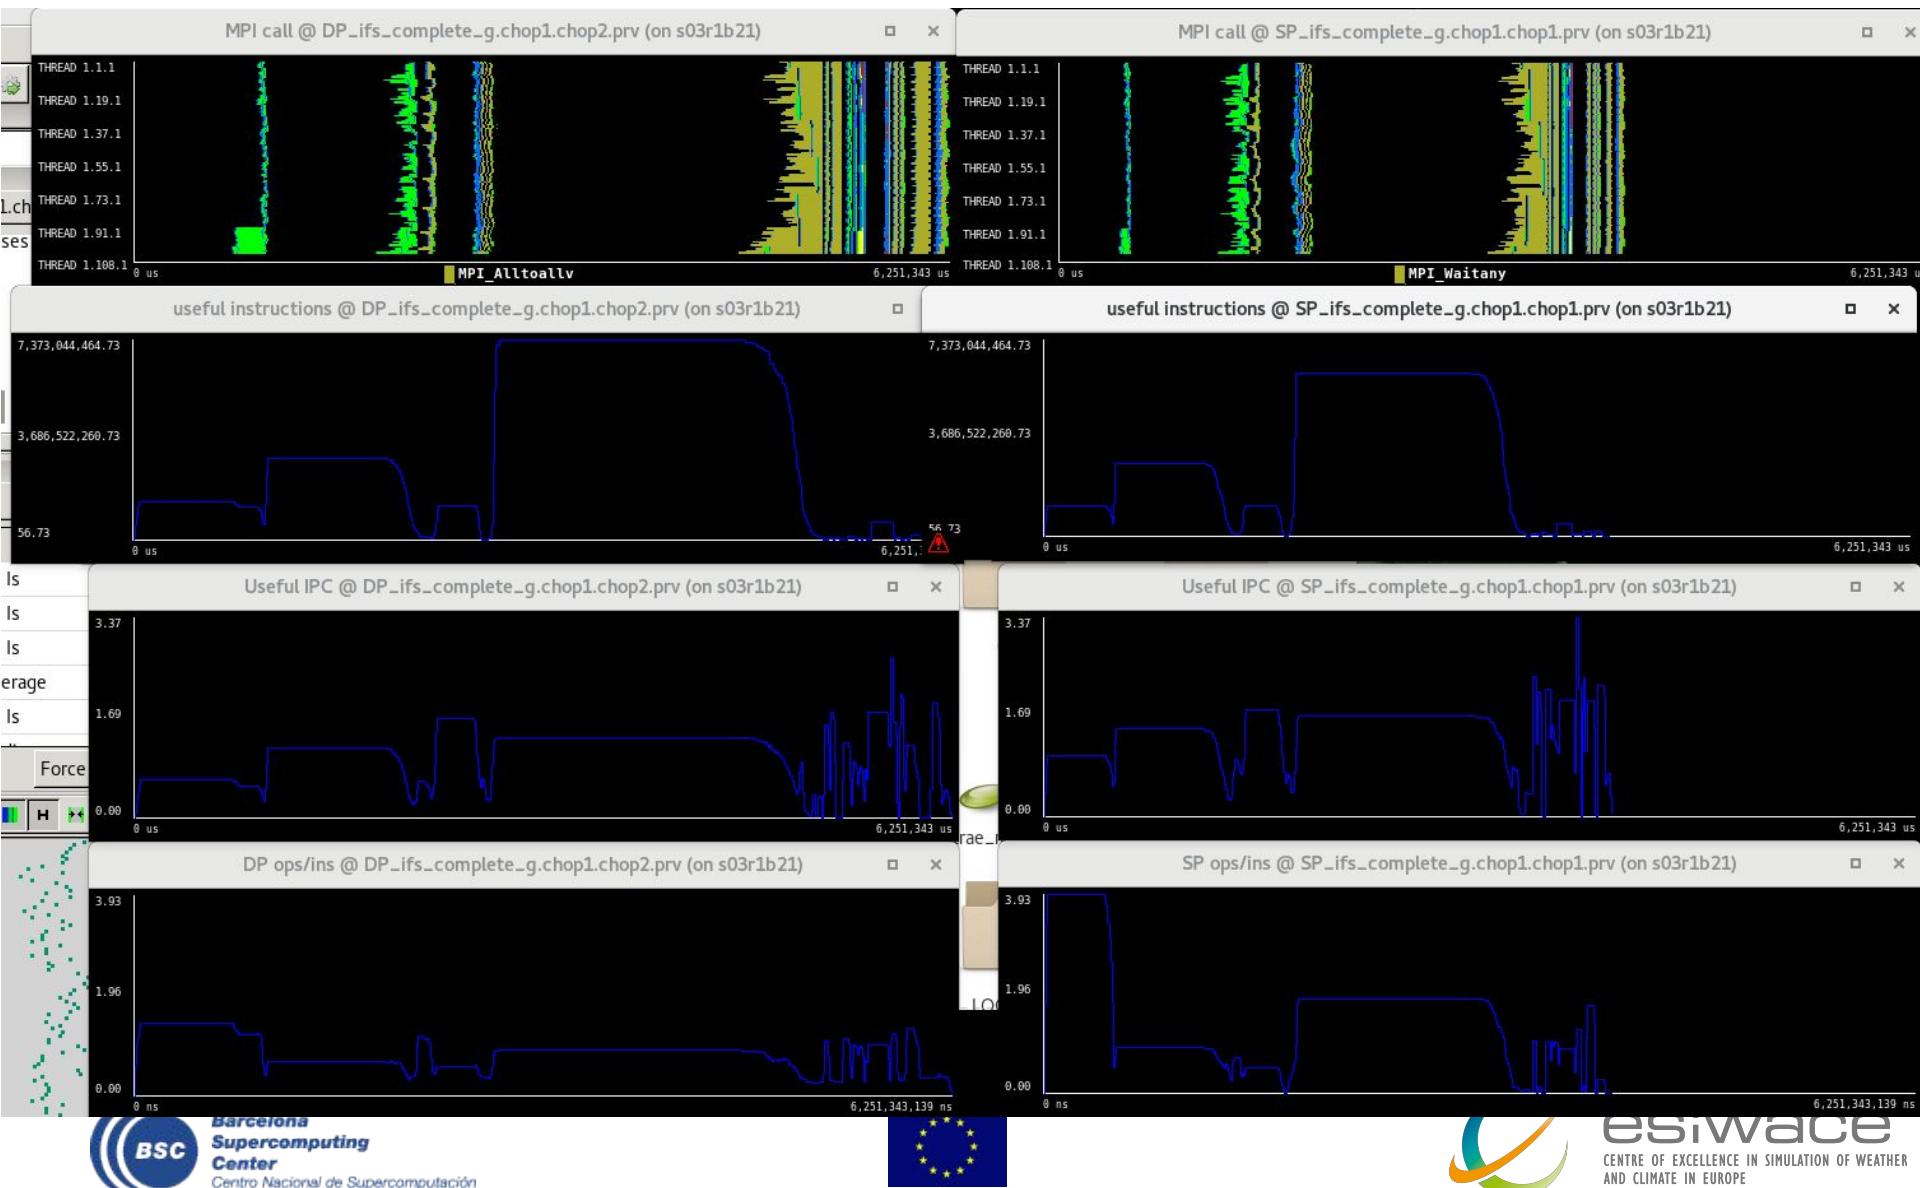
- PAPI counters collected during the execution
- Some of them are based on other native PAPI counters and derived from the base metrics

	Derived
Instructions	
Cycles	
Useful Duration	X
Useful Instructions	X
Useful IPC	X
Loads	
Stores	
L3/L2/L1_Total_Misses	
L3/L2/L1_MISS_RATIO	X
FP_OPS	
FP_TOT_INS	
INS_VEC	X

# PAPI Counters

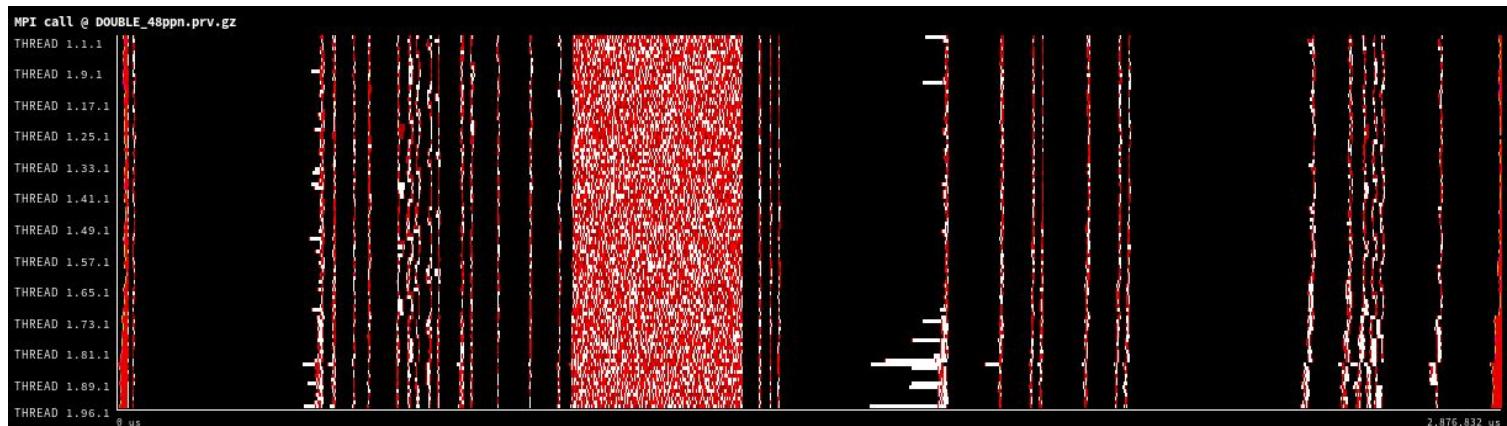


# PAPI Counters

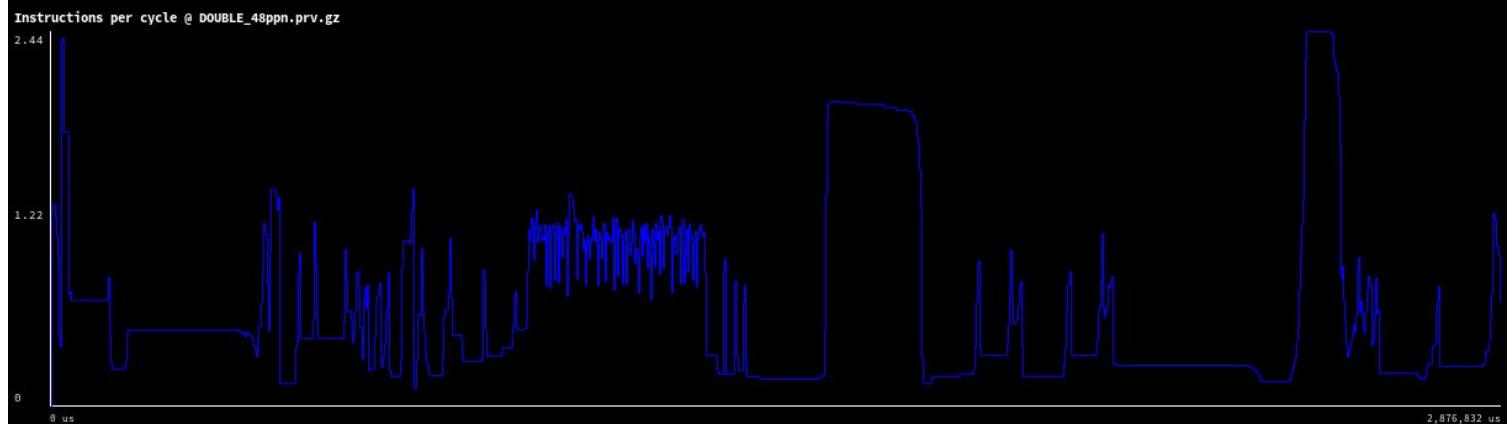


# PAPI Counters

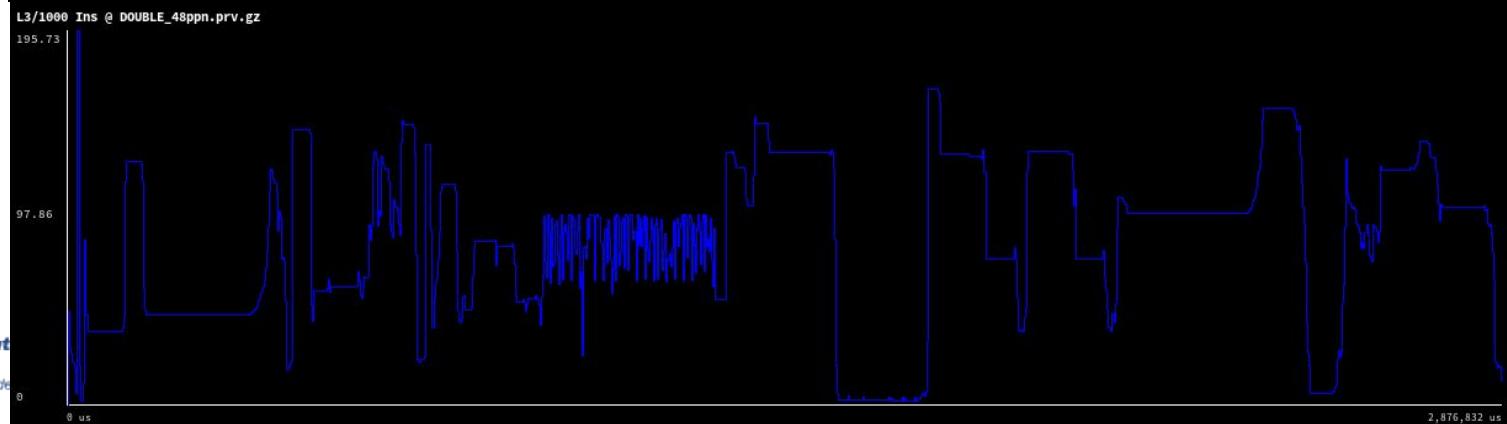
MPI Events



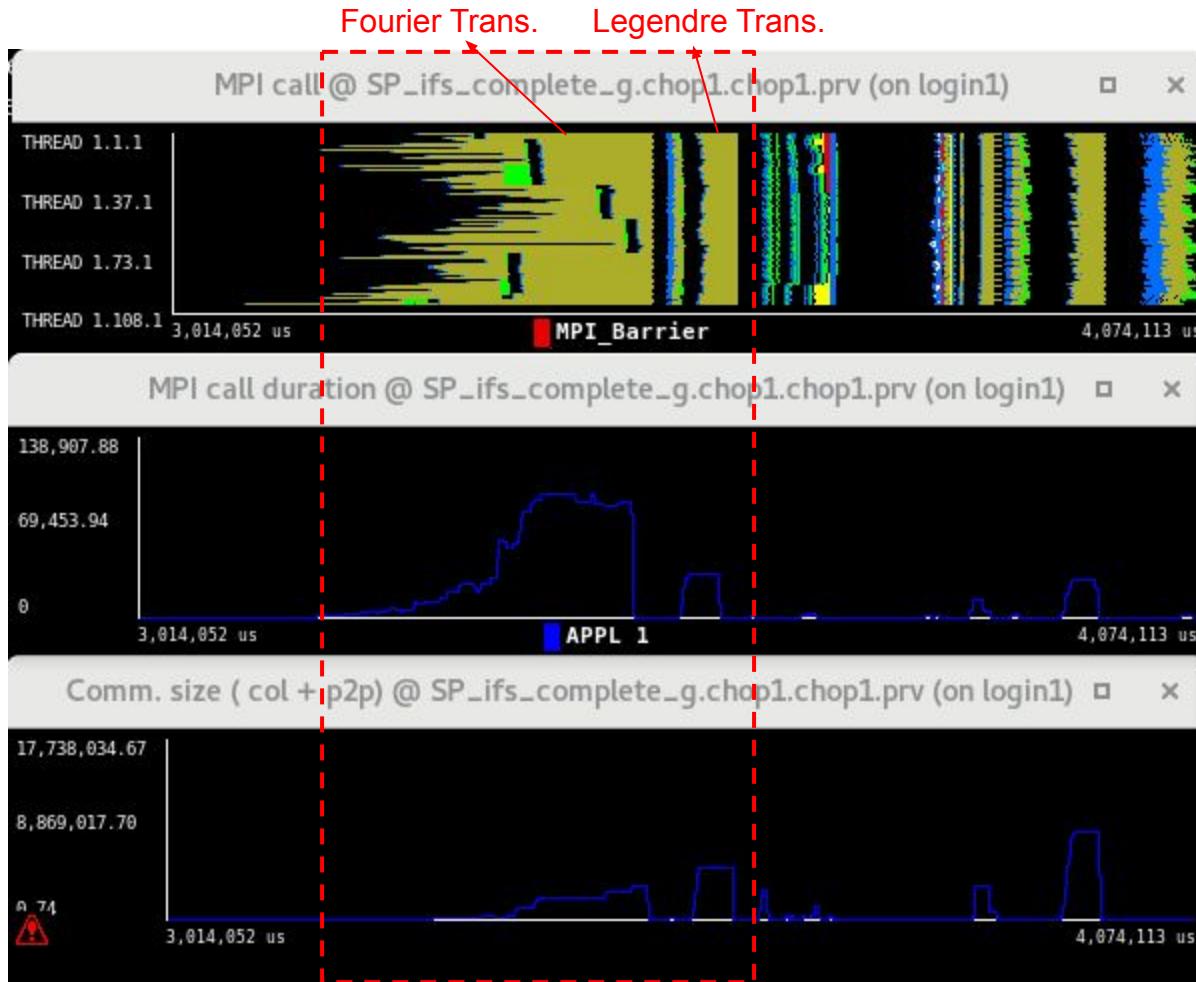
IPC



L1 Misses  
per 1000  
INS

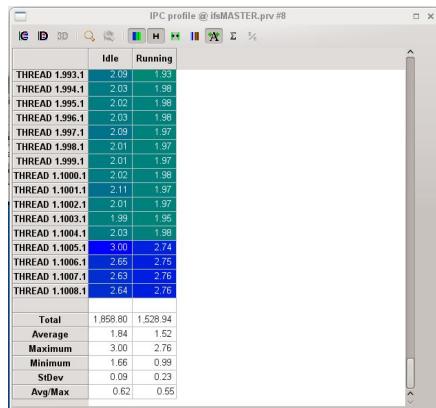


# MPI evaluation

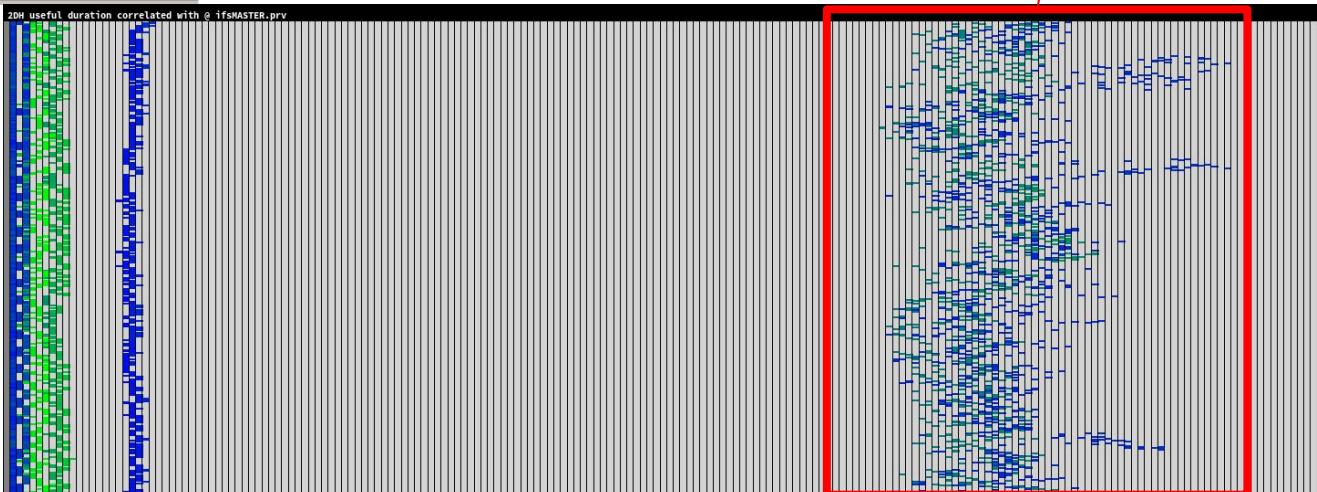
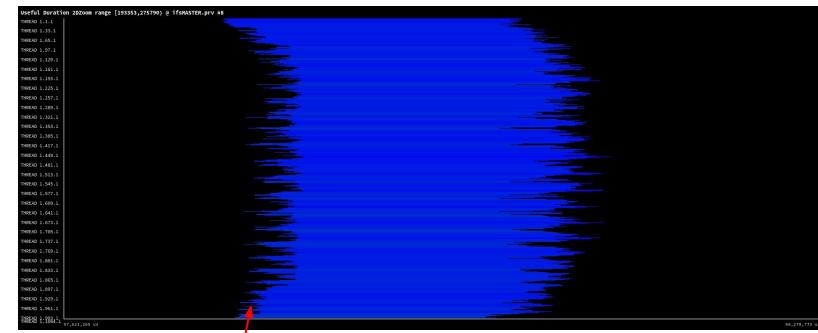


# MPI evaluation

- IPC less than 1 for calculation areas?
- Are there load imbalance regions?

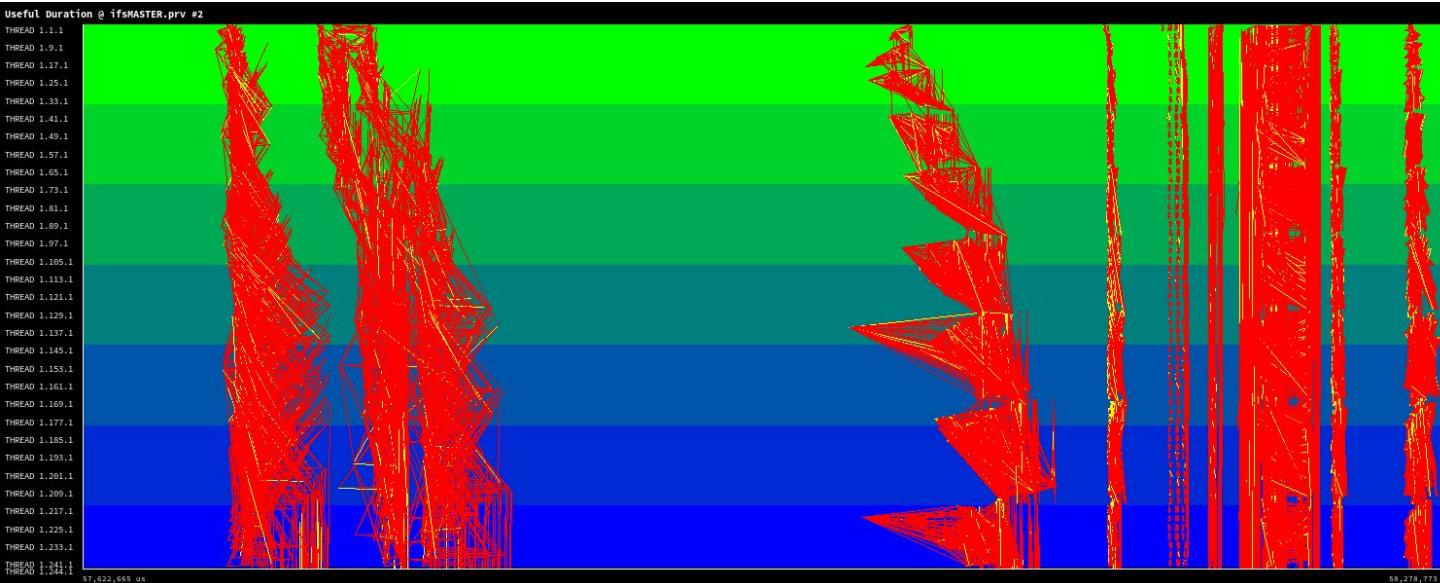


IPC\_Profile



# MPI evaluation

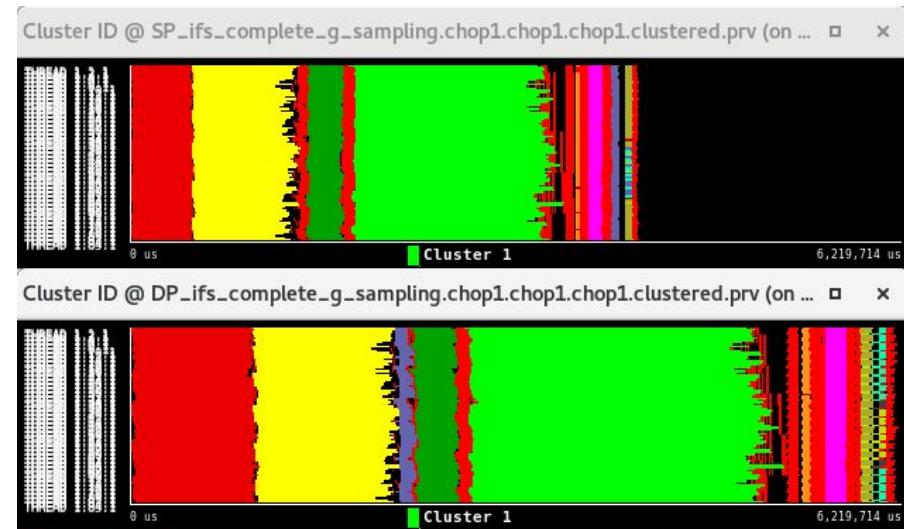
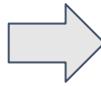
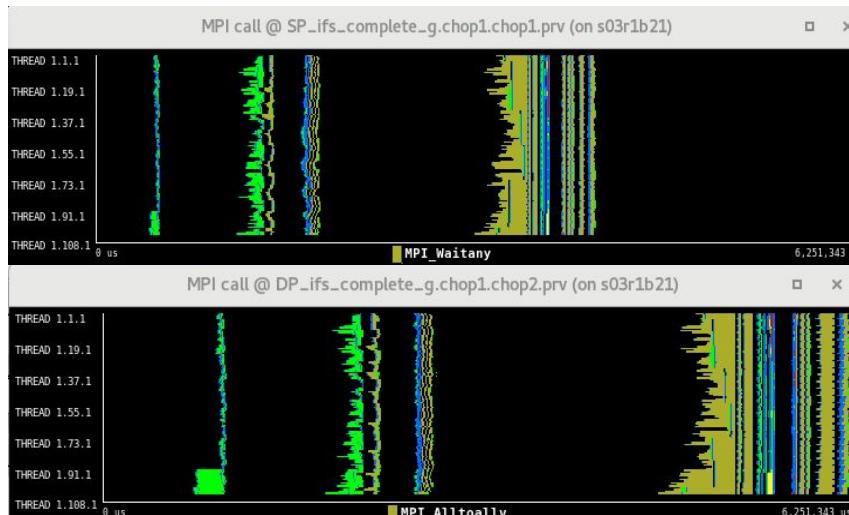
- Are MPI communications efficient according to the map affinity?



Affinity per node

# Clustering Tool

Applying Clustering for an automatic profiling analysis

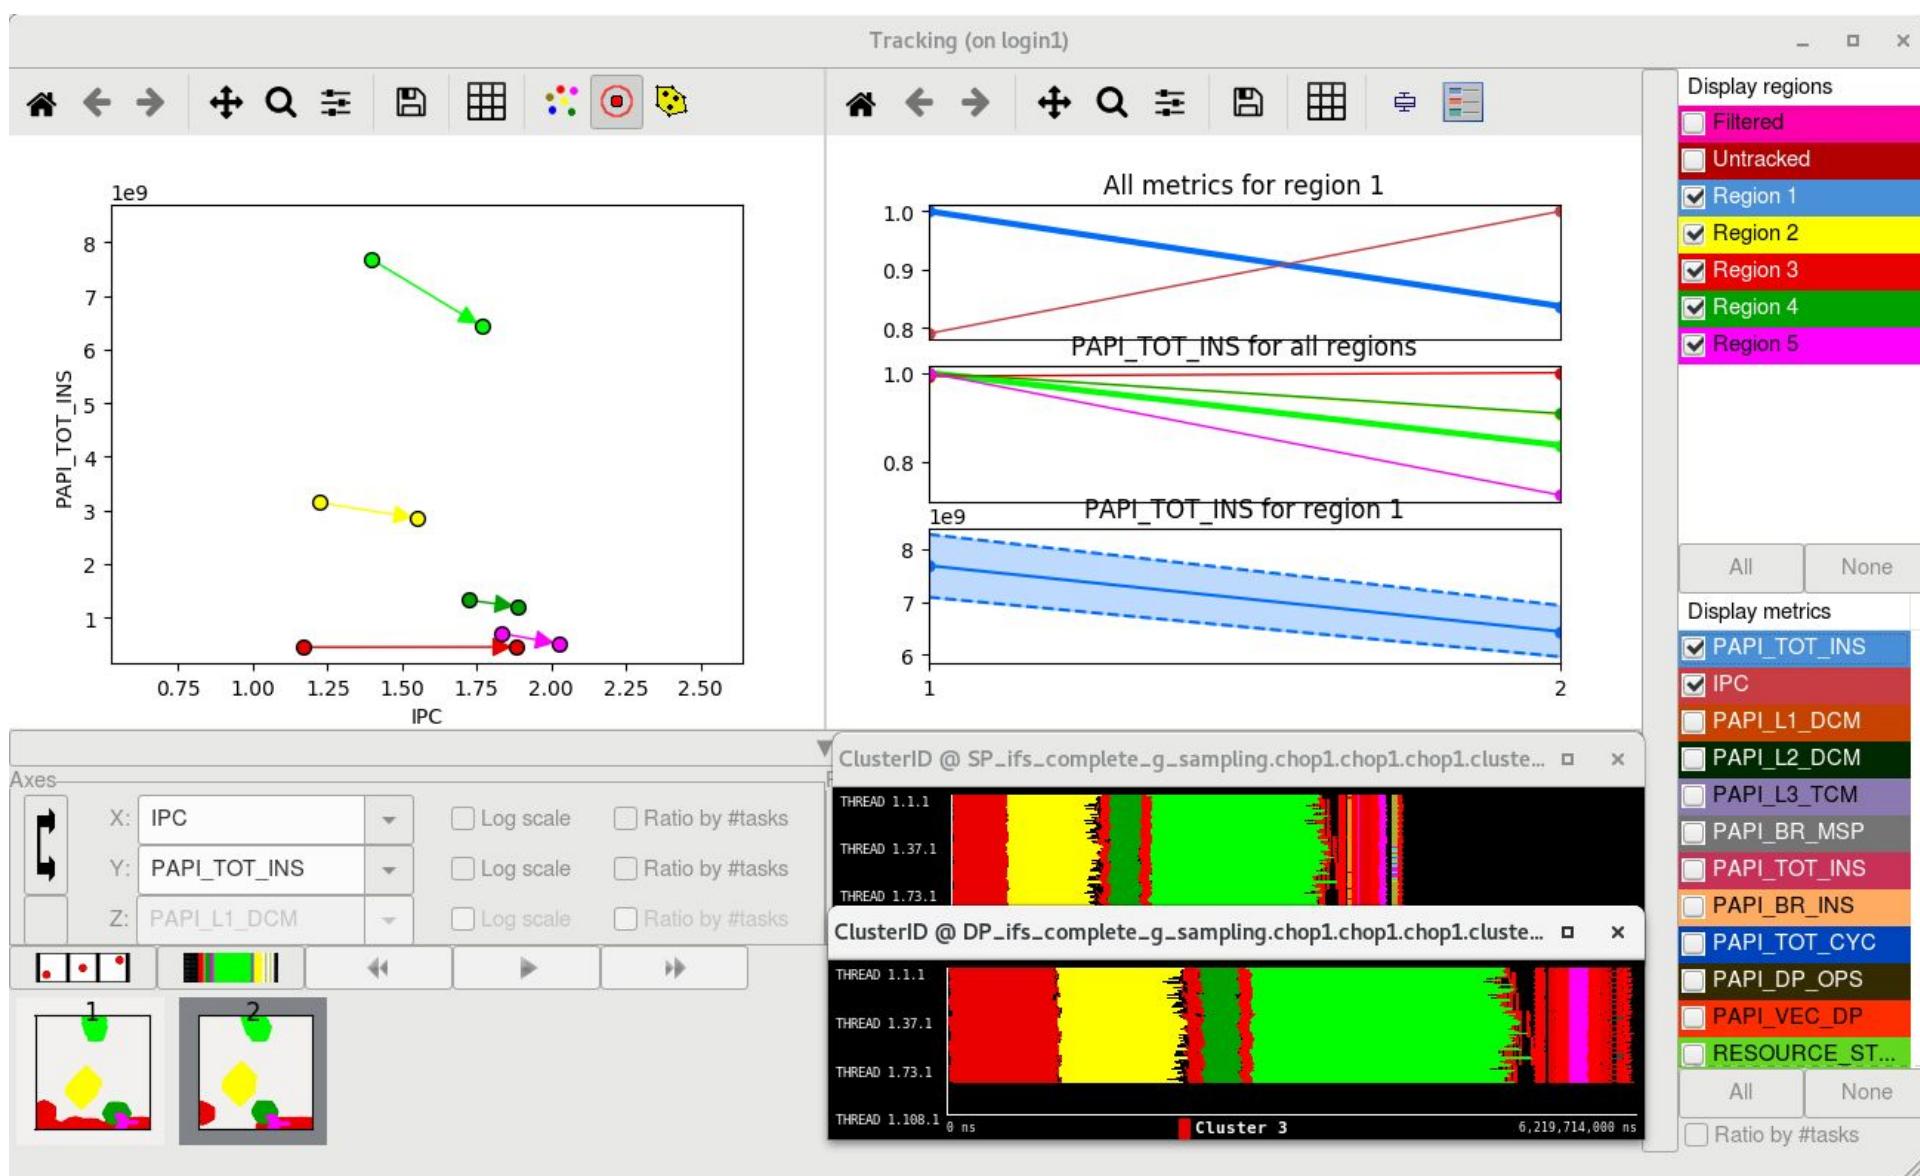


- Characterizes computing bursts that are similar and groups them into clusters
- Allows to study the behavior of the clusters separately, identify patterns, etc.

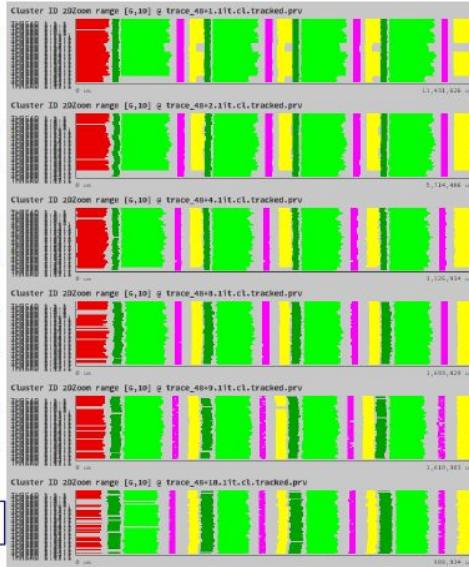
# Tracking Tool

- A friendly way to quantify and visualize the evolution of the clusters among several traces
- The tool has 2 parts
  - Recognition algorithm of “who-is-who”, based on heuristics
  - A visualization GUI
- Examples analyzing multiple traces
  - Scaling number of MPI/OpenMP resources (64 – 128 – 256...)
  - Testing different microarchitecture features
  - Changing the problem size
  - Trying different compiler optimizations

# Tracking Tool

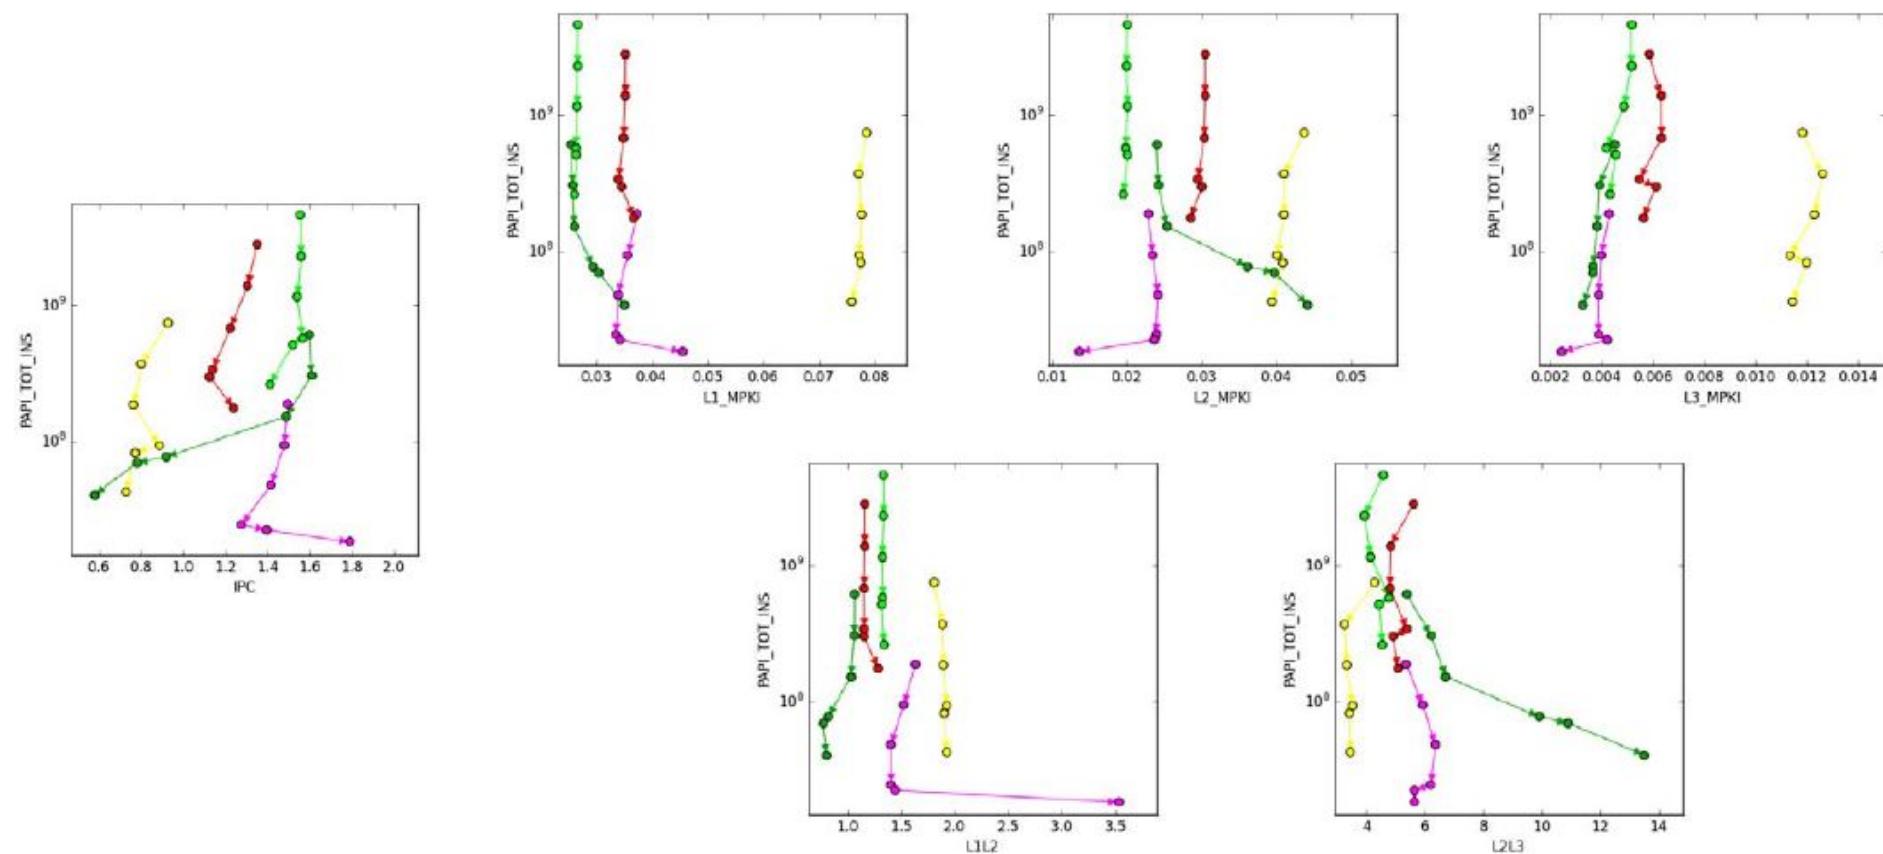


48x1



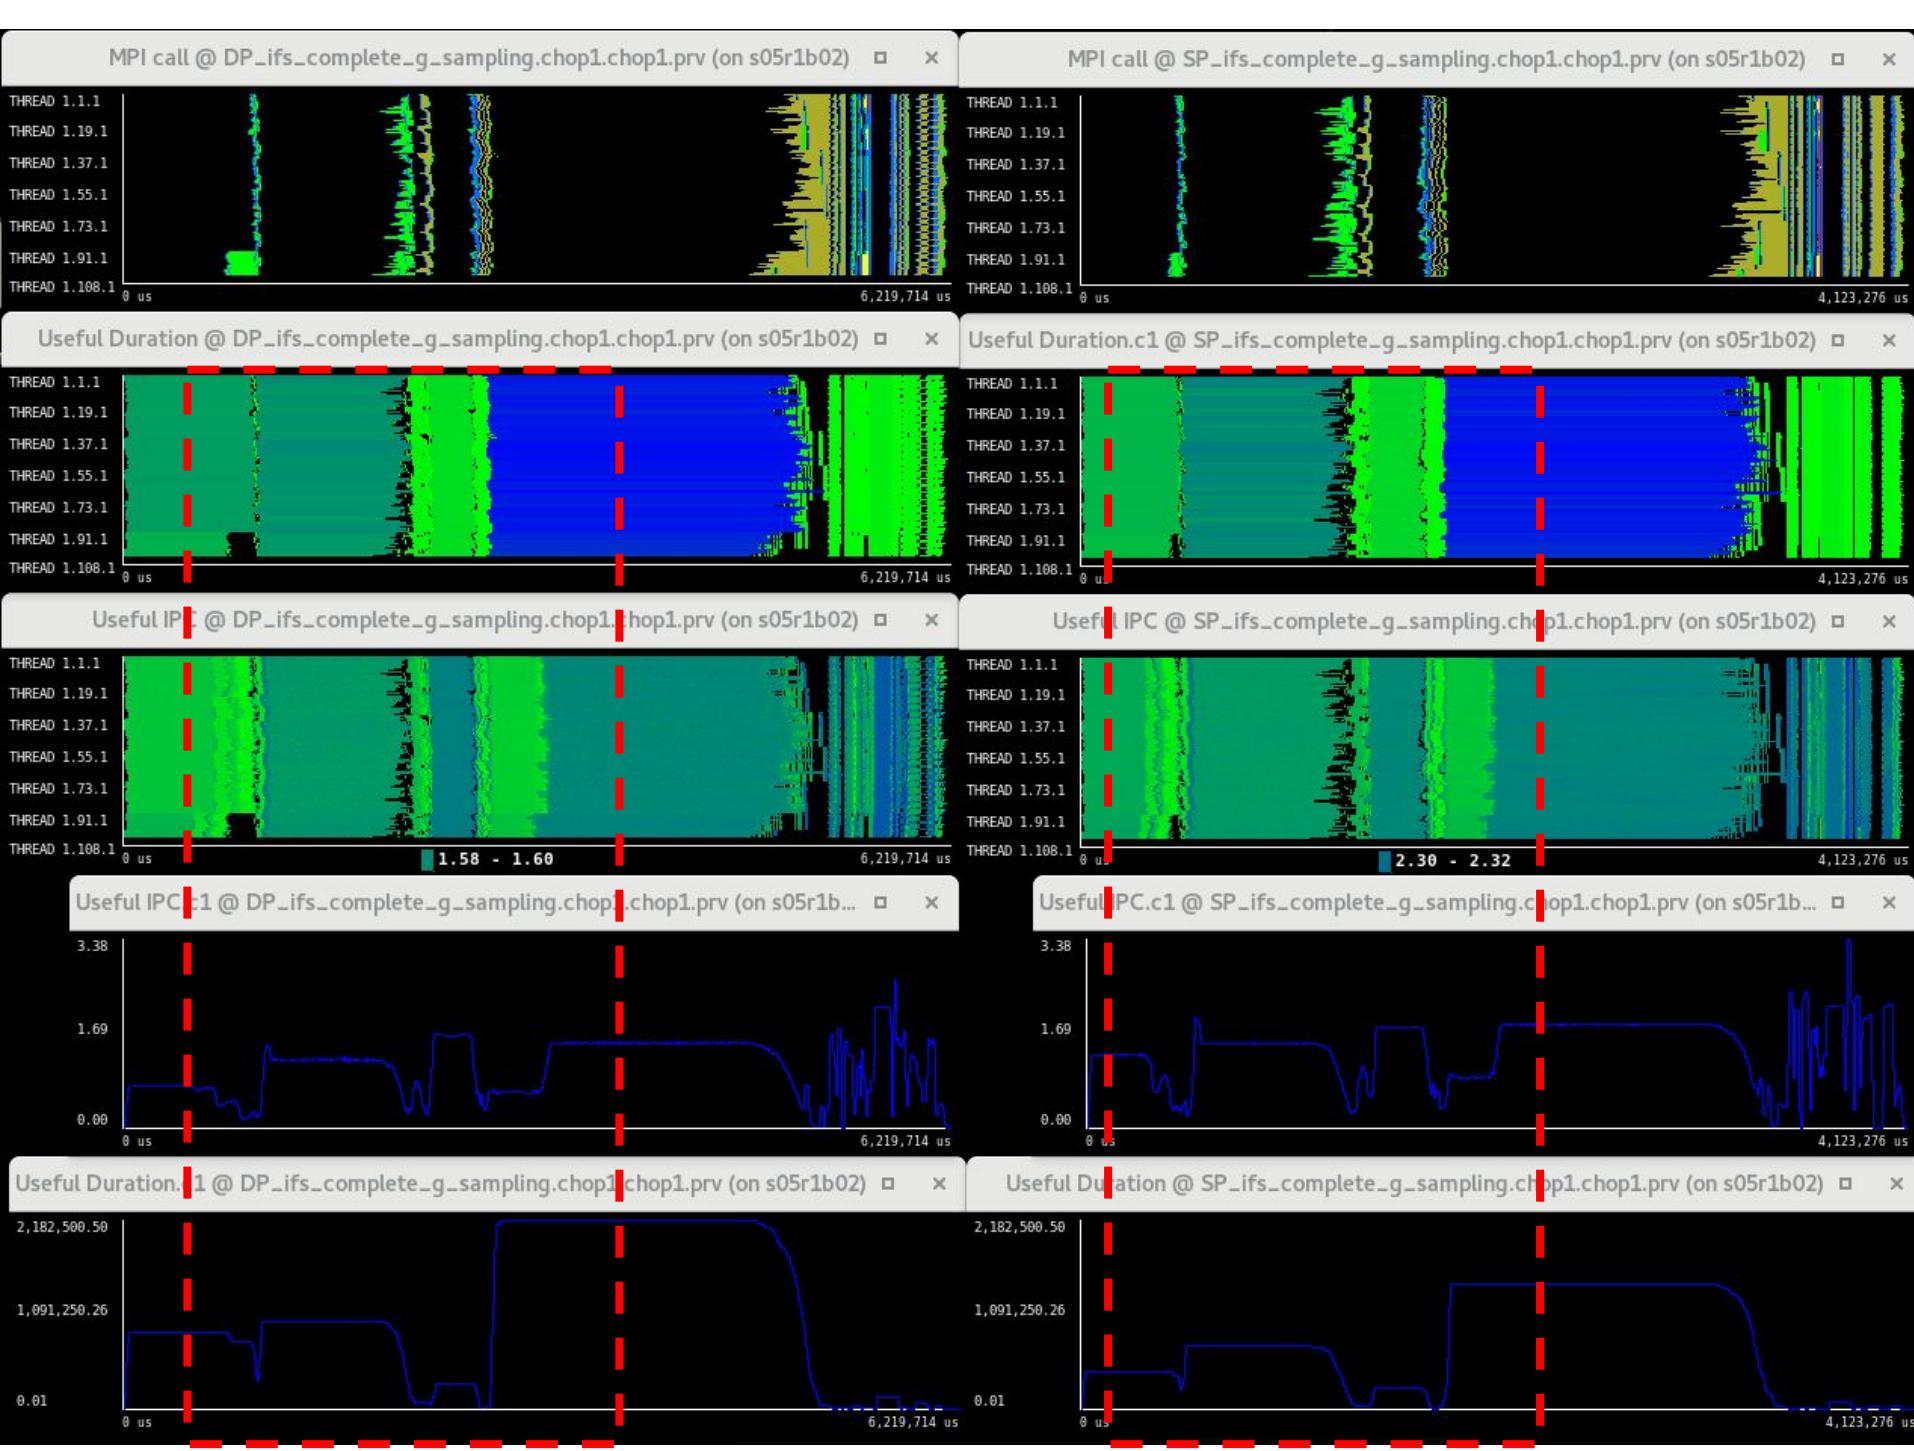
# Tracking Tool

## Tracking IFS MPI+OMP Strong Scaling



# Sampling Tool

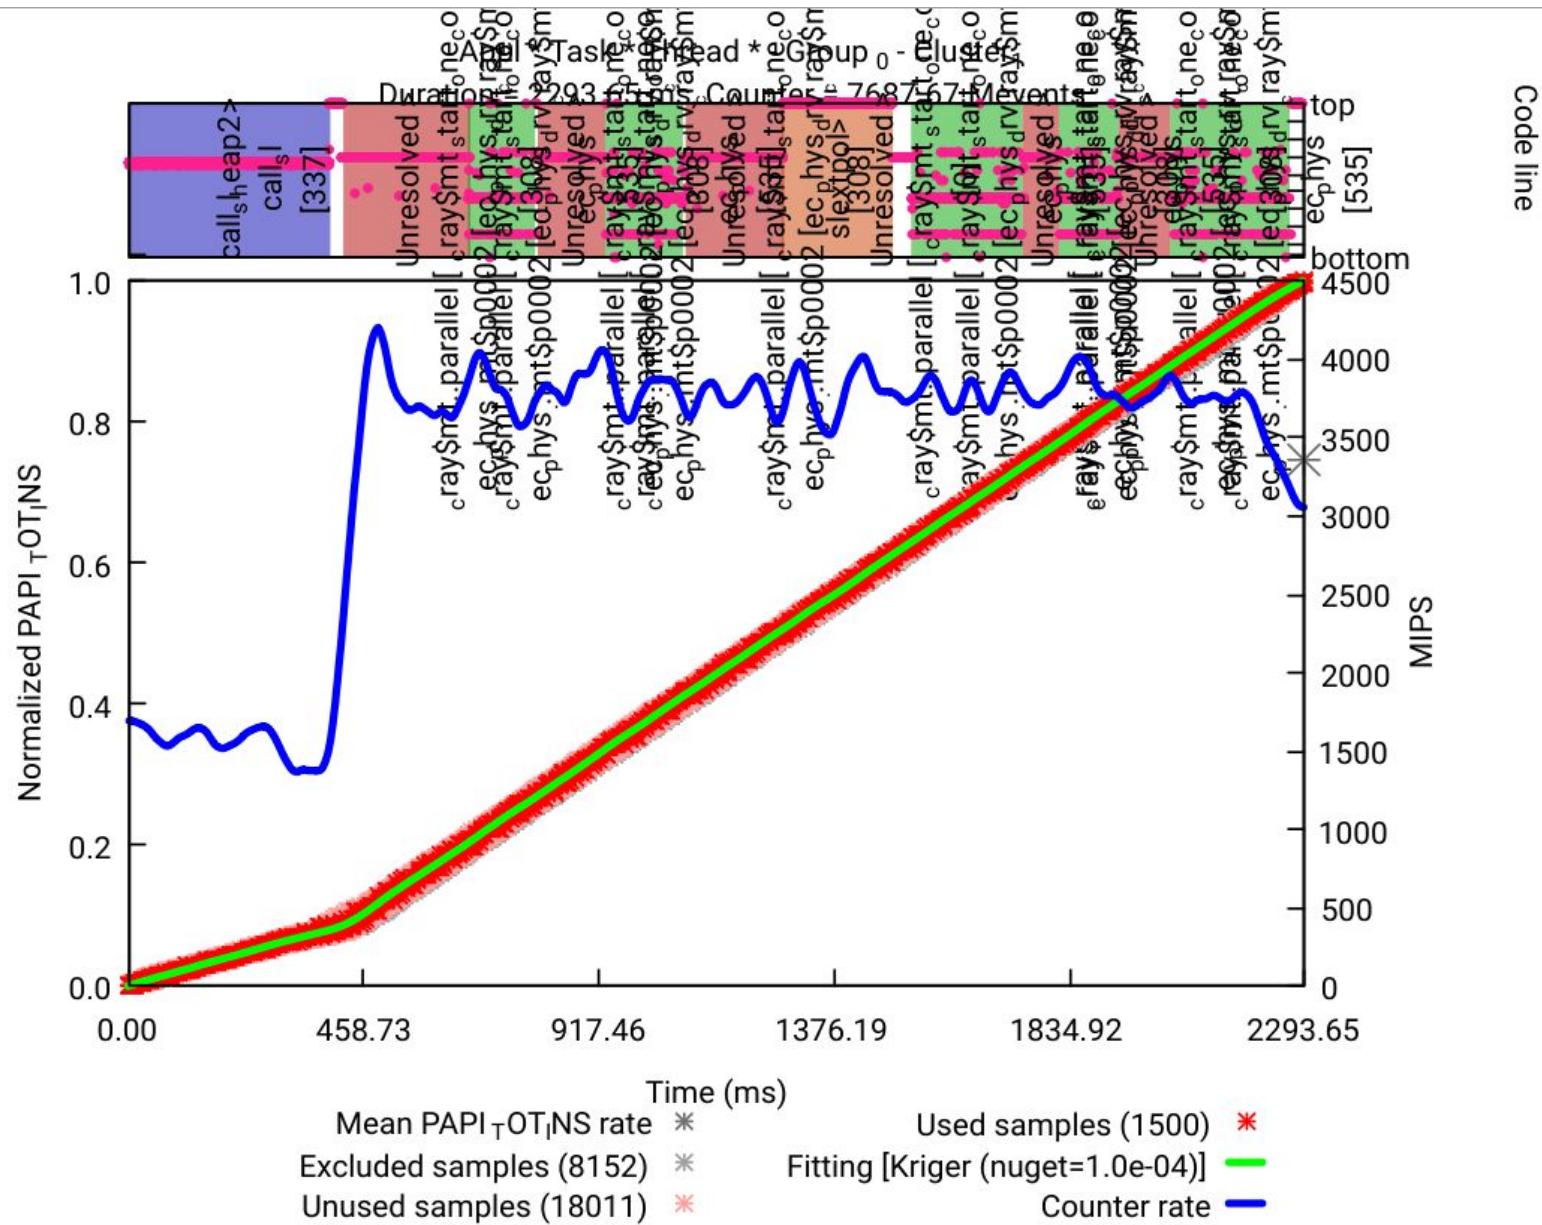
- Extrae can be configured to capture performance metrics on a periodic basis using alarm signals and specifying period and variability (10 and 2 respectively for IFS and NEMO tests).
- This means that we will capture samples every 10 ms with a random variability of 2 ms.
- Every sample contains processor performance counters (where every PAPI counter is referred at configured time) and callstack information.



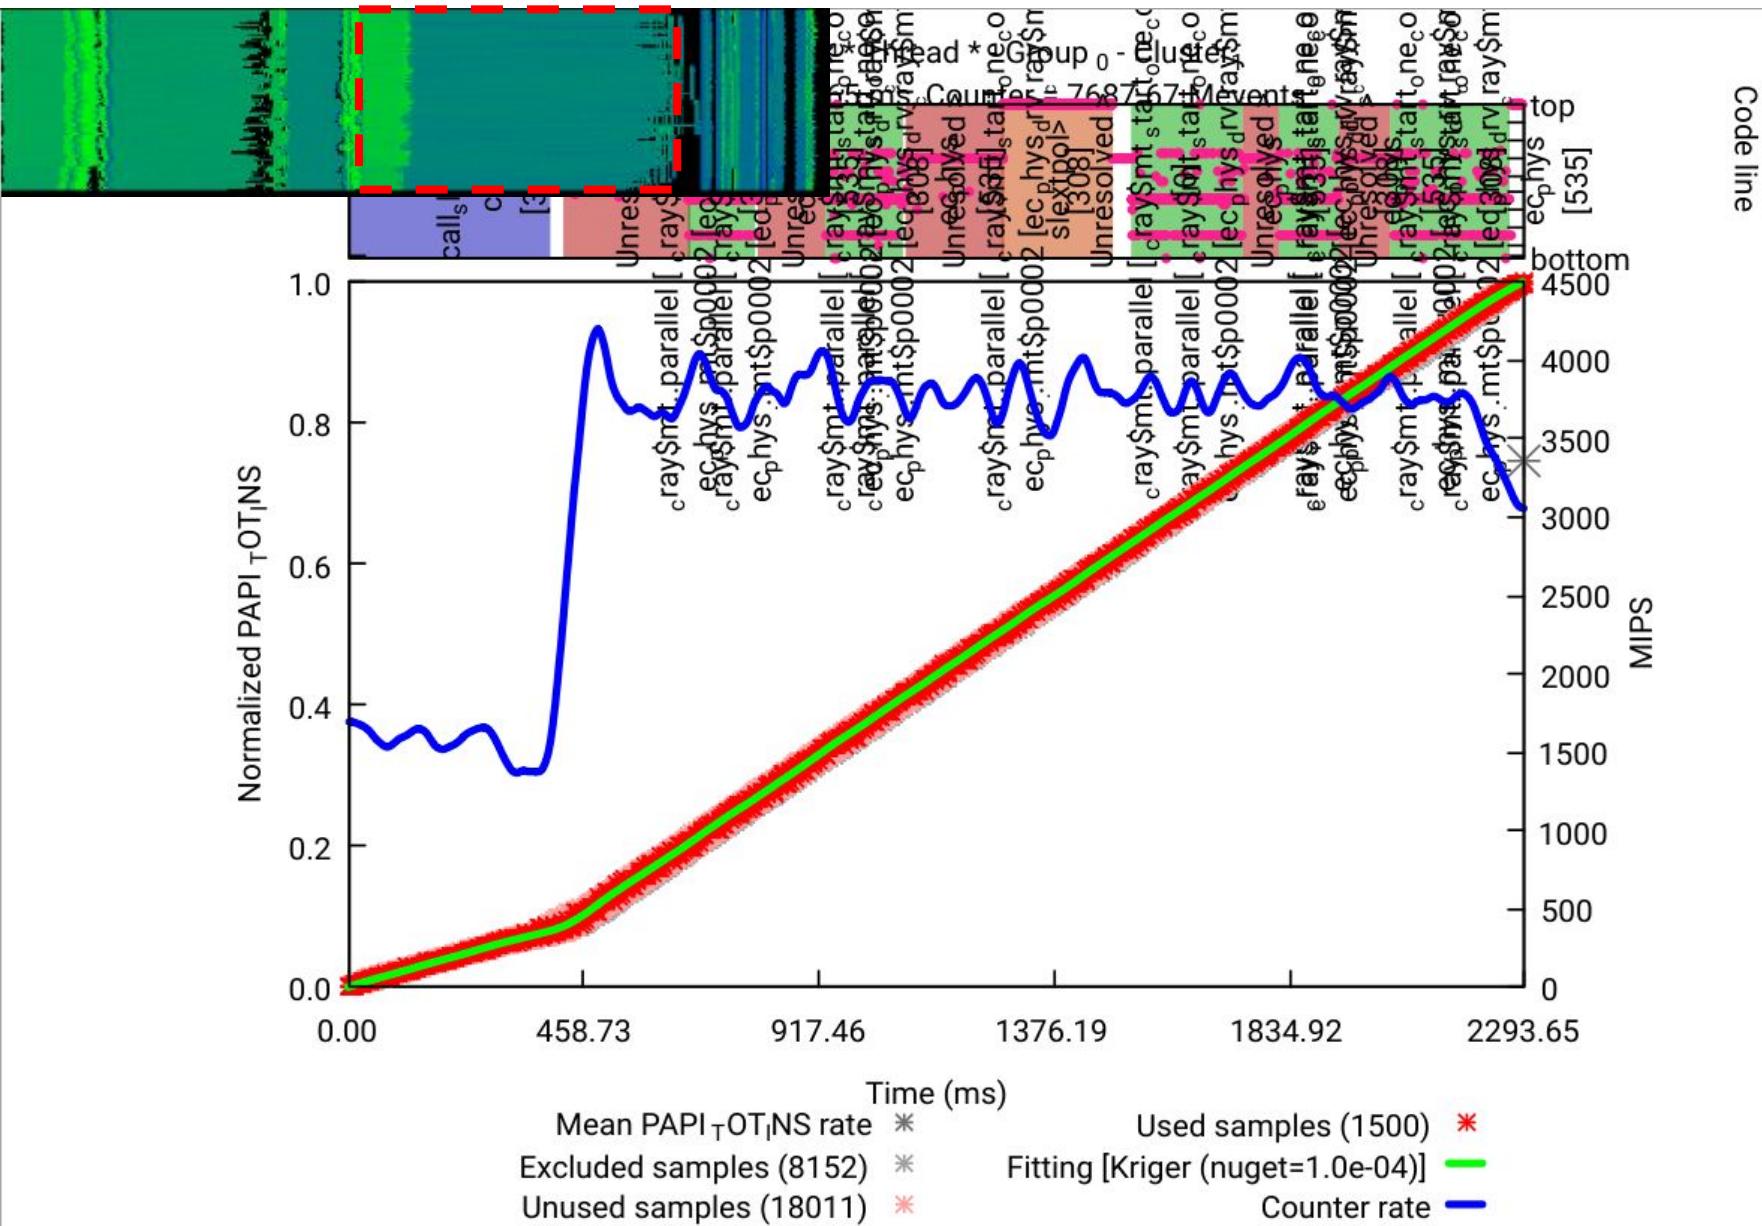
# Folding Tool

- Combine instrumentation and sampling to provide **instantaneous performance metrics**, source code and memory references. This mechanism receives a trace-file and generates plots showing the fine evolution of the performance.
- The samples collected are gathered from scattered computing regions into a synthetic region by preserving their relative time within their original region so that the sampled information determines how the performance evolves within the region.
- The performance evolution is connected to source code and memory references at the same time.

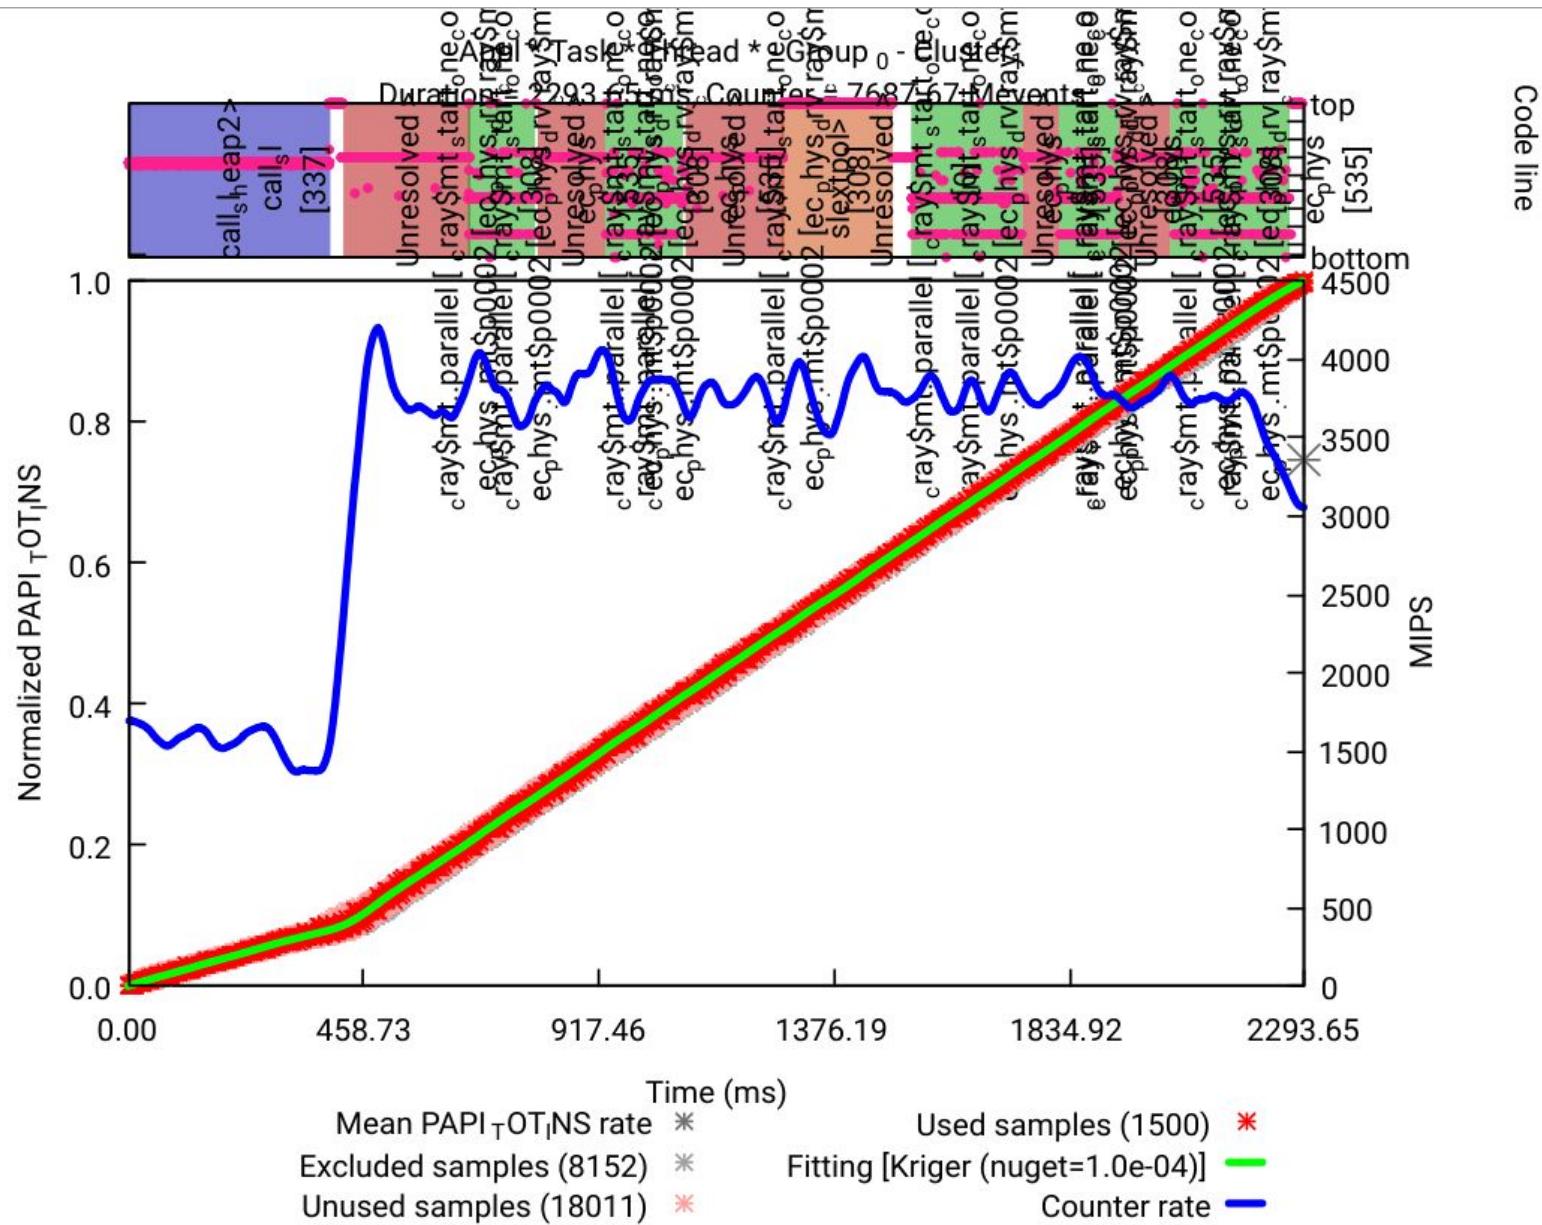
# Folding Tool



# Folding Tool



# Folding Tool



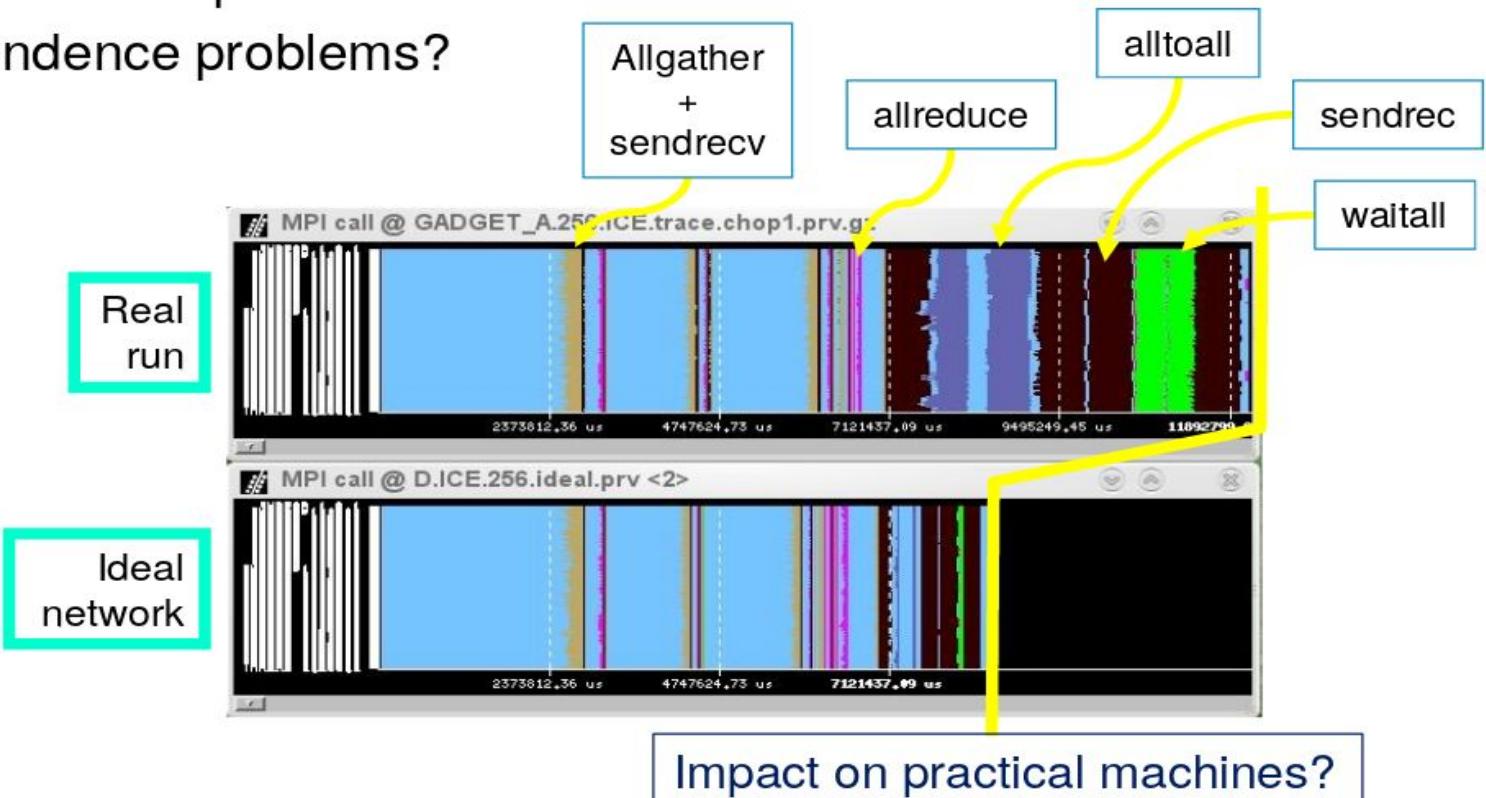
# Folding Tool



# DIMEMAS Tool

The impossible machine:  $BW = \infty$ ,  $L = 0$

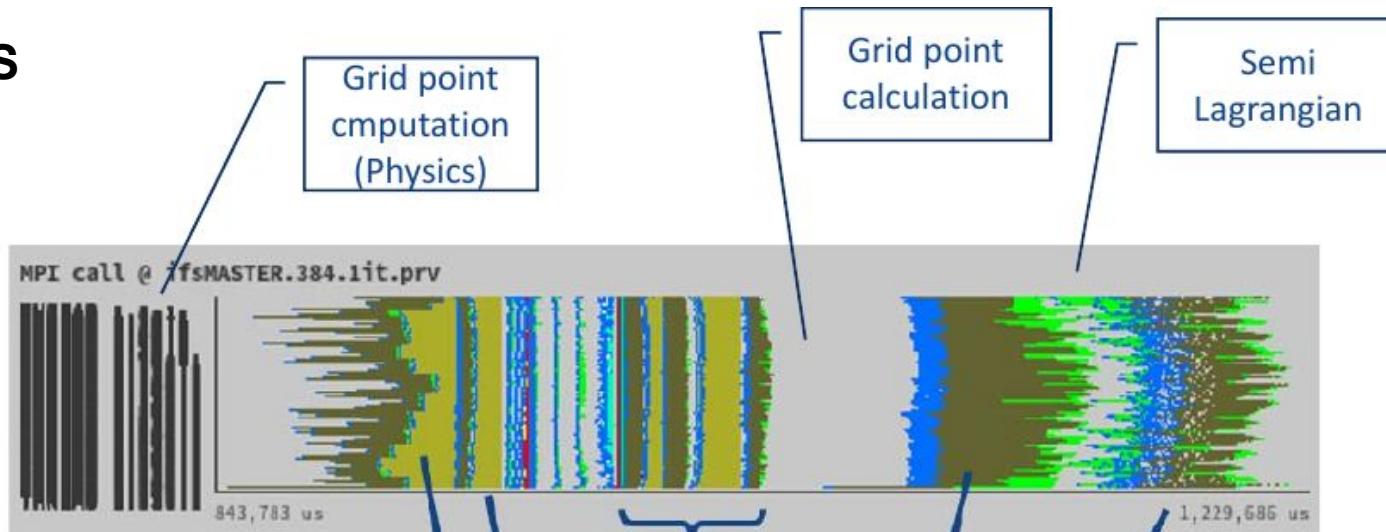
- Actually describes/characterizes intrinsic application behavior
  - Load balance problems?
  - Dependence problems?



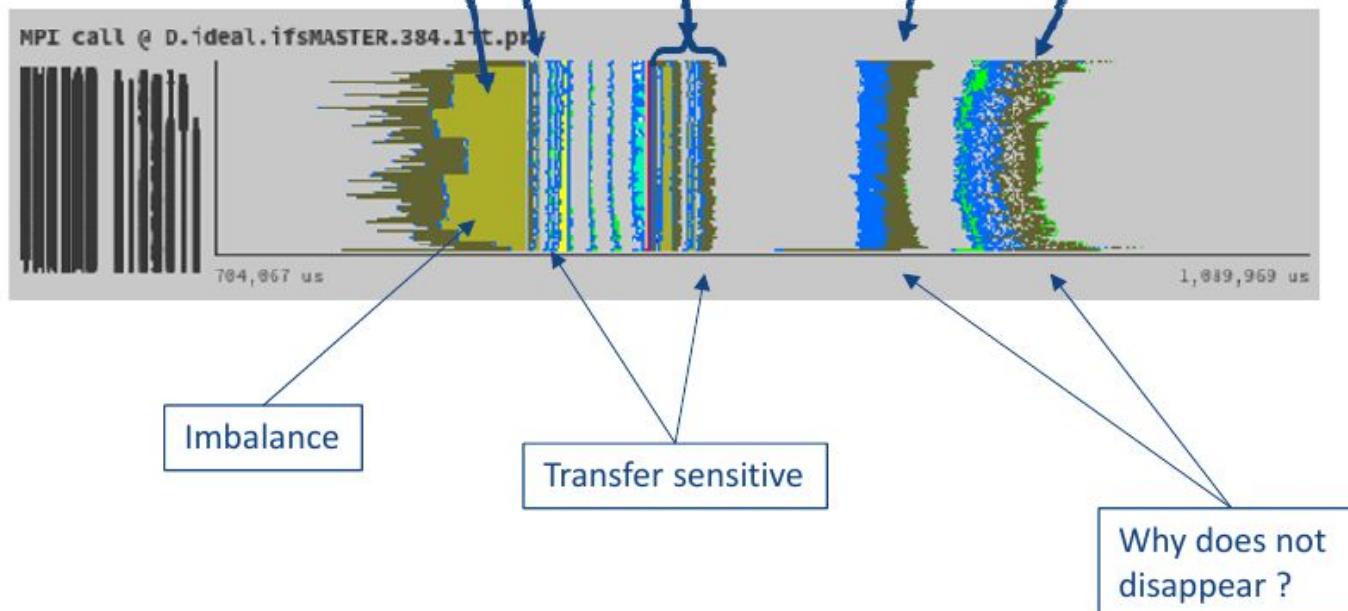
# DIMEMAS Tool

Ideal Network for IFS execution

- Actual run



- Ideal network



# Profiling Methodology

- Area of study
- Deployment efficiency
- Benchmarking
- Profiling analysis
- **Validation**
  - Reproducibility Test
  - Validation Test

# Validation

Reproducibility Test: Are your results comparable to the EC-Earth community results?

The Test proposed:

20-yr long, 5-member, Forcing Fixed Cmip and Amip simulations

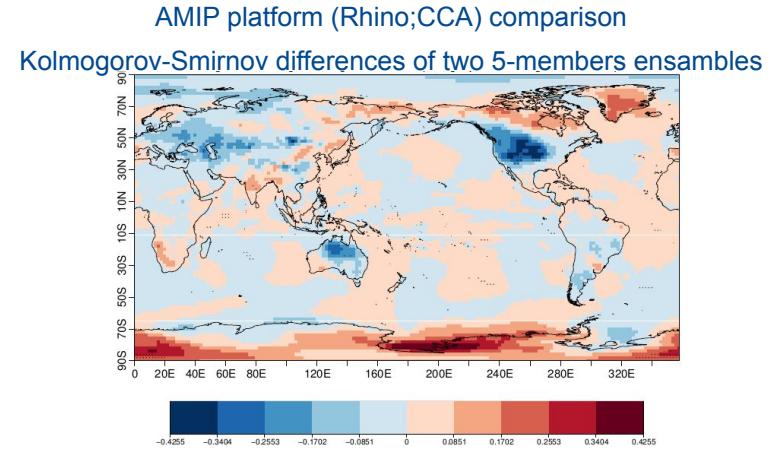
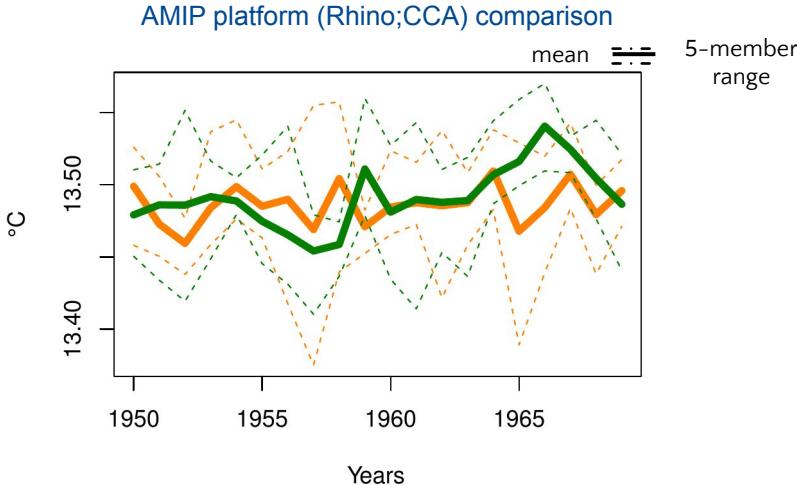
Allows to look at impact of machine on mean state/bias (not possible in the case of 1-yr simulations)

Allows to measure differences due hardware as compared to internal variability

Working under stationary conditions removes possible dependence of hardware impact on the mean state

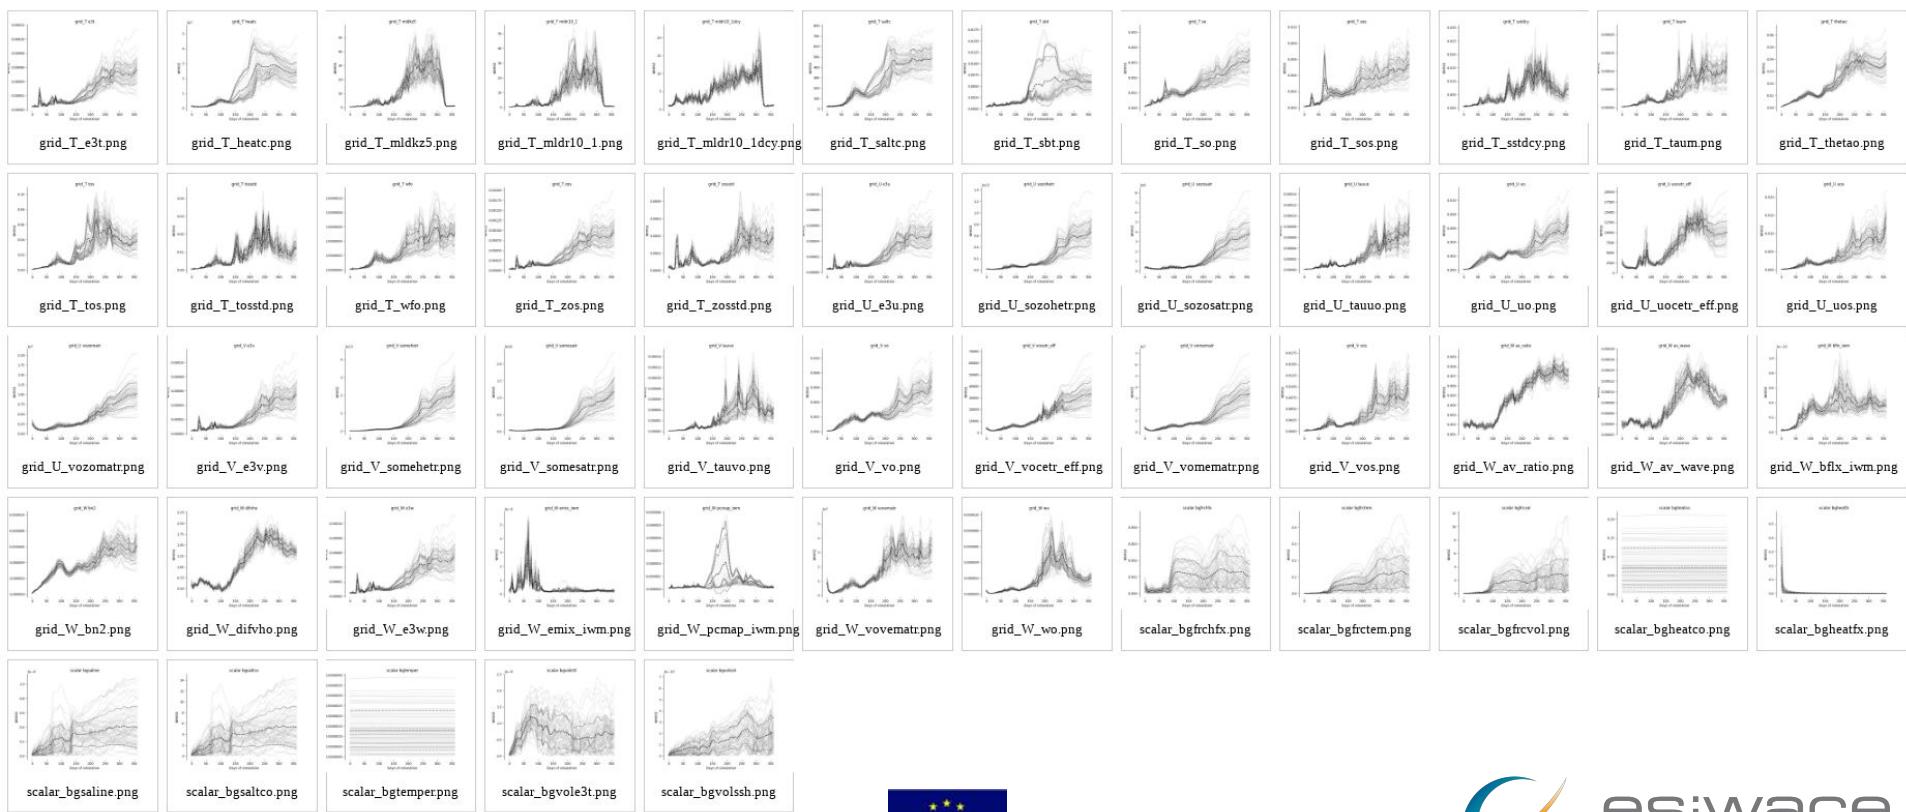
Addresses the problem from a global point-of-view; suitable to give recommendations for CMIP6

The results comparing platforms or configurations:



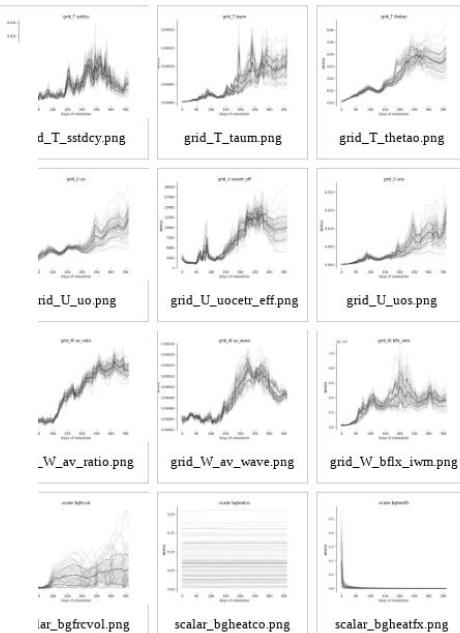
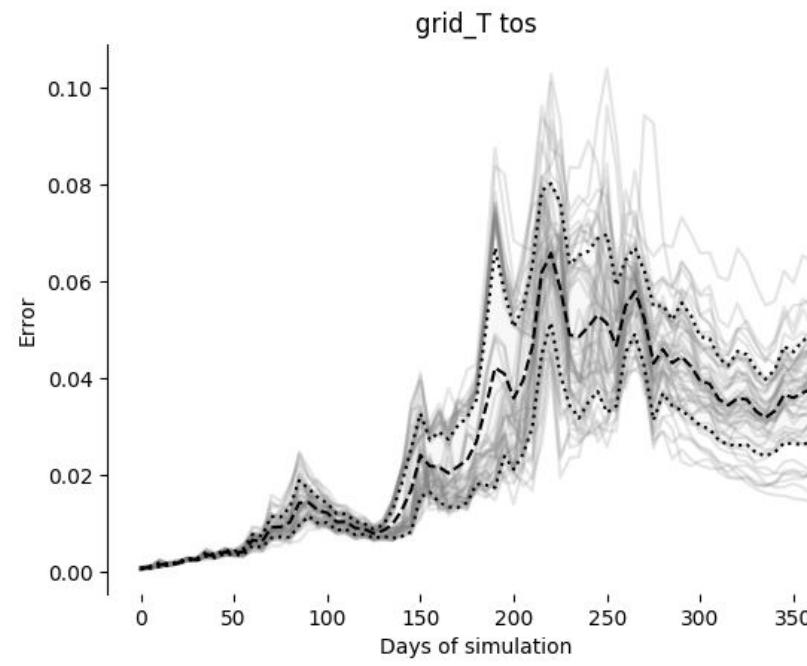
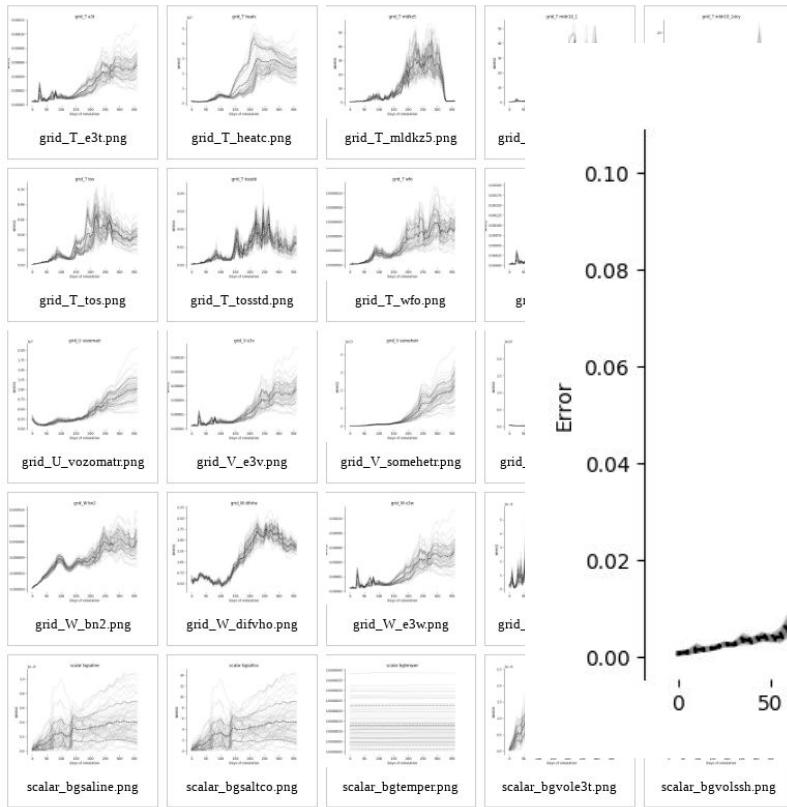
# Validation Test (NEMO)

- Initial conditions perturbed with white noise in the 3D temperature field.
- Evaluating 53 output variables.



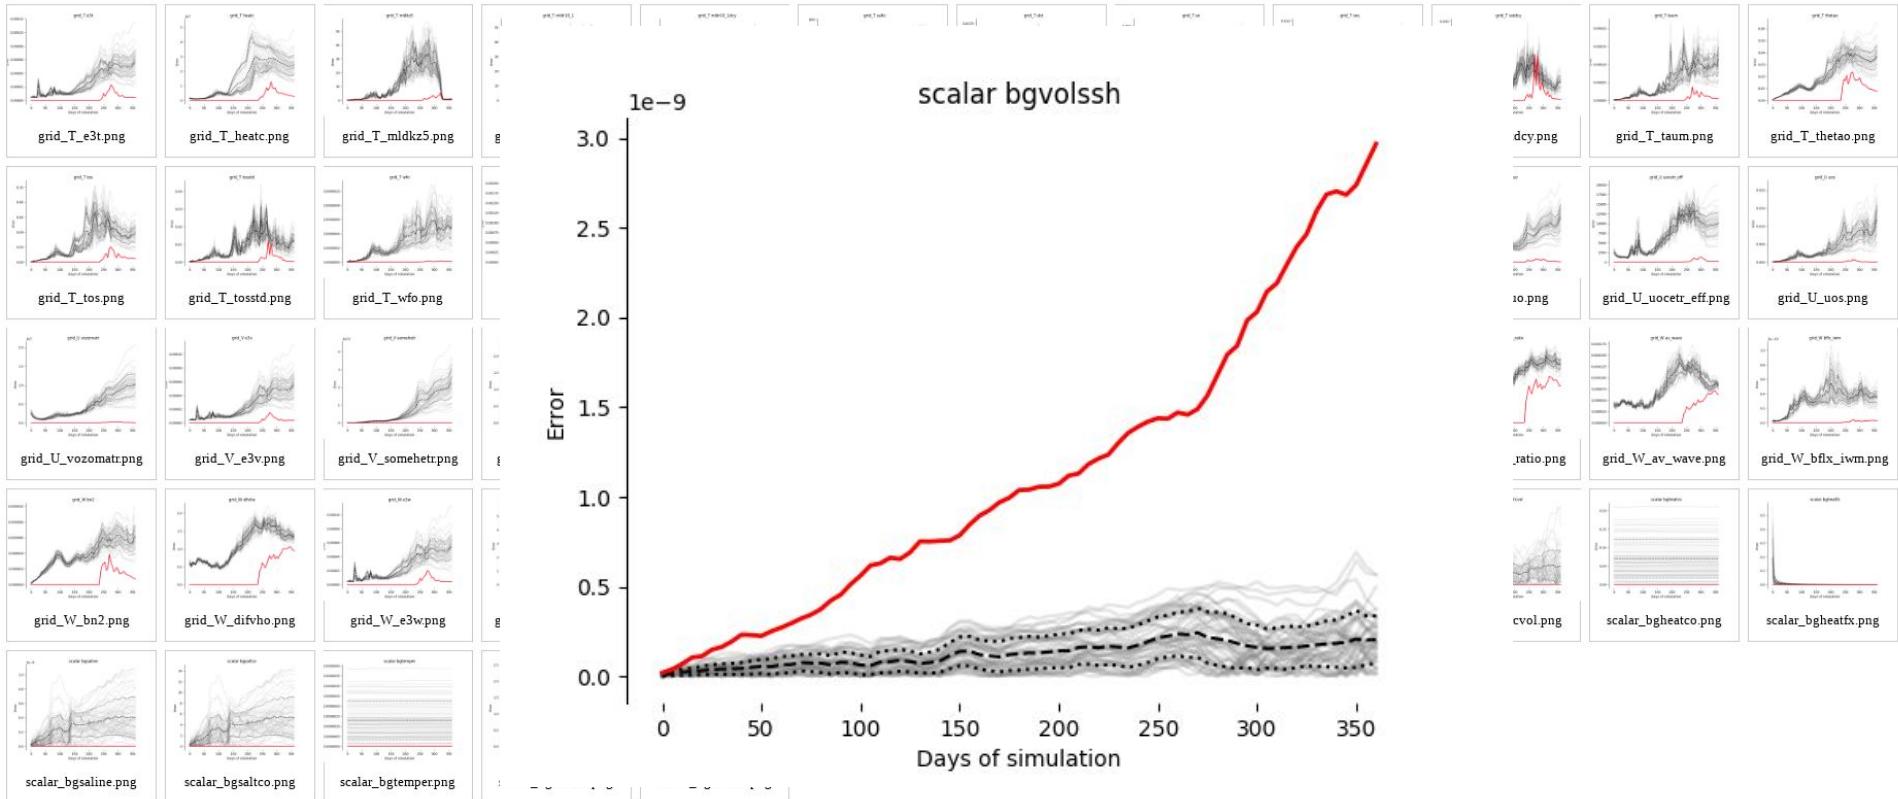
# Validation Test (NEMO)

- Initial conditions perturbed with white noise in the 3D temperature field.
- Evaluating 53 output variables.



# Validation Test (NEMO)

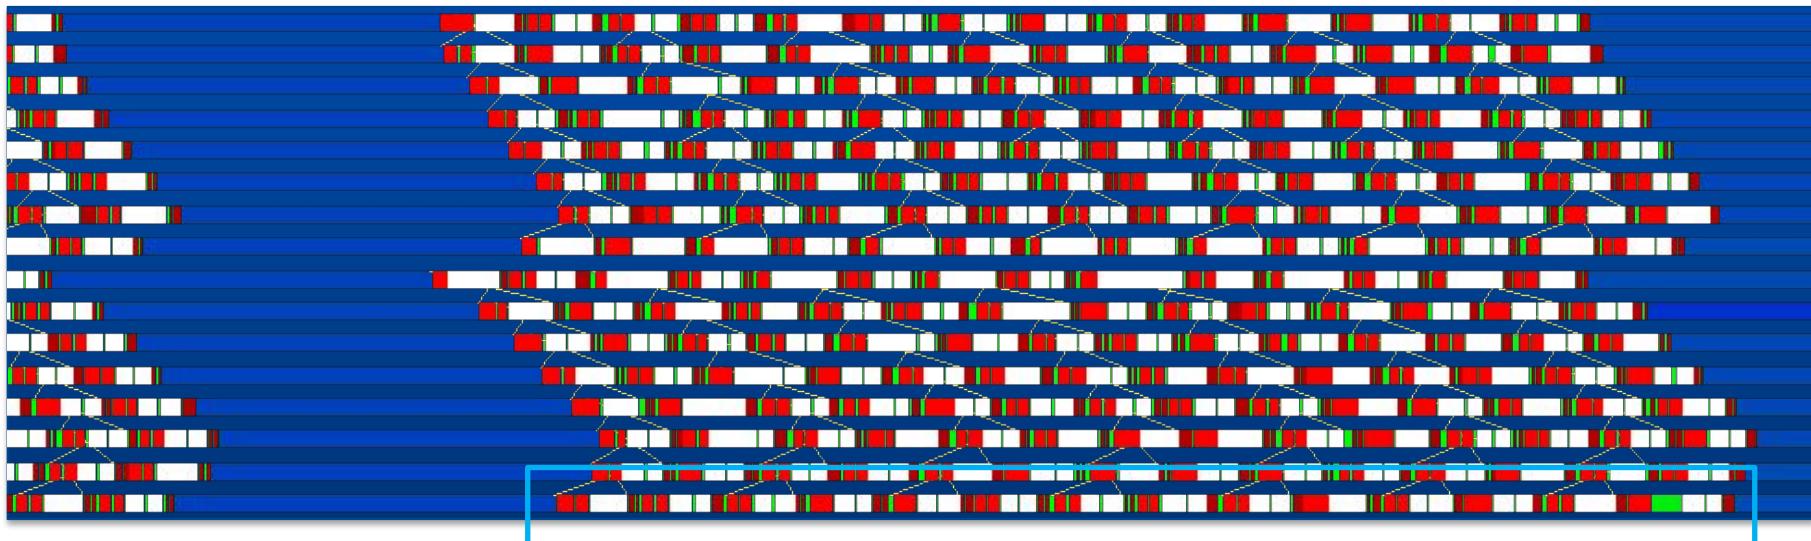
Example: Compiling with -xHost



# Examples



# Examples



Border Exchange

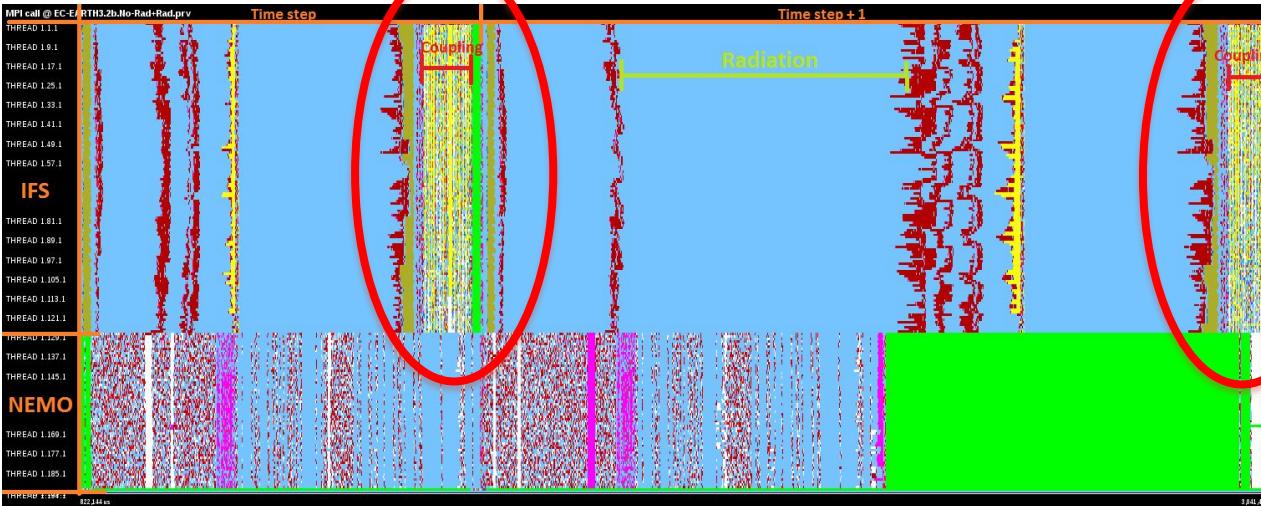
# Examples

- Diagnostic for NEMO:
  - Scalability is constrained by:
    - 1) Algorithms with too much communication
    - 2) Sub-optimal implementation
- Actions taken
  - Improve communication implementation to reduce number of point-to-point messages
  - Reduce number of collectives

# Examples

- First studies showed that IFS-NEMO coupling was not a big issue

IFS: 128 cores

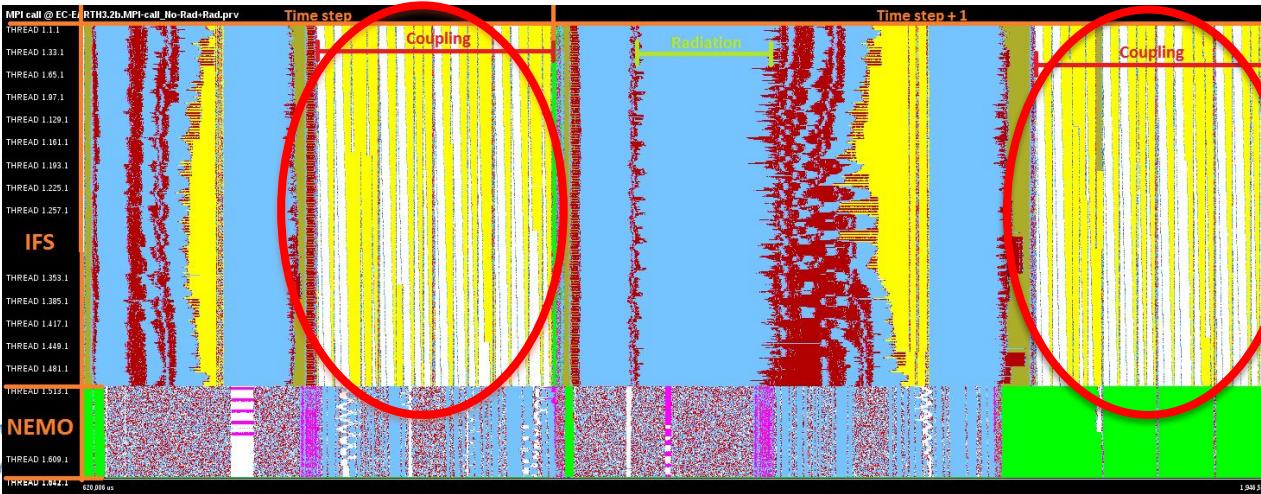


NEMO: 128 cores

Coupling time  
is increased

- But it seems that it is when increasing number of cores

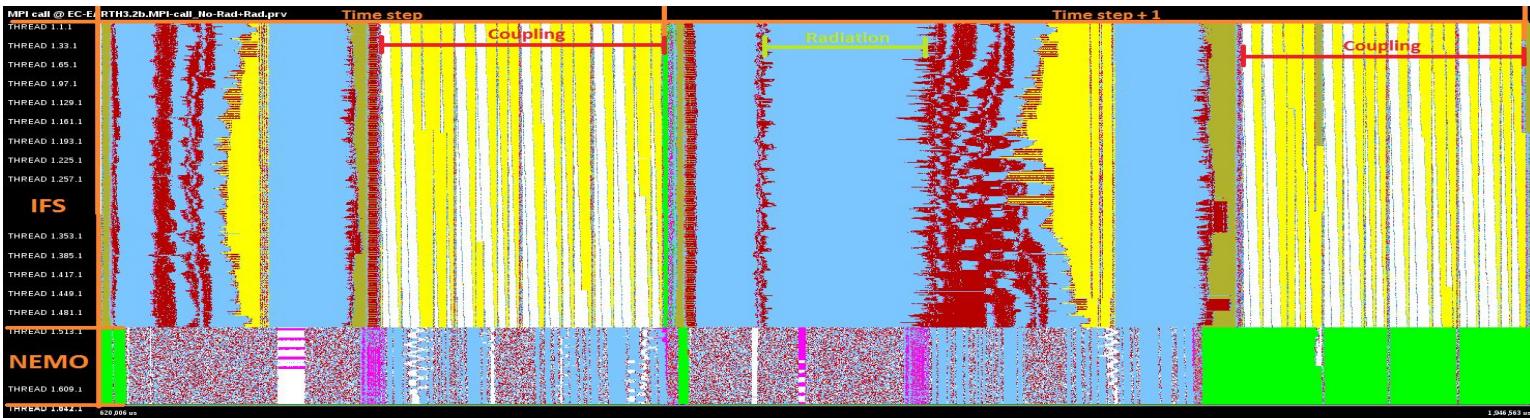
IFS: 512 cores



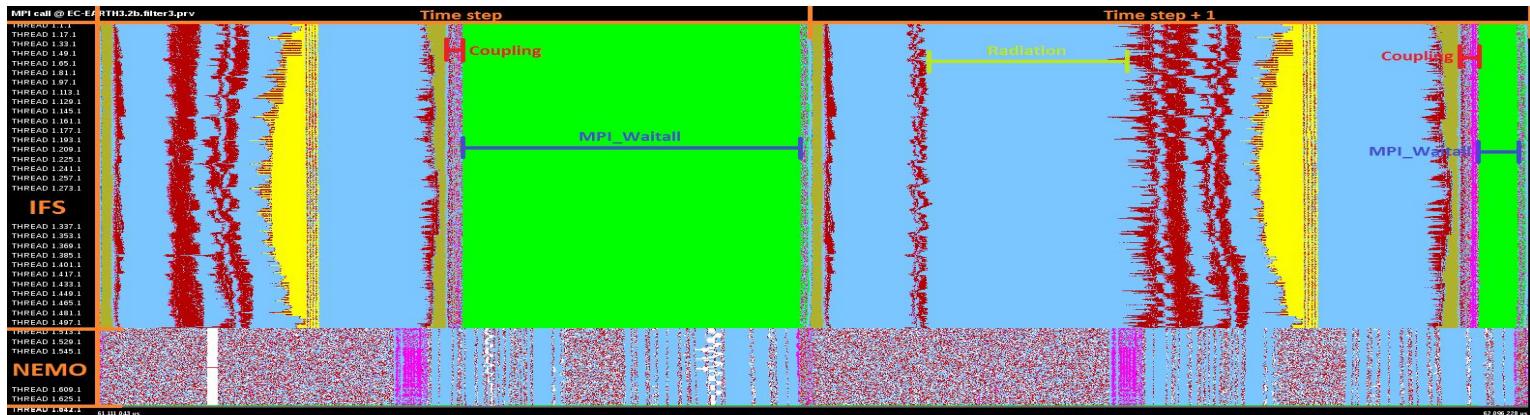
NEMO: 128 cores

# Examples

- BSC has been working successfully with the EC-Earth Technical Working Group to improve the execution of the model



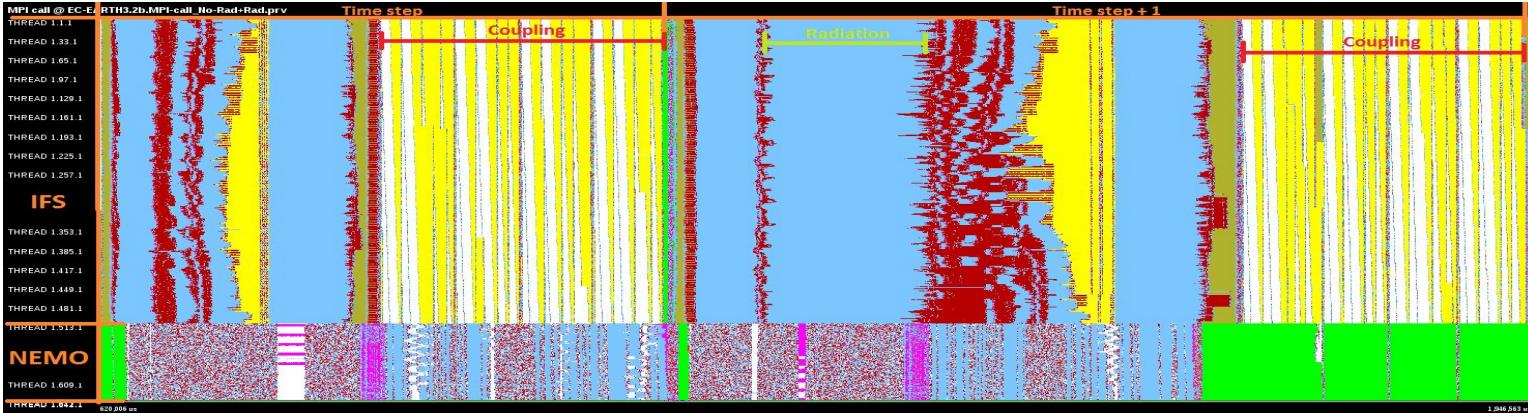
- A success case: coupling field gathering and OPT option of OASIS coupler for global conservative transformations



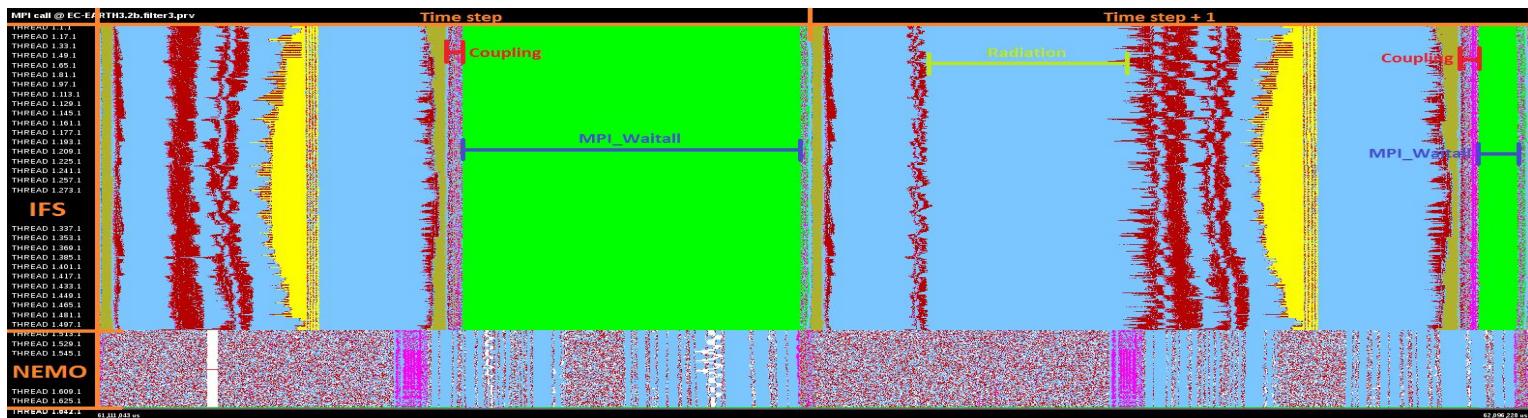
- With these optimizations, up to 90% improvement in coupling process can be achieved
- These improvements are now in trunk EC-Earth 3.2.2, substantially benefiting our CMIP6 simulations

# Examples

- BSC has been working successfully with the EC-Earth Technical Working Group to improve the execution of the model



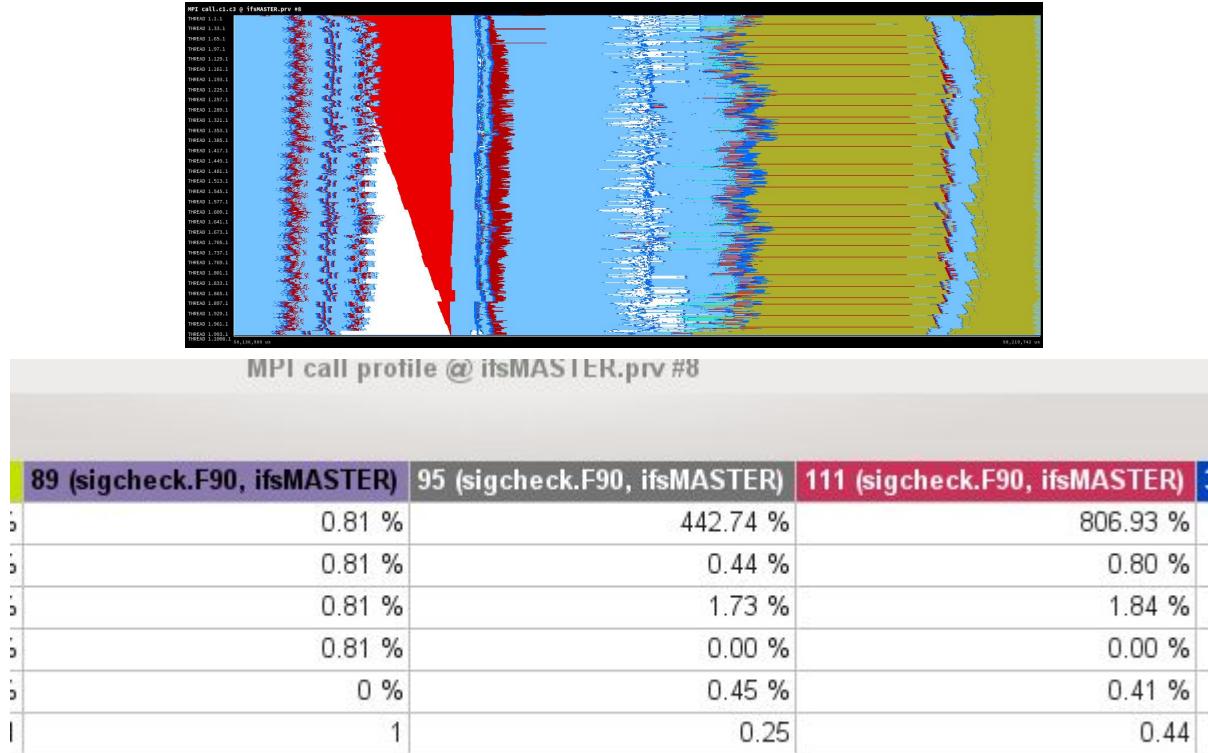
- A success case: coupling field gathering and OPT option of OASIS coupler for global conservative transformations



- With these optimizations, up to 90% improvement in coupling process can be achieved
- These improvements are now in trunk EC-Earth 3.2.2, substantially benefiting our CMIP6 simulations

# Examples

- Synchronal point to point communication could be a bottleneck even for only one message from one master to hundreds of slaves
  - Sigcheck method

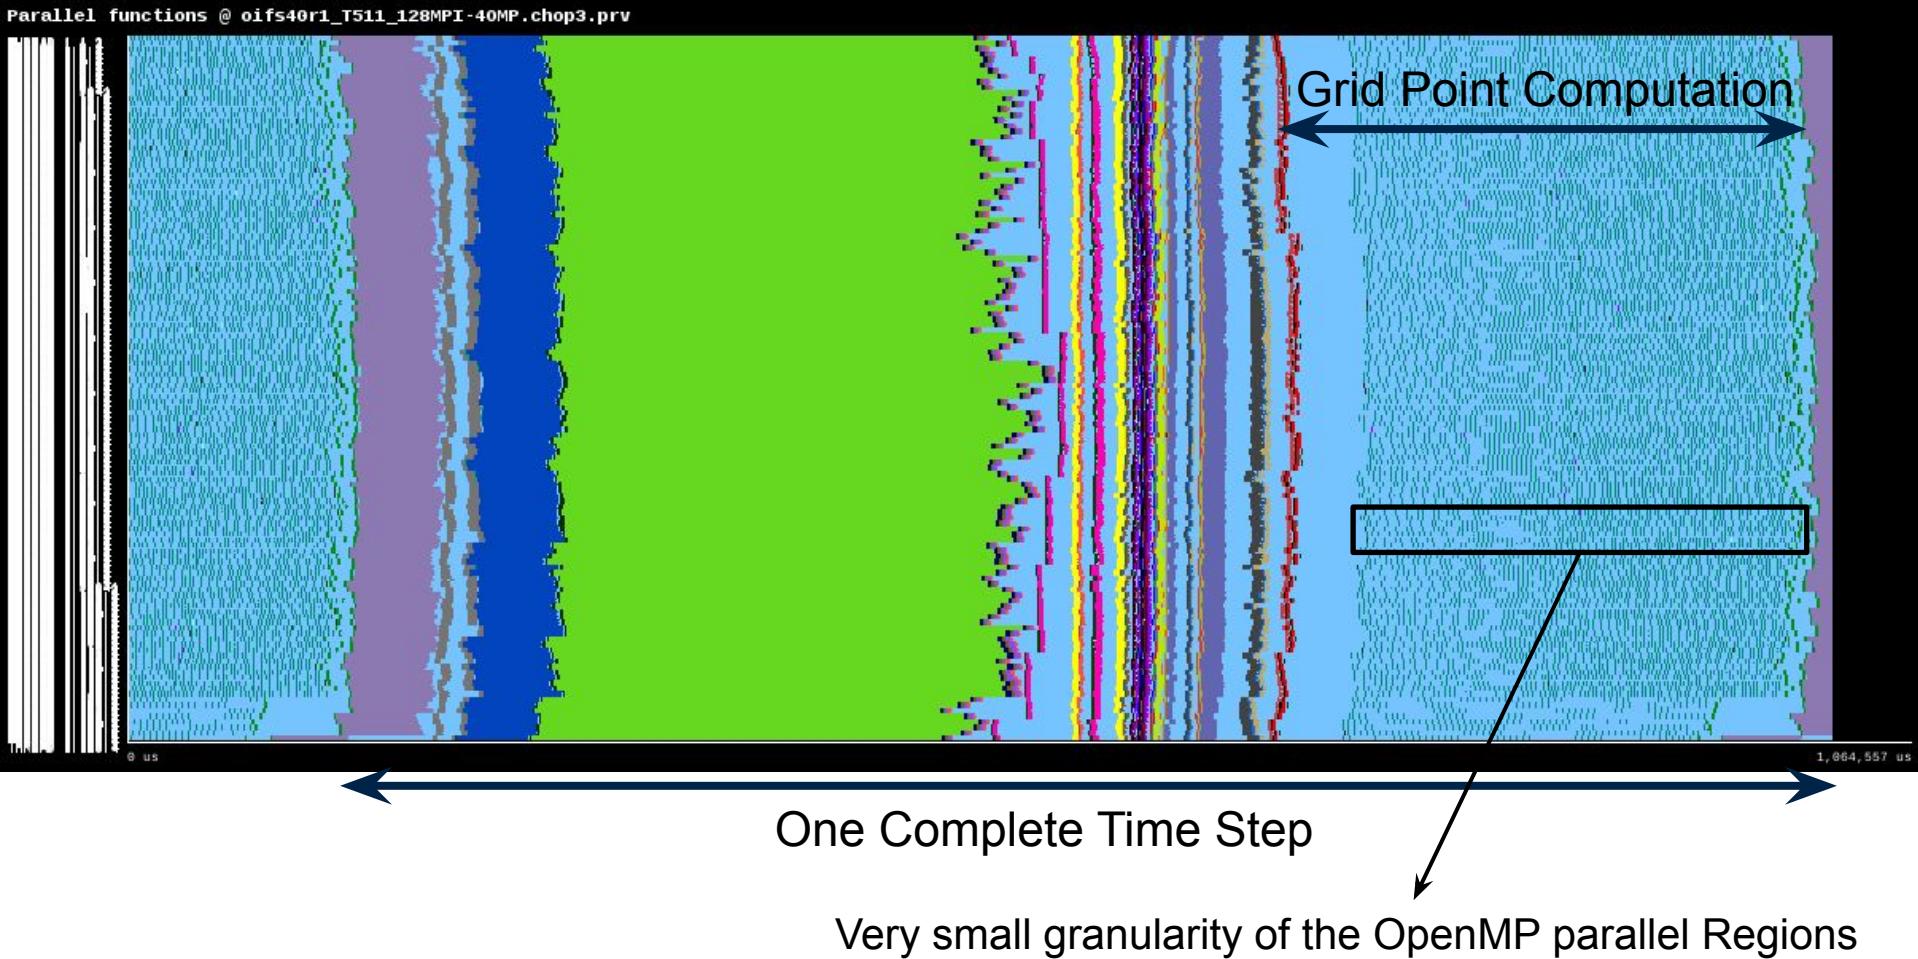


- Using one asynchronal collective communication this time is reduced almost to 0

# Examples

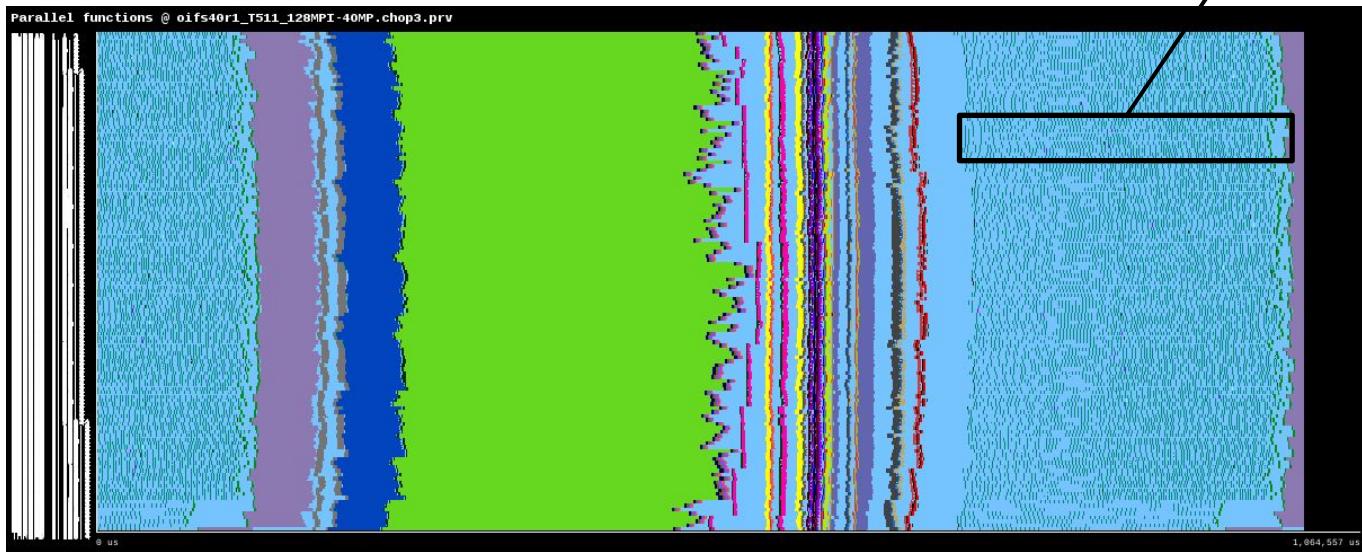
- Hybrid Test (128 MPI+4 OpenMP, Total: 512)

## OpenMP Parallel Regions

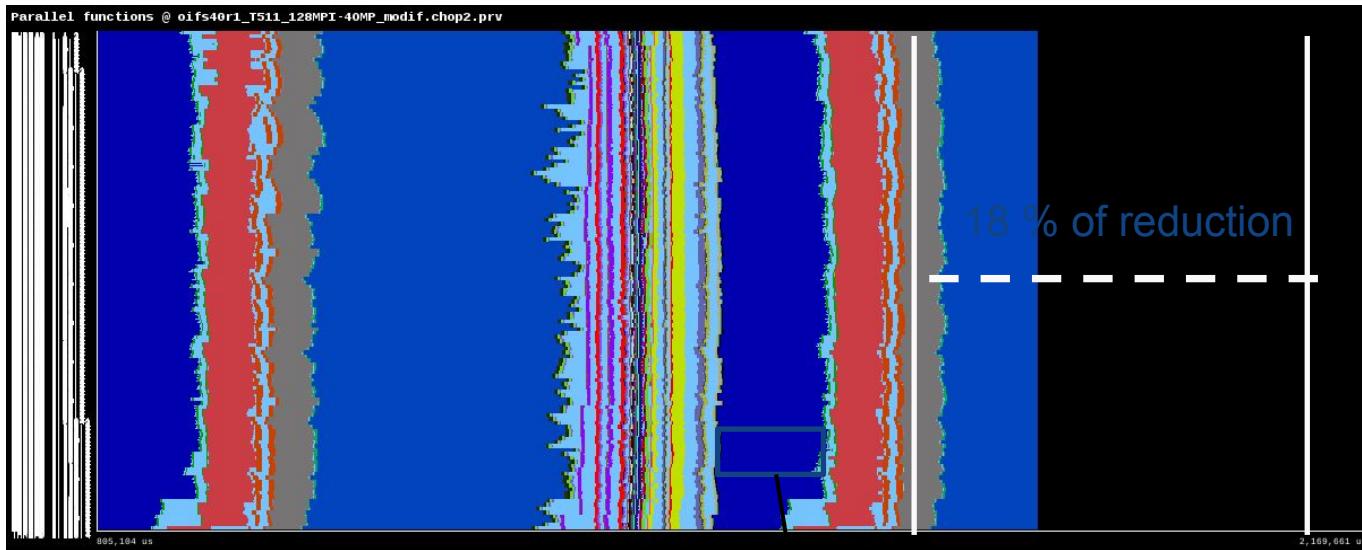


# Examples

128 MPI processes  
and 4 OpenMP  
threads per process



128 MPI processes  
and 4 OpenMP  
threads per process



Only one coarse OpenMP parallel region



**Barcelona  
Supercomputing  
Center**  
*Centro Nacional de Supercomputación*



EXCELENCIA  
SEVERO  
OCHOA



**esiwace**

CENTRE OF EXCELLENCE IN SIMULATION OF WEATHER  
AND CLIMATE IN EUROPE

# Thank you

*The research leading to these results has received funding from the EU H2020 Framework Programme under grant agreement H2020 GA 675191.*

*The content of this presentation reflects only the author's view. The European Commission is not responsible for any use that may be made of the information it contains.*

[mario.acosta@bsc.es](mailto:mario.acosta@bsc.es)